



Malloc: Gear Up

Overview

- Dynamic Storage Allocator
 - Create storage as needed
 - Use a doubly-linked explicit free list to manage free blocks
- Functions You Will Write
 - `mm_init()`, `mm_malloc()`, `mm_free()`, `mm_realloc()`
 - `mm_check_heap()`

Roadmap

Install: `cs0330_install_malloc`


Coding

- `mm_init()`
- `mm_check_heap()`
- `mm_malloc()`
 - `mm_extend_heap()`
- `mm_free()`
 - `coalesce()`
- `mm_realloc()`

To Run: `./mdriver -V`



Blocks: Structure



Header: Contains the **size** of the block (note: size is inclusive of the header and footer) and a **boolean** representing whether the block is allocated or free

Payload: For allocated blocks, this is where the data is stored

Footer: Same as the header. We use it to coalesce free blocks.

Free Blocks: Structure

8	Size	0
8	Next	
8	Prev	
Div 8	Payload and Padding	
8	Size	0

Next and Prev: Free-list pointers

- **Next:** Points to the next free block in the free list
- **Prev:** Points to the previous free block in the free list

In our implementation, the free list is a **circularly, doubly linked list**. What does this mean for prev and next pointers in an empty list? single-item list? n-item list?

Blocks: Struct



```
typedef struct block {  
    size_t size;  
    size_t payload[0];  
} block_t;
```



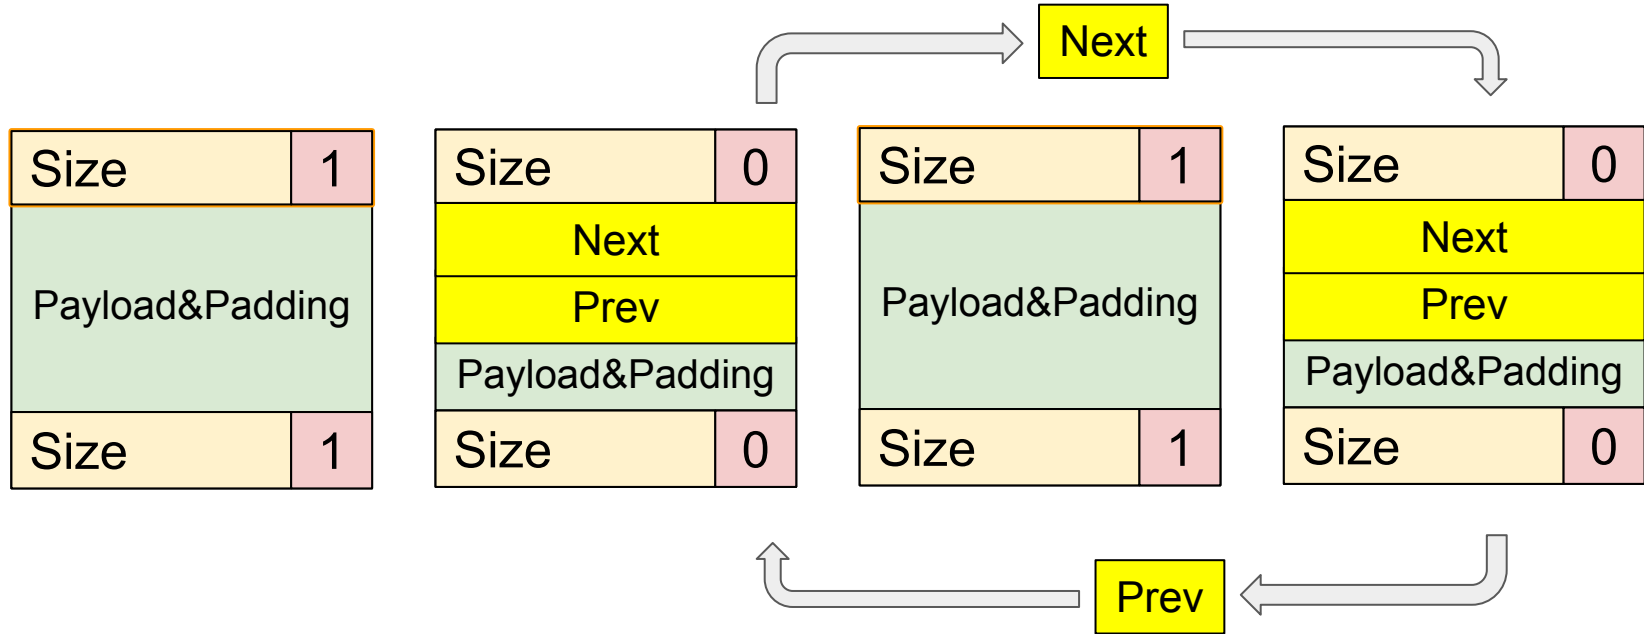
- The `payload[0]` means that the array can be of any length
- Use the inline functions defined for you in `mminline.h` to manipulate blocks.

Discussion

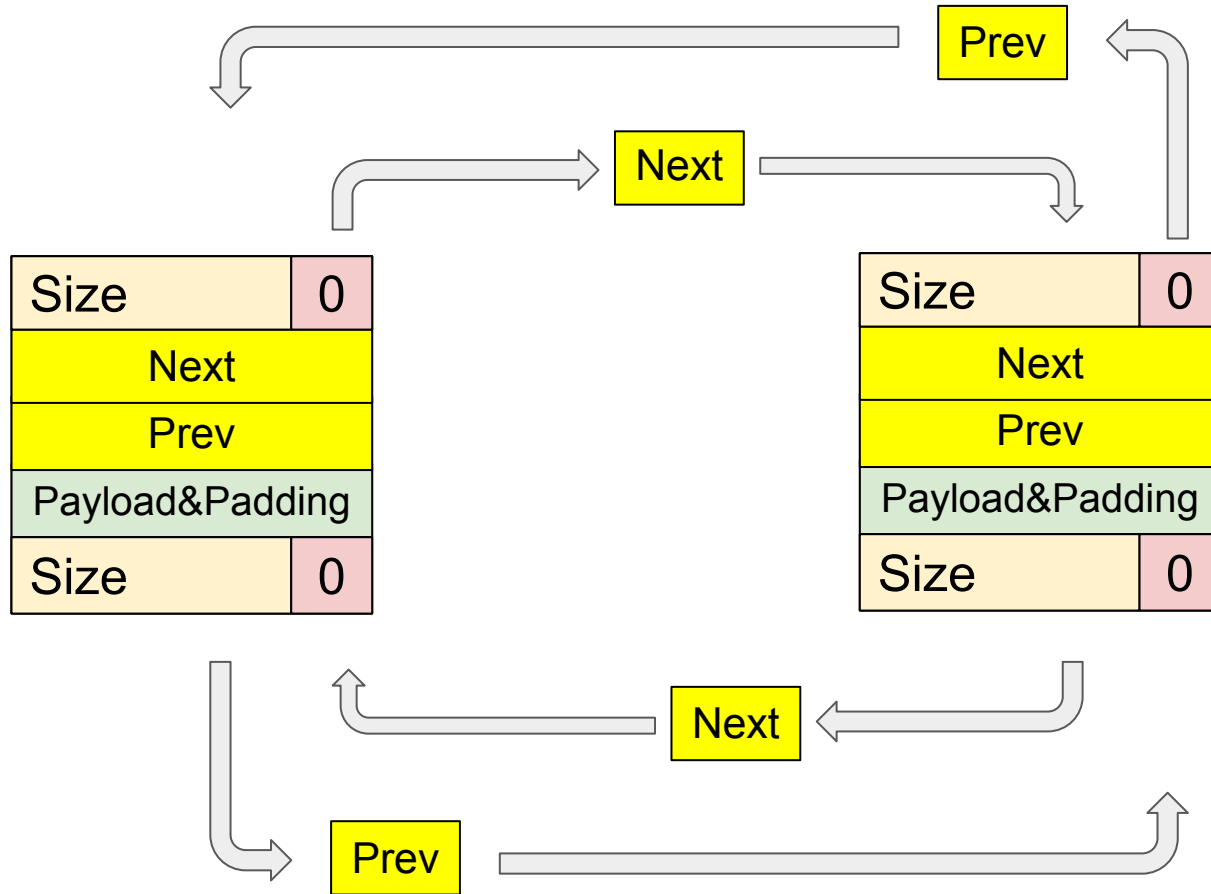
What do you think the minimum block size is?



The Free List, a visual overview



The Free List, a visual overview



Malloc

Implementing a successful malloc!



What functions do you need to implement?

`mm_init(void)`: initializes the dynamic storage allocator

`mm_malloc(size_t size)`: allocates a block of memory, returning a pointer to that block's payload

`mm_free(void *ptr)`: frees a block of memory. This block can now be reused!

`mm_realloc(void *ptr, size_t size)`: reallocates a memory block to update it with a new size

`mm_check_heap(void)`: checks the heap for consistency

`mm_init(void)`

- Start by building prologue and epilogue blocks
- What type should these be?
- Should these be allocated or free?
- What other qualities about these blocks should you set?
- Remember to error check!
- Are there other objects that need to be set?

`mm_extend_heap(size_t size)`

What does extending the heap mean?

- Create a new free block
 - Break more memory in the heap
 - Error check
 - Set free block's attributes
 - Check: Is this free block next to other free blocks? If so, what should we do?
- Overwrite the epilogue
 - Reset the epilogues attributes
- Examine the input
 - What happens if the size input isn't a multiple of 8?
 - What can we do to ensure that it will be?



`mm_malloc(size_t size)`

- Find the first free block in the list
- What do you do if you can't find a fit?
- Once you find a fit, decide whether or not you have to split the free block you want to allocate
- Check for edge cases (very important!!!!)
 - What if the caller tries to allocate memory of size 0?
 - What if the size isn't aligned to 8 bytes?
 - What if you have to split the block?

`mm_free(void *ptr)`

- Retrieve the pointer to the header of the block
- Mark the header and footer as free (using the last bit in each)
- Remember to coalesce if necessary!



`mm_realloc(void *ptr, size_t size)`

What could you consider in implementing realloc?

- Cover the edges
 - What if the pointer is null?
 - What does it mean if the size is 0?
- Create a block for the pointer
 - What would it mean if this block was free? Allocated?
- Sizing: the original block had a size, the user has requested a new size
 - How should we keep track of these?
 - Is the requested size appropriately formatted? How do we know? What can we do if it isn't?
- Neighboring blocks
 - Are these free or allocated? What does that mean for our new block?

`mm_realloc(void *ptr, size_t size)`

What could you consider in implementing realloc?

- Backing up our data to the stack
 - Max memory realloc.rep asks for is 615000
 - Using memcpy: What is the size we should copy over? How could we save space?
- Set our block free!
- Maintaining the free list
 - How can we avoid undefined behavior when splitting or coalescing?

`mm_realloc(void *ptr, size_t size)`

Fit

- How do we determine if we fit?
- If we do fit, what do we do?
 - Get the new address
 - What should we consider about the size that's left over?
 - What if we're at the same address as before? What should we copy?
- If we don't fit, what do we do?
 - Take from the free list
 - There are many different ways to do this!

mm_check_heap(void)

- What would be good to check in this function?
 - Is everything aligned properly?
 - How can we keep track of the free blocks?
 - Is every free block actually in the free list?
 - Check free list has valid free blocks: is every block in the free list marked as free?
 - Are there any contiguous free blocks that somehow escaped coalescing?
 - Do the pointers in the free list point to valid free blocks?
 - Can you access the prologue and epilogue blocks properly?
- Use assert statements
- **Must have a viable mm_check_heap() before a TA can help you debug heap corruption bugs**

Where should I start?

Understand what happens to the blocks using the inline functions

Answer the questions posed in these slides

Coding


- Start with `mm_init`, `mm_extend_heap`
- Move on to `malloc`, `free`, `realloc`, and `coalesce`
- Implement `mm_check_heap` as you go along!!



Helpful GDB Tips!

Look at what is pointed to by the pointer `blk` by typing `print *blk` or `p *blk` in GDB

- This dereferences the pointer and lets you see the size and allocation status
- If it's allocated, the size will show up as a multiple of 4 plus 1, otherwise just a multiple of 4



REPL

- mdriver comes with a REPL you can use to test your malloc
- See section 4 of the handout!
- Run with `./mdriver -r`

```
Welcome to the Malloc REPL. (Enter 'help' to see available commands.)
> help
commands:
malloc <index> <size>    mallocs the block at <index> to a size <amount>
realloc <index> <size>  reallocs the block at <index> to <amount>
free <index>             frees block at <index>
check_heap              calls mm_check_heap() to test validity
print                   prints the heap
print -b <index>        prints the status of the block at <index>
quit                    quits repl

> malloc 1 99
> malloc 2 100
> print
heap size: 544
prologue                block at 0x7f7934a91010    size 16
free block              block at 0x7f7934a91020    size 272    Next: 0x7f7934a91020
block[2] allocated      block at 0x7f7934a91130    size 120
block[1] allocated      block at 0x7f7934a911a8    size 120
epilogue                block at 0x7f7934a91220    size 16

> print -b 2
block[2] allocated      block at 0x7f7934a91130    size 120
> free 2
```

Traces

A trace file is an ASCII file. It begins with a 4-line header:

```
<sugg_heapsize> /*suggested heap size (unused)*/  
<num_ids>        /*number of request id's*/  
<num_ops>        /*number of requests  
<operations>
```

2000

6

12

1

a 0 2040

a 1 2040

f 1

a 2 48

a 3 4072

f 3

a 4 4072

f 0

f 2

a 5 4072

f 4

f 5

Traces, continued

The header is followed by `num_ops` text lines. Each line denotes either an `allocate[a]`, `realloc[r]`, or `free[f]` request. The `<alloc_id>` is an integer that uniquely identifies an `allocate` or `realloc` request.

```
a <id> <bytes>    /*ptr_<id> = malloc(<bytes>)*/  
r <id> <bytes>    /*realloc(ptr_<id>, <bytes>)*/  
f <id>           /*free(ptr_<id>)*//
```

2000

6

12

1

a 0 2040

a 1 2040

f 1

a 2 48

a 3 4072

f 3

a 4 4072

f 0

f 2

a 5 4072

f 4

f 5

Malloc

You can do this,
and we're here
to help!

