

CS 33

Data Representation (Part 3)

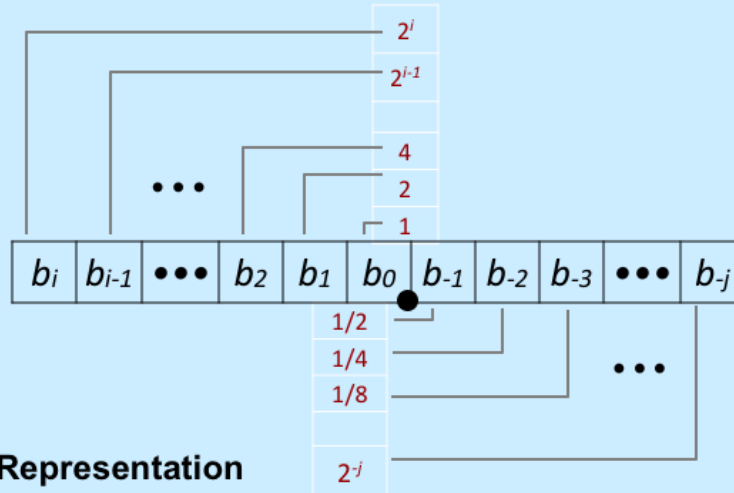
Many of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective.” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Fractional binary numbers

- What is 1011.101_2 ?

Supplied by CMU.

Fractional Binary Numbers



• Representation

- bits to right of “binary point” represent fractional powers of 2
- represents rational number:

$$\sum_{k=-j}^i b_k \times 2^k$$

Supplied by CMU.

Representable Numbers

- **Limitation #1**

- can exactly represent only numbers of the form $n/2^k$
 - » other rational numbers have repeating bit representations
- value representation
 - » 1/3 0.0101010101[01]...₂
 - » 1/5 0.001100110011[0011]...₂
 - » 1/10 0.0001100110011[0011]...₂

- **Limitation #2**

- just one setting of decimal point within the w bits
 - » limited range of numbers (very small values? very large?)

Supplied by CMU.

IEEE Floating Point

- **IEEE Standard 754**
 - established in 1985 as uniform standard for floating point arithmetic
 - » before that, many idiosyncratic formats
 - supported by all major CPUs
- **Driven by numerical concerns**
 - nice standards for rounding, overflow, underflow
 - hard to make fast in hardware
 - » numerical analysts predominated over hardware designers in defining standard

Supplied by CMU.

Floating-Point Representation


- Numerical Form:

$$(-1)^s M 2^E$$

- sign bit s determines whether number is negative or positive
- significand M normally a fractional value in range $[1.0, 2.0)$
- exponent E weights value by power of two

- Encoding

- MSB s is sign bit s
- exp field encodes E (but is not equal to E)
- frac field encodes M (but is not equal to M)




Precision options

- **Single precision: 32 bits**



- **Double precision: 64 bits**



- **Extended precision: 80 bits (Intel only)**



Supplied by CMU.

On x86 hardware, all floating-point arithmetic is done with 80 bits, then reduced to either 32 or 64 as required.

“Normalized” Values

- When: $\text{exp} \neq 000\dots 0$ and $\text{exp} \neq 111\dots 1$
- Exponent coded as *biased* value: $E = \text{Exp} - \text{Bias}$
 - *exp*: unsigned value exp
 - $\text{bias} = 2^{k-1} - 1$, where k is number of exponent bits
 - » single precision: 127 (Exp: 1...254, E: -126...127)
 - » double precision: 1023 (Exp: 1...2046, E: -1022...1023)
- Significand coded with implied leading 1: $M = 1.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
 - minimum when $\text{frac} = 000\dots 0$ ($M = 1.0$)
 - maximum when $\text{frac} = 111\dots 1$ ($M = 2.0 - \epsilon$)
 - get extra leading bit for “free”

Supplied by CMU.

Normalized Encoding Example

- Value: float $F = 15213.0$;

$$\begin{aligned} - 15213_{10} &= 11101101101101_2 \\ &= 1.1101101101101_2 \times 2^{13} \end{aligned}$$


- Significand

$$\begin{aligned} M &= 1.\underline{1101101101101}_2 \\ \text{frac} &= \underline{1101101101101}0000000000_2 \end{aligned}$$

- Exponent

$$\begin{aligned} E &= 13 \\ \text{bias} &= 127 \\ \text{exp} &= 140 = 10001100_2 \end{aligned}$$

- Result:



Supplied by CMU.

Denormalized Values

- Condition: $\text{exp} = 000\dots 0$
- Exponent value: $E = -\text{Bias} + 1$ (instead of $E = 0 - \text{Bias}$)
- Significand coded with implied leading 0:
 $M = 0.\text{xxx}\dots\text{x}_2$
 - $\text{xxx}\dots\text{x}$: bits of frac
- Cases
 - $\text{exp} = 000\dots 0, \text{frac} = 000\dots 0$
 - » represents zero value
 - » note distinct values: $+0$ and -0 (why?)
 - $\text{exp} = 000\dots 0, \text{frac} \neq 000\dots 0$
 - » numbers closest to 0.0
 - » equispaced


Supplied by CMU.

Special Values

- **Condition: $\text{exp} = 111\dots 1$**
- **Case: $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$**
 - represents value ∞ (infinity)
 - operation that overflows
 - both positive and negative
 - e.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$
- **Case: $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$**
 - not-a-number (NaN)
 - represents case when no numeric value can be determined
 - e.g., $\text{sqrt}(-1)$, $\infty - \infty$, $\infty \times 0$

Supplied by CMU.

Visualization: Floating-Point Encodings



Supplied by CMU.

Tiny Floating-Point Example



- **8-bit Floating Point Representation**
 - the sign bit is in the most significant bit
 - the next four bits are the exponent, with a bias of 7
 - the last three bits are the *frac*
- **Same general form as IEEE Format**
 - normalized, denormalized
 - representation of 0, NaN, infinity

Supplied by CMU.

Dynamic Range (Positive Only)

	s	exp	frac	E	Value	
Denormalized numbers	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 * 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 * 1/64 = 2/512$	
	...					
	0	0000	110	-6	$6/8 * 1/64 = 6/512$	
	0	0000	111	-6	$7/8 * 1/64 = 7/512$	largest denorm
Normalized numbers	0	0001	000	-6	$8/8 * 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 * 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 * 1/2 = 14/16$	
	0	0110	111	-1	$15/8 * 1/2 = 15/16$	closest to 1 below
	0	0111	000	0	$8/8 * 1 = 1$	
	0	0111	001	0	$9/8 * 1 = 9/8$	closest to 1 above
	0	0111	010	0	$10/8 * 1 = 10/8$	
	...					
	0	1110	110	7	$14/8 * 128 = 224$	
	0	1110	111	7	$15/8 * 128 = 240$	largest norm
	0	1111	000	n/a	inf	

Supplied by CMU.


Distribution of Values

- 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- bias is $2^{3-1}-1 = 3$



- Notice how the distribution gets denser toward zero.




Supplied by CMU.

Distribution of Values (close-up view)

- 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3




Supplied by CMU.

Quiz 1

- 6-bit IEEE-like format

- e = 3 exponent bits
- f = 2 fraction bits
- bias is 3



What number is represented by 0 011 10?

- a) 12
- b) 1.5
- c) .5
- d) none of the above

Floating-Point Operations: Basic Idea

- $x \oplus_f y = \text{Round}(x \oplus y)$
- $x \otimes_f y = \text{Round}(x \otimes y)$
- **Basic idea**
 - first **compute exact result**
 - make it fit into desired precision
 - » possibly overflow if exponent too large
 - » possibly **round to fit into frac**

Supplied by CMU.

Rounding

- Rounding modes (illustrated with \$ rounding)


	\$1.40	\$1.60	\$1.50	\$2.50	-\$1.50
towards zero	\$1	\$1	\$1	\$2	-\$1
round down ($-\infty$)	\$1	\$1	\$1	\$2	-\$2
round up ($+\infty$)	\$2	\$2	\$2	\$3	-\$1
nearest integer	\$1	\$2	?	?	?
nearest even (default)	\$1	\$2	\$2	\$2	-\$2

Supplied by CMU.

Creating a Floating Point Number

- **Steps**

- normalize to have leading 1
- round to fit within fraction
- postnormalize to deal with effects of rounding



1 4-bits 3-bits

- **Case study**


- convert 8-bit unsigned numbers to tiny floating-point format

example numbers

128	10000000
13	00001101
33	00010001
35	00010011
138	10001010
63	00111111

Supplied by CMU.

Normalize



- **Requirement**

- set binary point so that numbers of form 1.xxxxx
- adjust all to have leading one
 - » decrement exponent as shift left

<i>Value</i>	<i>Binary</i>	<i>Fraction</i>	<i>Exponent</i>
128	10000000	1.0000000	7
13	00001101	1.1010000	3
17	00010001	1.0001000	4
19	00010011	1.0011000	4
138	10001010	1.0001010	7
63	00111111	1.1111100	5

Supplied by CMU.

Rounding

1.BBG**RXXX**

Guard bit: LSB of result

Round bit: 1st bit removed

Sticky bit: OR of remaining bits

- Round-up conditions

- round = 1, sticky = 1 $\Rightarrow > 0.5$

- guard = 1, round = 1, sticky = 0 \Rightarrow round to even

<i>Value</i>	<i>Fraction</i>	<i>GRS</i>	<i>Incr?</i>	<i>Rounded</i>
128	1.000 0000	000	N	1.000
13	1.101 0000	100	N	1.101
17	1.000 1000	010	N	1.000
19	1.001 1000	110	Y	1.010
138	1.000 1010	011	Y	1.001
63	1.111 1100	111	Y	10.000

Supplied by CMU.

Postnormalize

- **Issue**

- rounding may have caused overflow
- handle by shifting right once & incrementing exponent

<i>Value</i>	<i>Rounded</i>	<i>Exp</i>	<i>Adjusted</i>	<i>Result</i>
128	1.000	7		128
13	1.101	3		13
17	1.000	4		16
19	1.010	4		20
138	1.001	7		134
63	10.000	5	1.000×2^6	64

Supplied by CMU.

Floating-Point Multiplication

- $(-1)^{s_1} M_1 2^{E_1} \times (-1)^{s_2} M_2 2^{E_2}$
- **Exact result:** $(-1)^s M 2^E$
 - sign s : $s_1 \wedge s_2$
 - significand M : $M_1 \times M_2$
 - exponent E : $E_1 + E_2$
- **Fixing**
 - if $M \geq 2$, shift M right, increment E
 - if E out of range, overflow (or underflow)
 - round M to fit $frac$ precision
- **Implementation**
 - biggest chore is multiplying significands

Supplied by CMU.

Note that to compute E , one must first convert exp_1 and exp_2 to E_1 and E_2 , then add them together and check for underflow or overflow (corresponding to $-\infty$ and $+\infty$), and then convert to exp .

Floating-Point Addition

- $(-1)^{s_1} M_1 2^{E_1} + (-1)^{s_2} M_2 2^{E_2}$


–assume $E_1 > E_2$

- **Exact result:** $(-1)^s M 2^E$

–sign s , significand M :

» result of signed align & add

–exponent E : E_1



- **Fixing**

–if $M \geq 2$, shift M right, increment E

–if $M < 1$, shift M left k positions, decrement E by k

–overflow if E out of range

–round M to fit **frac** precision

Supplied by CMU.

Note that, by default, overflow results in either $+\infty$ or $-\infty$.

Floating Point in C

- **C guarantees two levels**

- float single precision
- double double precision

- **Conversions/casting**

- casting between int, float, and double changes bit representation
- double/float → int
 - » truncates fractional part
 - » like rounding toward zero
 - » not defined when out of range or NaN: generally sets to TMin
- int → double
 - » exact conversion, as long as int has ≤ 53 -bit word size
- int → float
 - » will round according to rounding mode

Supplied by CMU.

Quiz 2

Suppose f , declared to be a `float`, is assigned the largest possible floating-point positive value (other than $+\infty$). What is the value of $g = f + 1.0$?

- a) f
- b) $+\infty$
- c) NAN
- d) 0

Float is not Rational ...

- **Floating addition**

- commutative: $a +^f b = b +^f a$

- » yes!

- associative: $a +^f (b +^f c) = (a +^f b) +^f c$

- » no!

- $2 +^f (1e20 +^f -1e20) = 2$

- $(2 +^f 1e20) +^f -1e20 = 0$

Note that the floating-point numbers in this and the next two slides are expressed in base 10, not base 2.

Float is not Rational ...

- **Multiplication**

- commutative: $a *^f b = b *^f a$

- » yes!

- associative: $a *^f (b *^f c) = (a *^f b) *^f c$

- » no!

- $1e20 *^f (1e20 *^f 1e-20) = 1e20$

- $(1e20 *^f 1e20) *^f 1e-20 = +\infty$

Float is not Rational ...

- **More ...**

- **multiplication distributes over addition:**

- $a *^f (b +^f c) = (a *^f b) +^f (a *^f c)$

- » no!

- » $1e20 *^f (1e20 +^f -1e20) = 0$

- » $(1e20 *^f 1e20) +^f (1e20 *^f -1e20) = \text{NaN}$

- **loss of significance:**

- $x = y + 1$

- $z = 2 / (x - y)$

- $z == 2?$

- » not necessarily!

- consider $y = 1e20$