

CS 33

Data Representation (Part 4)

Floating-Point Operations: Basic Idea


- $x +_f y = \text{Round}(x + y)$
- $x \times_f y = \text{Round}(x \times y)$
- Basic idea
 - first **compute exact result**
 - make it fit into desired precision
 - » possibly overflow if exponent too large
 - » possibly **round to fit into** frac

Rounding

- Rounding modes (illustrated with \$ rounding)

| | \$1.40 | \$1.60 | \$1.50 | \$2.50 | −\$1.50 |
|--|---------------|---------------|---------------|---------------|----------------|
| towards zero | \$1 | \$1 | \$1 | \$2 | −\$1 |
| round down ($-\infty$) | \$1 | \$1 | \$1 | \$2 | −\$2 |
| round up ($+\infty$) | \$2 | \$2 | \$2 | \$3 | −\$1 |
| nearest integer | \$1 | \$2 | ? | ? | ? |
| nearest even (default) | \$1 | \$2 | \$2 | \$2 | −\$2 |

Creating a Floating Point Number



- **Steps**

- normalize to have leading 1
- round to fit within fraction
- postnormalize to deal with effects of rounding


- **Case study**

- convert 8-bit unsigned numbers to tiny floating-point format

example numbers

| | |
|-----|----------|
| 128 | 10000000 |
| 13 | 00001101 |
| 33 | 00010001 |
| 35 | 00010011 |
| 138 | 10001010 |
| 63 | 00111111 |

Normalize



- **Requirement**

- set binary point so that numbers of form 1.xxxxx
- adjust all to have leading one
 - » decrement exponent as shift left

| Value | Binary | Fraction | Exponent |
|-------|----------|-----------|----------|
| 128 | 10000000 | 1.0000000 | 7 |
| 13 | 00001101 | 1.1010000 | 3 |
| 17 | 00010001 | 1.0001000 | 4 |
| 19 | 00010011 | 1.0011000 | 4 |
| 138 | 10001010 | 1.0001010 | 7 |
| 63 | 00111111 | 1.1111100 | 5 |

Rounding

1.BBG**RXXX**

Guard bit: LSB of result

Round bit: 1st bit removed

Sticky bit: OR of remaining bits

- Round-up conditions

- round = 1, sticky = 1 $\Rightarrow > 0.5$
- guard = 1, round = 1, sticky = 0 \Rightarrow round up to even

| Value | Fraction | GRS | Incr? | Rounded |
|-------|-------------------|-----|-------|---------|
| 128 | 1.000 0000 | 000 | N | 1.000 |
| 13 | 1.101 0000 | 100 | N | 1.101 |
| 17 | 1.000 1000 | 010 | N | 1.000 |
| 19 | 1.001 1000 | 110 | Y | 1.010 |
| 138 | 1.000 1010 | 011 | Y | 1.001 |
| 63 | 1.111 1100 | 111 | Y | 10.000 |

Postnormalize

- **Issue**
 - rounding may have caused overflow
 - handle by shifting right once & incrementing exponent

| Value | Rounded | Exp | Adjusted | Result |
|-------|---------|-----|---------------|--------|
| 128 | 1.000 | 7 | | 128 |
| 13 | 1.101 | 3 | | 13 |
| 17 | 1.000 | 4 | | 16 |
| 19 | 1.010 | 4 | | 20 |
| 138 | 1.001 | 7 | | 134 |
| 63 | 10.000 | 5 | $1.000 * 2^6$ | 64 |

Floating-Point Multiplication

- $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$
- **Exact result:** $(-1)^s M 2^E$
 - sign s : $s1 \wedge s2$
 - significand M : $M1 \times M2$
 - exponent E : $E1 + E2$
- **Fixing**
 - if $M \geq 2$, shift M right, increment E
 - if E out of range, overflow (or underflow)
 - round M to fit `frac` precision
- **Implementation**
 - biggest chore is multiplying significands

Floating-Point Addition

- $(-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$


–assume $E1 > E2$

- **Exact result:** $(-1)^s M 2^E$

–sign s , significand M :

» result of signed align & add

–exponent E : $E1$



- **Fixing**

–if $M \geq 2$, shift M right, increment E

–if $M < 1$, shift M left k positions, decrement E by k

–overflow if E out of range

–round M to fit **frac** precision

Floating Point

- **Single precision (float)**



– range: $\pm 1.8 \times 10^{-38}$ – $\pm 3.4 \times 10^{38}$, ~7 decimal digits

- **Double Precision (double)**



– range: $\pm 2.23 \times 10^{-308}$ – $\pm 1.8 \times 10^{308}$, ~16 decimal digits

Floating Point in C

- **Conversions/casting**

- casting between `int`, `float`, and `double` changes bit representation
- `double/float` \rightarrow `int`
 - » truncates fractional part
 - » like rounding toward zero
 - » not defined when out of range or NaN: generally sets to TMin
- `int` \rightarrow `double`
 - » exact conversion, as long as `int` has ≤ 53 -bit word size
- `int` \rightarrow `float`
 - » will round according to rounding mode

Quiz 1

Suppose f , declared to be a `float`, is assigned the largest possible floating-point positive value (other than $+\infty$). What is the value of $g = f + 1.0$?

- a) f
- b) $+\infty$
- c) NaN
- d) 0

Float is not Rational ...

- **Floating addition**
 - commutative: $a +^f b = b +^f a$
 - » yes!
 - associative: $a +^f (b +^f c) = (a +^f b) +^f c$
 - » no!
 - $2 +^f (1e38 +^f -1e38) = 2$
 - $(2 +^f 1e38) +^f -1e38 = 0$

Float is not Rational ...

- **Multiplication**

- **commutative:** $a *^f b = b *^f a$

- » **yes!**

- **associative:** $a *^f (b *^f c) = (a *^f b) *^f c$

- » **no!**

- $1e37 *^f (1e37 *^f 1e-37) = 1e37$

- $(1e37 *^f 1e37) *^f 1e-37 = +\infty$


Float is not Rational ...

- **More ...**
 - **multiplication distributes over addition:**
$$a *^f (b +^f c) = (a *^f b) +^f (a *^f c)$$
 - » **no!**
 - » $1e38 *^f (1e38 +^f -1e38) = 0$
 - » $(1e38 *^f 1e38) +^f (1e38 *^f -1e38) = \text{NaN}$
 - **loss of significance:**
$$x = y + 1$$
$$z = 2 / (x - y)$$
$$z == 2?$$
 - » **not necessarily!**
 - **consider $y = 1e38$**

CS 33

Intro to Machine Programming

Machine Model



Memory



Instructions


Data

or



**Instructions
are Data**

Processor: Some Details



Processor: Basic Operation

```
while (forever) {  
  fetch instruction IP points at  
  decode instruction  
  fetch operands  
  execute  
  store results  
  update IP and condition code  
}
```

Instructions ...

| | | | |
|----------------|-----------------|-----------------|------------|
| Op code | Operand1 | Operand2 | ... |
|----------------|-----------------|-----------------|------------|

Operands

- **Form**
 - immediate vs. reference
 - » value vs. address
- **How many?**
 - 3
 - » add a,b,c
 - $c = a + b$
 - 2
 - » add a,b
 - $b += a$

Operands (continued)

- **Accumulator**
 - special memory in the processor
 - » known as a *register*
 - » fast access
 - allows single-operand instructions
 - » add a
 - `acc += a`
 - » add b
 - `acc += b`

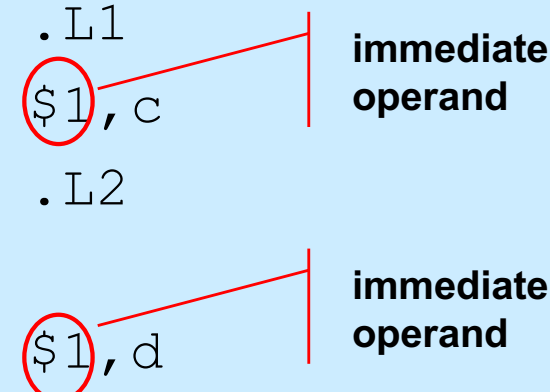
From C to Assembler ...

```
a = (b + c) * d;
```

```
mov    b,%acc
add    c,%acc
mul    d,%acc
mov    %acc,a
```

```
if (a<b)
    c = 1;
else
    d = 1;
```

```
cmp    a,b
jge     .L1
mov     $1,c
jmp     .L2
.L1
mov     $1,d
.L2
```



immediate operand

immediate operand

Condition Codes

- **Set of flags giving status of most recent operation:**
 - **zero flag**
 - » result was zero
 - **sign flag**
 - » for signed arithmetic interpretation: sign bit is set
 - **overflow flag**
 - » for signed arithmetic interpretation
 - **carry flag (generated by carry or borrow out of most-significant bit)**
 - » for unsigned arithmetic interpretation
- **Set implicitly by arithmetic instructions**
- **Set explicitly by compare instruction**
 - **cmp a,b**
 - » sets flags based on result of $b-a$

Examples (1)

- **Assume 32-bit arithmetic**
- **x is 0x80000000**
 - **TMIN if interpreted as two's-complement**
 - **2^{31} if interpreted as unsigned**
- **x-1 (0x7fffffff)**
 - **TMAX if interpreted as two's-complement**
 - **$2^{31}-1$ if interpreted as unsigned**
 - **zero flag is not set**
 - **sign flag is not set**
 - **overflow flag is set**
 - **carry flag is not set**

Examples (2)

- **x is 0xffffffff**
 - -1 if interpreted as two's-complement
 - UMAX ($2^{32}-1$) if interpreted as unsigned
- **x+1 (0x00000000)**
 - zero under either interpretation
 - zero flag is set
 - sign flag is not set
 - overflow flag is not set
 - carry flag is set

Examples (3)

- **x is 0xffffffff**
 - -1 if interpreted as two's-complement
 - UMAX ($2^{32}-1$) if interpreted as unsigned
- **x+2 (0x00000001)**
 - (+)1 under either interpretation
 - zero flag is not set
 - sign flag is not set
 - overflow flag is not set
 - carry flag is set

Quiz 2

- **Set of flags giving status of most recent operation:**
 - zero flag
 - » result was zero
 - sign flag
 - » for signed arithmetic interpretation: sign bit is set
 - overflow flag
 - » for signed arithmetic interpretation
 - carry flag (generated by carry or borrow out of most-significant bit)
 - » for unsigned arithmetic interpretation
- **Set explicitly by compare instruction**
 - `cmp a,b`
 - » sets flags based on result of `b-a`

Which flags are set to one by “`cmp 2,1`”?

- a) overflow flag only**
- b) carry flag only**
- c) sign and carry flags only**
- d) sign and overflow flags only**
- e) sign, overflow, and carry flags**

Jump Instructions

- **Unconditional jump**
 - just do it
- **Conditional jump**
 - to jump or not to jump determined by condition-code flags
 - field in the op code indicates how this is computed
 - in assembler language, simply say
 - » **je**
 - jump on equal
 - » **jne**
 - jump on not equal
 - » **jgt**
 - jump on greater than
 - » **etc.**

Addresses


```
int a, b, c, d;
```

```
int main() {  
    a = (b + c) * d;  
    ...  
}
```

```
mov    b, %acc  
add    c, %acc  
mul    d, %acc  
mov    %acc, a
```

```
mov    1004, %acc  
add    1008, %acc  
mul    1012, %acc  
mov    %acc, 1000
```

1012:
1008:
1004:
1000:



Memory

Addresses

```
int b;
```

```
int func(int c, int d) {  
    int a;  
    a = (b + c) * d;  
    ...  
}
```


```
mov    ?, %acc  
add    ?, %acc  
mul    ?, %acc  
mov    %acc, ?
```

- One copy of *b* for duration of program's execution
 - *b*'s address is the same for each call to *func*
- Different copies of *a*, *c*, and *d* for each call to *func*
 - addresses are different in each call

Relative Addresses


- **Absolute address**
 - actual location in memory
- **Relative address**
 - offset from some other location

- Blob's absolute address is 10000
- Datum's relative address (to Blob) is 100
 - its absolute address is 10100



Base Registers

```
mov $10000, %base  
mov $10, 100(%base)
```




Addresses

```
long b;
```

```
int func(long c, long d) {  
    long a;  
    a = (b + c) * d;  
    ...  
}
```

```
mov    1000,%acc  
add    -8(%base),%acc  
mul    -12(%base),%acc  
mov    %acc,-16(%base)
```




Quiz 3


Suppose the value in *base* is 10,000. What is the address of *c*?

- a) 9992
- b) 9996
- c) 10,004
- d) 10,008

```
mov    1000, %acc
add    -8(%base), %acc
mul    -12(%base), %acc
mov    %acc, -16(%base)
```



Registers



Registers vs. Memory

