

CS 33

Machine Programming (2)

Many of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

Jump Instructions

- **Unconditional jump**
 - just do it
- **Conditional jump**
 - to jump or not to jump determined by condition-code flags
 - field in the op code indicates how this is computed
 - in assembler language, simply say
 - » je
 - jump on equal
 - » jne
 - jump on not equal
 - » jgt
 - jump on greater than
 - » etc.

Jump instructions cause the processor to start executing instructions at some specified address. For conditional jump instructions, whether to jump or not is determined by the values of the condition codes. Fortunately, rather than having to specify explicitly those values, one may use mnemonics as shown in the slide.

We'll see examples of their use in the next lecture, when we start looking at x86 assembler instructions.


Addresses

```
int a, b, c, d;

int main() {
    a = (b + c) * d;
    ...
}
```

mov	b,%acc
add	c,%acc
mul	d,%acc
mov	%acc,a

mov	1004,%acc
add	1008,%acc
mul	1012,%acc
mov	%acc,1000



Memory

In the C code above, the assignment to *a* might be coded in assembler as shown in the box in the lower left. But this brings up the question, where are the values represented by *a*, *b*, *c*, and *d*? Variable names are part of the C language, not assembler. Let's assume that these global variables are located at addresses 1000, 1004, 1008, and 1012, as shown on the right. Thus correct assembler language would be as in the middle box, which deals with addresses, not variable names. Note that "mov 1004,%acc" means to copy the contents of location 1004 to the accumulator register; it does not mean to copy the integer 1004 into the register!

Beginning with this slide, whenever we draw pictures of memory, lower memory addresses are at the bottom, higher addresses are at the top. This is the opposite of how we've been drawing pictures of memory in previous slides.

Addresses

```
int b;  
  
int func(int c, int d) {  
    int a;  
    a = (b + c) * d;  
    ...  
}
```

```
mov    ?,%acc  
add    ?,%acc  
mul    ?,%acc  
mov    %acc,?
```


- One copy of *b* for duration of program's execution
 - *b*'s address is the same for each call to *func*
- Different copies of *a*, *c*, and *d* for each call to *func*
 - addresses are different in each call

Here we rearrange things a bit. *b* is a global variable, but *a* is a local variable within *func*, and *c* and *d* are arguments. The issue here is that the locations associated with *a*, *c*, and *d* will, in general, be different for each call to *func*. Thus we somehow must modify the assembler code to take this into account.

Relative Addresses

- **Absolute address**
 - actual location in memory
- **Relative address**
 - offset from some other location

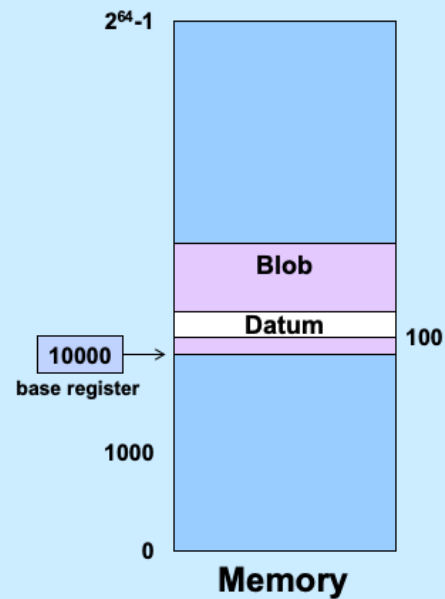
- Blob's absolute address is 10000
- Datum's relative address (to Blob) is 100
 - its absolute address is 10100



Note that both positive and negative offsets might be used.

Base Registers

```
mov $10000, %base  
mov $10, 100(%base)
```



Here we load the value 10,000 into the base register (recall that the “\$” means what follows is a literal value; a “%” sign means that what follows is the name of a register), then store the value 10 into the memory location 10100 (the contents of the base register plus 100): the notation $n(\%base)$ means the address obtained by adding n to the contents of the base register.

Addresses


```
long b;
```

```
int func(long c, long d) {  
    long a;  
    a = (b + c) * d;  
    ...  
}
```

```
mov    1000,%acc  
add    -8(%base),%acc  
mul    -12(%base),%acc  
mov    %acc,-16(%base)
```

base →

1000:



Memory


Here we return to our earlier example. We assume that, as part of the call to *func*, the base register is loaded with the address of the beginning of *func*'s current stack frame, and that the local variable *a* and the parameters *c* and *d* are located within the frame. Thus we refer to them by their offset from the beginning of the stack frame, which are assumed to be -16, -8, and -12. Since the stack grows from higher addresses to lower addresses, these offsets are negative. Note that the first assembler instruction copies the contents of location 1000 into %acc.

Quiz 1

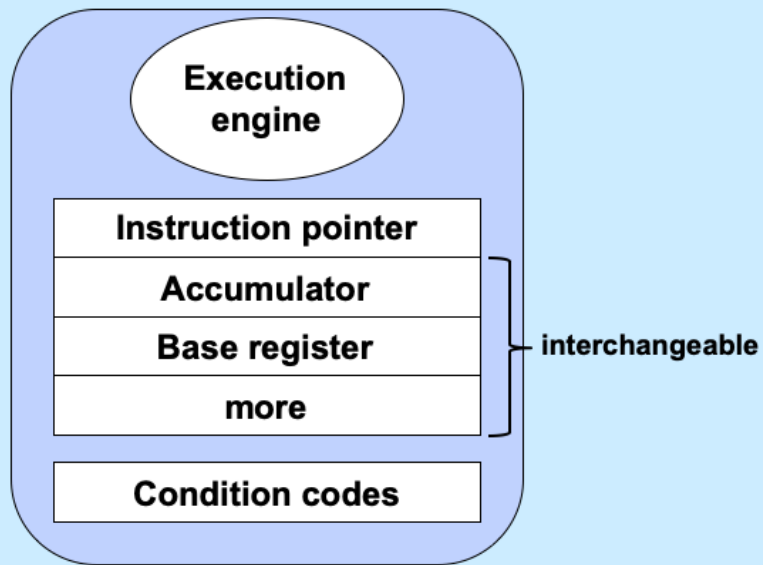
Suppose the value in *base* is 10,000. What is the address of *c*?

- a) 9992
- b) 9996
- c) 10,004
- d) 10,008

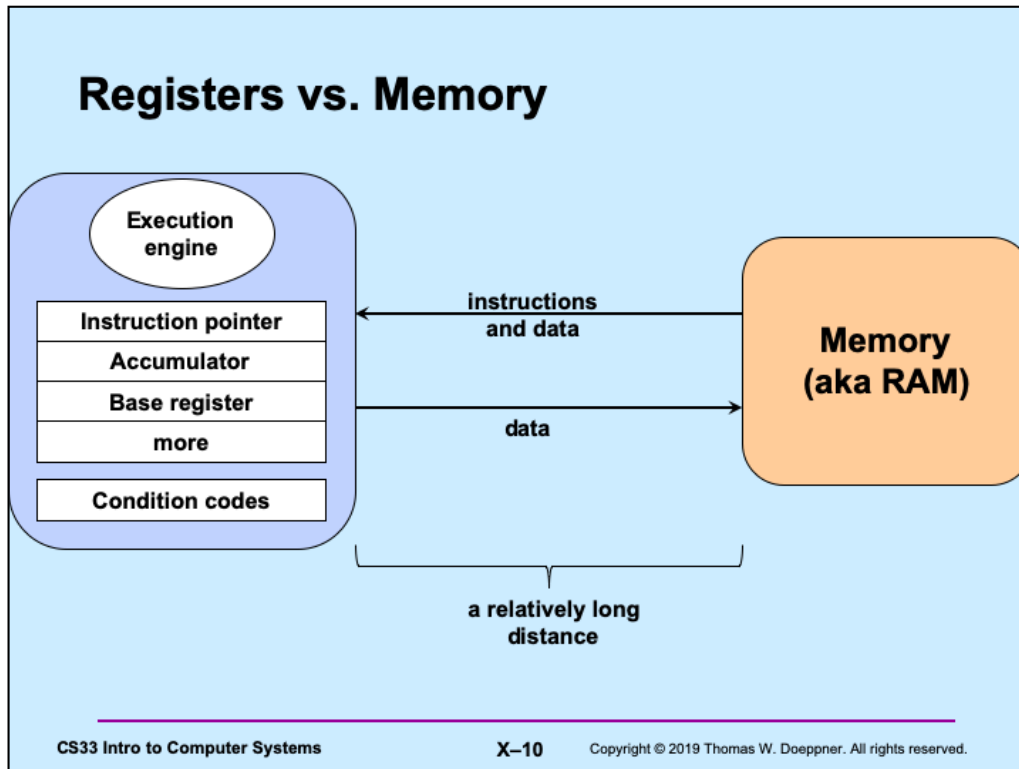
```
mov    1000,%acc
add    -8(%base),%acc
mul    -12(%base),%acc
mov    %acc,-16(%base)
```



Registers



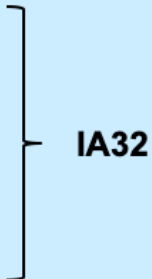
We've now seen four registers: the instruction pointer, the accumulator, the base register, and the condition codes. The accumulator is used to hold intermediate results for arithmetic; the base register is used to hold addresses for relative addressing. There's no particular reason why the accumulator can't be used as the base register and vice versa: thus they may be used interchangeably. Furthermore, it is useful to have more than two such dual-purpose registers. As we will see, the x86 architecture has eight such registers; the x86-64 architecture has 16.



Why do we make the distinction between registers and memory? Registers are in the processor itself and can be read from and written to very quickly. Memory is on separate hardware and takes much more time to access than registers do. Thus operations involving only registers can be executed very quickly, while significantly more time is required to access memory. Processors typically have relatively few registers (the IA-32 architecture has eight, the x86-64 architecture has 32; some other architectures have many more, perhaps as many as 256); memory is measured in gigabytes.

Note that memory access-time is mitigated by the use of on-processor caches, something that we will discuss in a few weeks.

Intel x86

- Intel created the 8008 (in 1972)
 - 8008 begat 8080
 - 8080 begat 8086
 - 8086 begat 8088
 - 8086 begat 286
 - 286 begat 386
 - 386 begat 486
 - 486 begat Pentium
 - Pentium begat Pentium Pro
 - Pentium Pro begat Pentium II
 - ad infinitum
- 
- IA32

The early computers of the x86 family had 16-bit words, starting with the 386, they supported 32-bit words.

2^{64}

- **2^{32} used to be considered a large number**
 - one couldn't afford 2^{32} bytes of memory, so no problem with that as an upper bound
- **Intel (and others) saw need for machines with 64-bit addresses**
 - devised IA64 architecture with HP
 - » became known as Itanium
 - » very different from x86
- **AMD also saw such a need**
 - developed 64-bit extension to x86, called x86-64
- **Itanium flopped**
- **x86-64 dominated**
- **Intel, reluctantly, adopted x86-64**

2^{32} = 4 gigabytes.

2^{64} = 16 exbibytes


All SunLab computers are x86-64.

Data Types on IA32 and x86-64

- **“Integer” data of 1, 2, or 4 bytes (plus 8 bytes on x86-64)**
 - **data values**
 - » whether signed or unsigned depends on interpretation
 - **addresses (untyped pointers)**
- **Floating-point data of 4, 8, or 10 bytes**
- **No aggregate types such as arrays or structures**
 - **just contiguously allocated bytes in memory**

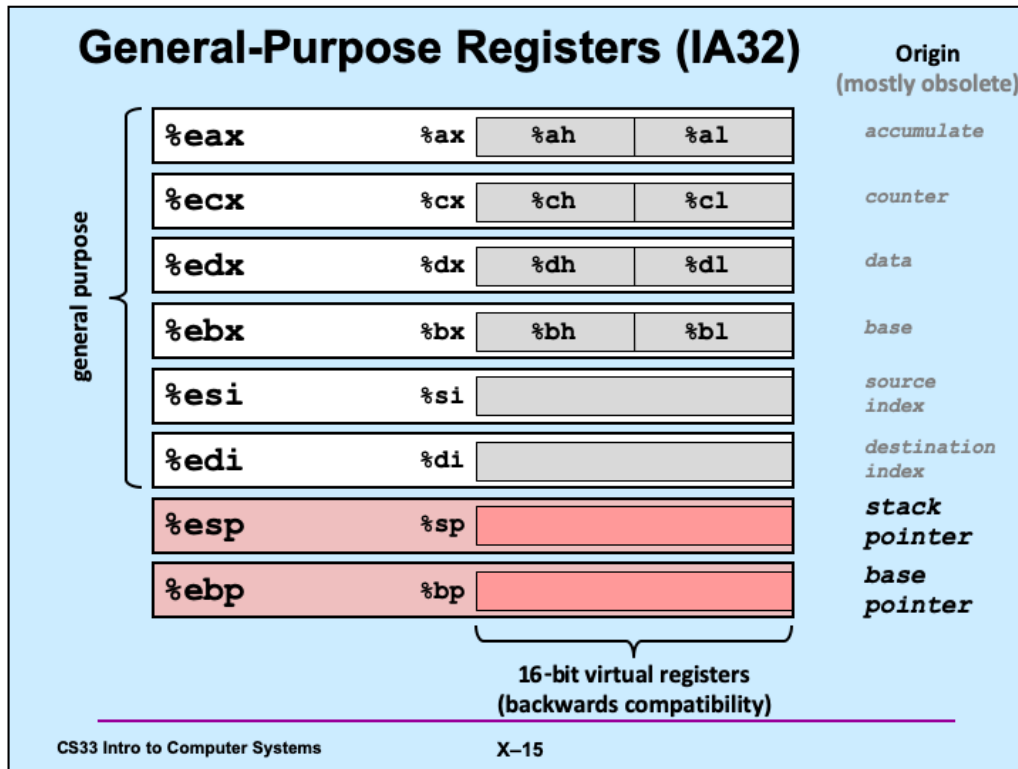
Supplied by CMU.

Operand Size



- **Rather than `mov ...`**
 - `movb`
 - `movs`
 - `movl`
 - `movq` (x86-64 only)

Most instructions come in three (on IA32) or four (on x86-64) forms, one for each possible operand size.



Supplied by CMU.

Moving Data: IA32

- Moving data

`movl source, dest`

- Operand types

- **Immediate:** constant integer data

- » example: `$0x400`, `$-533`
 - » like C constant, but prefixed with ``$'`
 - » encoded with 1, 2, or 4 bytes

- **Register:** one of 8 integer registers

- » example: `%eax`, `%edx`
 - » but `%esp` and `%ebp` reserved for special use
 - » others have special uses for particular instructions

- **Memory:** 4 consecutive bytes of memory at address given by register(s)

- » simplest example: `(%eax)`
 - » various other “address modes”

<code>%eax</code>
<code>%ecx</code>
<code>%edx</code>
<code>%ebx</code>
<code>%esi</code>
<code>%edi</code>
<code>%esp</code>
<code>%ebp</code>

Supplied by CMU.

Note that though `esp` and `ebp` have special uses, they may also be used in both source and destination operands.

Some assemblers (in particular, those of Intel and Microsoft) place the operands in the opposite order. Thus the example of the slide would be `addl %eax,8(%ebp)`. The order we use is that used by gcc, known as the “AT&T syntax” because it was used in the original Unix assemblers, written at Bell Labs, then part of AT&T.

movl Operand Combinations

	Source	Dest	Src, Dest	C Analog
movl	Imm	Reg	movl \$0x4,%eax	temp = 0x4;
		Mem	movl \$-147,(%eax)	*p = -147;
	Reg	Reg	movl %eax,%edx	temp2 = temp1;
		Mem	movl %eax,(%edx)	*p = temp;
	Mem	Reg	movl (%eax),%edx	temp = *p;

Cannot (normally) do memory-memory transfer with a single instruction

Supplied by CMU.

Simple Memory Addressing Modes

- **Normal (R) Mem[Reg[R]]**

– register R specifies memory address

```
movl (%ecx), %eax
```

- **Displacement D(R) Mem[Reg[R]+D]**

– register R specifies start of memory region

– constant displacement D specifies offset

```
movl 8(%ebp), %edx
```

Supplied by CMU.

If one thinks of there being an array of registers, then “Reg[R]” selects register “R” from this array.

Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp
movl  %esp,%ebp
pushl %ebx
```

} Set Up

```
movl  8(%ebp), %edx
movl 12(%ebp), %ecx
movl (%edx), %ebx
movl (%ecx), %eax
movl %eax, (%edx)
movl %ebx, (%ecx)
```

} Body

```
popl  %ebx
popl  %ebp
ret
```

} Finish

Supplied by CMU.


We discuss the “set up” and “finish” in a subsequent lecture. They have to do with facilitating the calling of functions.

Understanding Swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

Register	Value
%edx	xp
%ecx	yp
%ebx	t0
%eax	t1

```
movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0
```



Understanding Swap

%eax	
%edx	
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Offset	Address
		123	0x124
		456	0x120
			0x11c
			0x118
			0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%ebp	→ 0		0x104
	-4		0x100

```

movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0

```

Supplied by CMU.

Understanding Swap

%eax	
%edx	0x124
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address	
		123	0x124
		456	0x120
			0x11c
			0x118
	Offset		0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%ebp	→ 0		0x104
	-4		0x100

```
movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0
```

Supplied by CMU.

%eax	
%edx	0x124
%ecx	0x120
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

			Address
		123	0x124
		456	0x120
			0x11c
			0x118
	Offset		0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%ebp	→ 0		0x104
	-4		0x100

```
movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0
```

Supplied by CMU.

Understanding Swap

%eax	
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Offset	Address
		123	0x124
		456	0x120
			0x11c
			0x118
			0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%ebp	→ 0		0x104
	-4		0x100

```
movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0
```

Supplied by CMU.

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		123
		0x124
		456
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
		0x110
xp	8	0x124
		0x10c
		4
		Rtn adr
		0x108
%ebp	→ 0	
		0x104
		-4
		0x100

```

movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx  # ebx = *xp (t0)
movl (%ecx), %eax  # eax = *yp (t1)
movl %eax, (%edx)  # *xp = t1
movl %ebx, (%ecx)  # *yp = t0
    
```

Supplied by CMU.

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
	4	Rtn adr
%ebp	→ 0	0x108
	-4	0x104
		0x100

```

movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx   # ebx = *xp (t0)
movl (%ecx), %eax   # eax = *yp (t1)
movl %eax, (%edx)  # *xp = t1
movl %ebx, (%ecx)   # *yp = t0
    
```

Supplied by CMU.

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104


		Offset	Address
		456	0x124
		123	0x120
			0x11c
			0x118
			0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%ebp	→ 0		0x104
	-4		0x100

```
movl 8(%ebp), %edx # edx = xp
movl 12(%ebp), %ecx # ecx = yp
movl (%edx), %ebx # ebx = *xp (t0)
movl (%ecx), %eax # eax = *yp (t1)
movl %eax, (%edx) # *xp = t1
movl %ebx, (%ecx) # *yp = t0
```

Supplied by CMU.

Quiz 2

```
movl -4(%ebp), %eax
movl (%eax), %eax
movl (%eax), %eax
movl %eax, -8(%ebp)
```



Which C statements best describe the assembler code?

```
// a
int x;
int y;
y = x;
```

```
// b
int *x;
int y;
y = *x;
```

```
// c
int **x;
int y;
y = **x;
```

```
// d
int ***x;
int y;
y = ***x;
```

Complete Memory-Addressing Modes

- Most general form

$D(Rb, Ri, S)$ $Mem[Reg[Rb] + S * Reg[Ri] + D]$

- D: constant “displacement”
- Rb: base register: any of 8 integer registers
- Ri: index register: any, except for %esp
 - » unlikely you’d use %ebp either
- S: scale: 1, 2, 4, or 8

- Special cases

(Rb, Ri)	$Mem[Reg[Rb] + Reg[Ri]]$
$D(Rb, Ri)$	$Mem[Reg[Rb] + Reg[Ri] + D]$
(Rb, Ri, S)	$Mem[Reg[Rb] + S * Reg[Ri]]$
D	$Mem[D]$

Address-Computation Examples

<code>%edx</code>	<code>0xf000</code>
<code>%ecx</code>	<code>0x0100</code>

Expression	Address Computation	Address
<code>0x8(%edx)</code>	$0xf000 + 0x8$	<code>0xf008</code>
<code>(%edx,%ecx)</code>	$0xf000 + 0x0100$	<code>0xf100</code>
<code>(%edx,%ecx,4)</code>	$0xf000 + 4*0x0100$	<code>0xf400</code>
<code>0x80(,%edx,2)</code>	$2*0xf000 + 0x80$	<code>0x1e080</code>

Supplied by CMU.

Address-Computation Instruction

- **leal src, dest**
 - src is address mode expression
 - set *dest* to address denoted by expression
- **Uses**
 - computing addresses without a memory reference
 - » e.g., translation of `p = &x[i];`
 - computing arithmetic expressions of the form $x + k*y$
 - » $k = 1, 2, 4, \text{ or } 8$
- **Example**

```
int mul12(int x)
{
    return x*12;
}
```

Converted to ASM by compiler:

```
movl 8(%ebp), %eax    # get arg
leal (%eax,%eax,2), %eax # t <- x+x*2
sall $2, %eax          # return t<<2
```

Supplied by CMU.

Note that a function returns a value by putting it in %eax.

Quiz 3

What value ends up in %ecx?

```
movl $1000,%eax
movl $1,%ebx
movl 2(%eax,%ebx,4),%ecx
```

- a) 0x02030405
- b) 0x05040302
- c) 0x06070809
- d) 0x09080706

1009:	0x09
1008:	0x08
1007:	0x07
1006:	0x06
1005:	0x05
1004:	0x04
1003:	0x03
1002:	0x02
1001:	0x01
1000:	0x00

%eax →

Hint:



x86-64 General-Purpose Registers						
	%rax	%eax		%r8	%r8d	a5
	%rbx	%ebx		%r9	%r9d	a6
a4	%rcx	%ecx		%r10	%r10d	
a3	%rdx	%edx		%r11	%r11d	
a2	%rsi	%esi		%r12	%r12d	
a1	%rdi	%edi		%r13	%r13d	
	%rsp	%esp		%r14	%r14d	
	%rbp	%ebp		%r15	%r15d	

- Extend existing registers to 64 bits. Add 8 new ones.
- No special purpose for %ebp/%rbp

CS33 Intro to Computer Systems

X-33

Supplied by CMU.

Note that %ebp/%rbp may be used as a base register as on IA32, but they don't have to be used that way. This will become clearer when we explore how the runtime stack is accessed. The convention on Linux is for the first 6 arguments of a function to be in registers %rdi, %rsi, %rdx, %rcx, %r8, and %r9. The return value of a function is put in %rax.

Note also that each register, in addition to having a 32-bit version, also has an 8-bit (one-byte) version. For the numbered registers, it's, for example, %r10b. For the other registers it's the same as for IA32.

32-bit Instructions on x86-64

- **addl 4(%rdx), %eax**
 - memory address must be 64 bits
 - operands (in this case) are 32-bit
 - » result goes into %eax
 - lower half of %rax
 - upper half is filled with zeroes

On x86-64, for instructions with 32-bit (long) operands that produce 32-bit results going into a register, the register must be a 32-bit register; the higher-order 32 bits are filled with zeroes.

Bytes

- **Each register has a byte version**
 - e.g., %r10: %r10b
- **Needed for byte instructions**
 - `movb (%rax, %rsi), %r10b`
 - sets *only* the low byte in %r10
 - » other seven bytes are unchanged
- **Alternatives**
 - `movzbq (%rax, %rsi), %r10`
 - » copies byte to low byte of %r10
 - » zeroes go to higher bytes
 - `movsbq (%rax, %rsi), %r10`
 - » copies byte to low byte of %r10
 - » sign is extended to all higher bits

Note that using single-byte versions of registers has a different behavior from using 4-byte versions of registers. Putting data into the latter using `mov` causes the upper bytes to be zeroed. But with the byte versions, putting data into them does not affect the upper bytes.

32-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp
movl  %esp,%ebp
pushl %ebx
```

} Set
Up

```
movl  8(%ebp), %edx
movl  12(%ebp), %ecx
movl  (%edx), %ebx
movl  (%ecx), %eax
movl  %eax, (%edx)
movl  %ebx, (%ecx)
```

} Body

```
popl  %ebx
popl  %ebp
ret
```

} Finish

Supplied by CMU.

Note that for the IA32 architecture, arguments are passed on the stack.

64-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
movl    (%rdi), %edx
movl    (%rsi), %eax
movl    %eax, (%rdi)
movl    %edx, (%rsi)
```

ret

} Set
Up

} Body

} Finish

- **Arguments passed in registers**
 - first (xp) in %rdi, second (yp) in %rsi
 - 64-bit pointers
- **No stack operations required**
- **32-bit data**
 - data held in registers %eax and %edx
 - movl operation

Supplied by CMU.

No more than six arguments can be passed in registers. If there are more than six arguments (which is unusual), then remaining arguments are passed on the stack, and referenced via %rsp.

64-bit code for long int swap

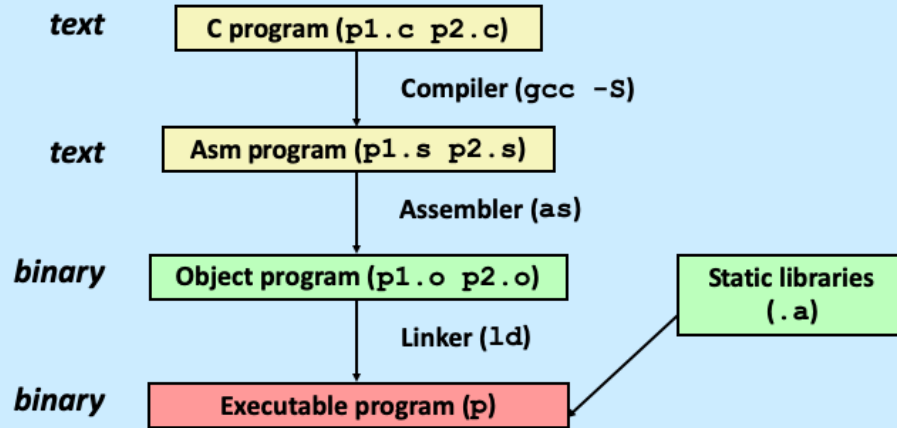
<pre>void swap(long *xp, long *yp) { long t0 = *xp; long t1 = *yp; *xp = t1; *yp = t0; }</pre>	<pre>swap_1: movq (%rdi), %rdx movq (%rsi), %rax movq %rax, (%rdi) movq %rdx, (%rsi) ret</pre>	<p>} Set Up</p> <p>} Body</p> <p>} Finish</p>
----------------------------------------------------------------------------------------------------------------	--------------------------------------------------------------------------------------------------------------------------------	-----------------------------------------------

- **64-bit data**

- data held in registers `%rax` and `%rdx`
- `movq` operation
 - » “q” stands for quad-word

Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -O1 p1.c p2.c -o p`
 - » use basic optimizations (`-O1`)
 - » put resulting binary in file `p`



Supplied by CMU.

Note that normally one does not ask `gcc` to produce assembler code, but instead it compiles C code directly into machine code (producing an object file). Note also that the `gcc` command actually invokes a script; the compiler (also known as `gcc`) compiles code into either assembler code or machine code; if necessary, the assembler (`as`) assembles assembler code into object code. The linker (`ld`) links together multiple object files (containing object code) into an executable program.

Example

```
int sum(int a, int b) {  
    return(a+b);  
}
```


Object Code

Code for `sum`

`0x401040 <sum>:`

`0x55`

`0x89`

`0xe5`

`0x8b`

`0x45`

`0x0c`

`0x03`

`0x45`

`0x08`

`0x5d`

`0xc3`

- **Total of 11 bytes**
- **Each instruction:
1, 2, or 3 bytes**
- **Starts at address
0x401040**

• Assembler


- translates `.s` into `.o`
- binary encoding of each instruction
- nearly-complete image of executable code
- missing linkages between code in different files

• Linker

- resolves references between files
- combines with static run-time libraries
 - » e.g., code for `printf`
- some libraries are *dynamically linked*
 - » linking occurs when program begins execution

Supplied by CMU.

Instruction Format



This is taken from Intel 64 and IA-32 Architecture Software Developer's Manual, Volume 2: Instruction Set Reference; Order Number 325462-043US, Intel Corporation, May 2012 (<https://software.intel.com/en-us/download/intel-64-and-ia-32-architectures-sdm-combined-volumes-1-2a-2b-2c-2d-3a-3b-3c-3d-and-4>)

Disassembling Object Code

Disassembled

```
080483c4 <sum>:  
80483c4: 55          push    %ebp  
80483c5: 89 e5       mov     %esp,%ebp  
80483c7: 8b 45 0c    mov     0xc(%ebp),%eax  
80483ca: 03 45 08    add     0x8(%ebp),%eax  
80483cd: 5d          pop     %ebp  
80483ce: c3          ret
```

- **Disassembler**

`objdump -d <file>`

- useful tool for examining object code
- analyzes bit pattern of series of instructions
- produces approximate rendition of assembly code
- can be run on either executable or object (.o) file

Alternate Disassembly

Object

```
0x401040:  
0x55  
0x89  
0xe5  
0x8b  
0x45  
0x0c  
0x03  
0x45  
0x08  
0x5d  
0xc3
```

Disassembled

Dump of assembler code for function sum:

```
0x080483c4 <sum+0>:    push    %ebp  
0x080483c5 <sum+1>:    mov     %esp,%ebp  
0x080483c7 <sum+3>:    mov     0xc(%ebp),%eax  
0x080483ca <sum+6>:    add     0x8(%ebp),%eax  
0x080483cd <sum+9>:    pop     %ebp  
0x080483ce <sum+10>:   ret
```

- **Within gdb debugger**

`gdb <file>`

`disassemble sum`

– `disassemble procedure`

`x/11xb sum`

– `examine the 11 bytes starting at sum`

Supplied by CMU.

How Many Instructions are There?

- We cover ~30
- Implemented by Intel:
 - 80 in original 8086 architecture
 - 7 added with 80186
 - 17 added with 80286
 - 33 added with 386
 - 6 added with 486
 - 6 added with Pentium
 - 1 added with Pentium MMX
 - 4 added with Pentium Pro
 - 8 added with SSE
 - 8 added with SSE2
 - 2 added with SSE3
 - 14 added with x86-64
 - 10 added with VT-x
 - 2 added with SSE4a
- Total: 198
- Doesn't count:
 - floating-point instructions
 - » ~100
 - SIMD instructions
 - » lots
 - AMD-added instructions
 - undocumented instructions

The source for this is http://en.wikipedia.org/wiki/X86_instruction_listings, viewed on 6/20/2017, which comes with the caveat that it may be out of date.

Some Arithmetic Operations

- **Two-operand instructions:**

Format	Computation	
<code>addl</code>	<code>Src, Dest</code>	<code>Dest = Dest + Src</code>
<code>subl</code>	<code>Src, Dest</code>	<code>Dest = Dest - Src</code>
<code>imull</code>	<code>Src, Dest</code>	<code>Dest = Dest * Src</code>
<code>sall</code>	<code>Src, Dest</code>	<code>Dest = Dest << Src</code>
<code>sarl</code>	<code>Src, Dest</code>	<code>Dest = Dest >> Src</code>
<code>shrl</code>	<code>Src, Dest</code>	<code>Dest = Dest >> Src</code>
<code>xorl</code>	<code>Src, Dest</code>	<code>Dest = Dest ^ Src</code>
<code>andl</code>	<code>Src, Dest</code>	<code>Dest = Dest & Src</code>
<code>orl</code>	<code>Src, Dest</code>	<code>Dest = Dest Src</code>

Also called `shll`

Arithmetic

Logical

- watch out for argument order!
- no distinction between signed and unsigned int (why?)

Supplied by CMU.

Note that for shift instructions, the `Src` operand (which is the size of the shift) must either be a immediate operand or be a designator for a one-byte register (e.g., `%cl` – see the slide on general-purpose registers for IA32).

Some Arithmetic Operations

- **One-operand Instructions**

<code>incl</code>	<code>Dest</code>	$= \text{Dest} + 1$
<code>decl</code>	<code>Dest</code>	$= \text{Dest} - 1$
<code>negl</code>	<code>Dest</code>	$= -\text{Dest}$
<code>notl</code>	<code>Dest</code>	$= \sim\text{Dest}$

- **See book for more instructions**

Supplied by CMU.

Arithmetic Expression Example

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

```
arith:
    leal    (%rdi,%rsi), %eax
    addl    %edx, %eax
    leal    (%rsi,%rsi,2), %edx
    sall    $4, %edx
    leal    4(%rdi,%rdx), %ecx
    imull   %ecx, %eax
    ret
```

Supplied by CMU, but converted to x86-64.

Understanding arith

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

```
leal  (%rdi,%rsi), %eax
addl  %edx, %eax
leal  (%rsi,%rsi,2), %edx
sall  $4, %edx
leal  4(%rdi,%rdx), %ecx
imull %ecx, %eax
ret
```

%rdx	z
%rsi	y
%rdi	x

Supplied by CMU, but converted to x86-64.

Understanding arith

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

%rdx	z
%rsi	y
%rdi	x

```
leal  (%rdi,%rsi), %eax    # eax = x+y    (t1)
addl  %edx, %eax          # eax = t1+z    (t2)
leal  (%rsi,%rsi,2), %edx  # edx = 3*y    (t4)
sall  $4, %edx            # edx = t4*16    (t4)
leal  4(%rdi,%rdx), %ecx   # ecx = x+4+t4  (t5)
imull %ecx, %eax          # eax *= t5     (rval)
ret
```

Supplied by CMU, but converted to x86-64.

By convention, the first three arguments to a procedure are placed in registers rdi, rsi, and rdx, respectively. Note that, also by convention, procedures put their return values in register eax/rax.

Observations about arith

```
int arith(int x, int y, int z)
{
    int t1 = x+y;
    int t2 = z+t1;
    int t3 = x+4;
    int t4 = y * 48;
    int t5 = t3 + t4;
    int rval = t2 * t5;
    return rval;
}
```

- Instructions in different order from C code
- Some expressions might require multiple instructions
- Some instructions might cover multiple expressions

```
leal    (%rdi,%rsi), %eax    # eax = x+y      (t1)
addl    %edx, %eax          # eax = t1+z      (t2)
leal    (%rsi,%rsi,2), %edx  # edx = 3*y      (t4)
sall    $4, %edx            # edx = t4*16     (t4)
leal    4(%rdi,%rdx), %ecx   # ecx = x+4+t4   (t5)
imull   %ecx, %eax          # eax *= t5      (rval)
ret
```

Supplied by CMU, but converted to x86-64.

Another Example

```
int logical(int x, int y)
{
    int t1 = x^y;
    int t2 = t1 >> 17;
    int mask = (1<<13) - 7;
    int rval = t2 & mask;
    return rval;
}
```

$2^{13} = 8192, 2^{13} - 7 = 8185$

<code>xorl %esi, %edi</code>	<code># edi = x^y</code>	<code>(t1)</code>
<code>sarl \$17, %edi</code>	<code># edi = t1>>17</code>	<code>(t2)</code>
<code>movl %edi, %eax</code>	<code># eax = edi</code>	
<code>andl \$8185, %eax</code>	<code># eax = t2 & mask</code>	<code>(rval)</code>

Supplied by CMU, but converted to x86-64.

Quiz 4

- What is the final value in `%ecx`?

```
xorl %ecx, %ecx  
incl %ecx  
sall %cl, %ecx # %cl is the low byte of %ecx  
addl %ecx, %ecx
```

- a) 2
- b) 4
- c) 8
- d) indeterminate