

CS 33

Memory Hierarchy II

Most of the slides in this lecture are either from or adapted from slides provided by the authors of the textbook “Computer Systems: A Programmer’s Perspective,” 2nd Edition and are provided from the website of Carnegie-Mellon University, course 15-213, taught by Randy Bryant and David O’Hallaron in Fall 2010. These slides are indicated “Supplied by CMU” in the notes section of the slides.

What's Inside A Disk Drive?





Image courtesy of Seagate Technology

Disk Geometry


- Disks consist of **platters**, each with two **surfaces**
- Each surface consists of concentric rings called **tracks**
- Each track consists of **sectors** separated by **gaps**



Supplied by CMU.

Disk Geometry (Multiple-Platter View)

- Aligned tracks form a cylinder



Supplied by CMU.

Disk Capacity

- **Capacity**: maximum number of bits that can be stored
 - capacity expressed in units of gigabytes (GB), where $1 \text{ GB} = 2^{30} \text{ Bytes} \approx 10^9 \text{ Bytes}$
- Capacity is determined by these technology factors:
 - **recording density** (bits/in): number of bits that can be squeezed into a 1 inch segment of a track
 - **track density** (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment
 - **areal density** (bits/in²): product of recording and track density
- Modern disks partition tracks into disjoint subsets called **recording zones**
 - each track in a zone has the same number of sectors, determined by the circumference of innermost track
 - each zone has a different number of sectors/track

Supplied by CMU.

Computing Disk Capacity

$$\text{Capacity} = (\# \text{ bytes/sector}) \times (\text{avg. } \# \text{ sectors/track}) \times$$
$$(\# \text{ tracks/surface}) \times (\# \text{ surfaces/platter}) \times$$
$$(\# \text{ platters/disk})$$

Example:


- 512 bytes/sector
- 600 sectors/track (on average)
- 40,000 tracks/surface
- 2 surfaces/platter
- 5 platters/disk

$$\begin{aligned}\text{Capacity} &= 512 \times 600 \times 40000 \times 2 \times 5 \\ &= 122,880,000,000 \\ &= 113.88 \text{ GB}\end{aligned}$$

Supplied by CMU.


Note that $1\text{GB} = 2^{30}$ bytes.

Disk Operation (Single-Platter View)




Supplied by CMU.

Disk Operation (Multi-Platter View)



Supplied by CMU.

Disk Structure: Top View of Single Platter




Surface organized into tracks

Tracks divided into sectors

Supplied by CMU.


Disk Access



Head in position above a track

Supplied by CMU.


Disk Access



Rotation is counter-clockwise


Supplied by CMU.

Disk Access – Read



About to read blue sector


Disk Access – Read



After **BLUE**
read

After reading blue sector


Disk Access – Read



After **BLUE**
read


Red request scheduled next

Disk Access – Seek




Seek to red's track

Disk Access – Rotational Latency



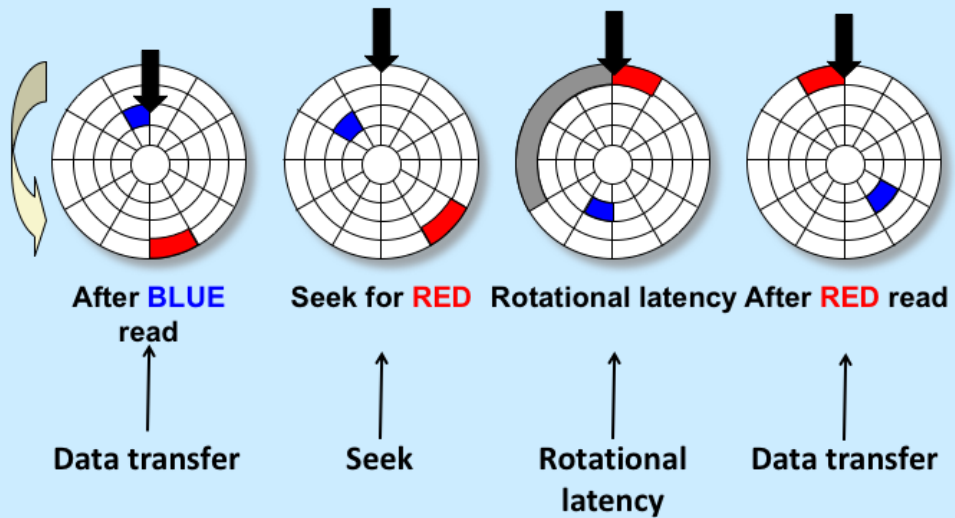
Wait for red sector to rotate around

Disk Access – Read



Complete read of red

Disk Access – Service Time Components



Supplied by CMU.

Disk Access Time

- Average time to access some target sector approximated by :
 - $T_{\text{access}} = T_{\text{avg seek}} + T_{\text{avg rotation}} + T_{\text{avg transfer}}$
- **Seek time** ($T_{\text{avg seek}}$)
 - time to position heads over cylinder containing target sector
 - typical $T_{\text{avg seek}}$ is 3–9 ms
- **Rotational latency** ($T_{\text{avg rotation}}$)
 - time waiting for first bit of target sector to pass under r/w head
 - typical rotation speed $R = 7200$ RPM
 - $T_{\text{avg rotation}} = \frac{1}{2} \times \frac{1}{R} \times 60 \text{ sec/1 min}$
- **Transfer time** ($T_{\text{avg transfer}}$)
 - time to read the bits in the target sector
 - $T_{\text{avg transfer}} = \frac{1}{R} \times \frac{1}{(\text{avg \# sectors/track})} \times 60 \text{ secs/1 min}$

Supplied by CMU.

Disk Access Time Example

- **Given:**
 - rotational rate = 7,200 RPM
 - average seek time = 9 ms
 - avg # sectors/track = 600
- **Derived:**
 - Tavg rotation = $1/2 \times (60 \text{ secs}/7200 \text{ RPM}) \times 1000 \text{ ms/sec} = 4 \text{ ms}$
 - Tavg transfer = $60/7200 \text{ RPM} \times 1/600 \text{ sects/track} \times 1000 \text{ ms/sec} = 0.014 \text{ ms}$
 - Taccess = 9 ms + 4 ms + 0.014 ms
- **Important points:**
 - access time dominated by seek time and rotational latency
 - first bit in a sector is the most expensive, the rest are free
 - SRAM access time is about 4 ns/doubleword, DRAM about 60 ns
 - » disk is about 40,000 times slower than SRAM
 - » 2,500 times slower than DRAM

Supplied by CMU.

Quiz 1

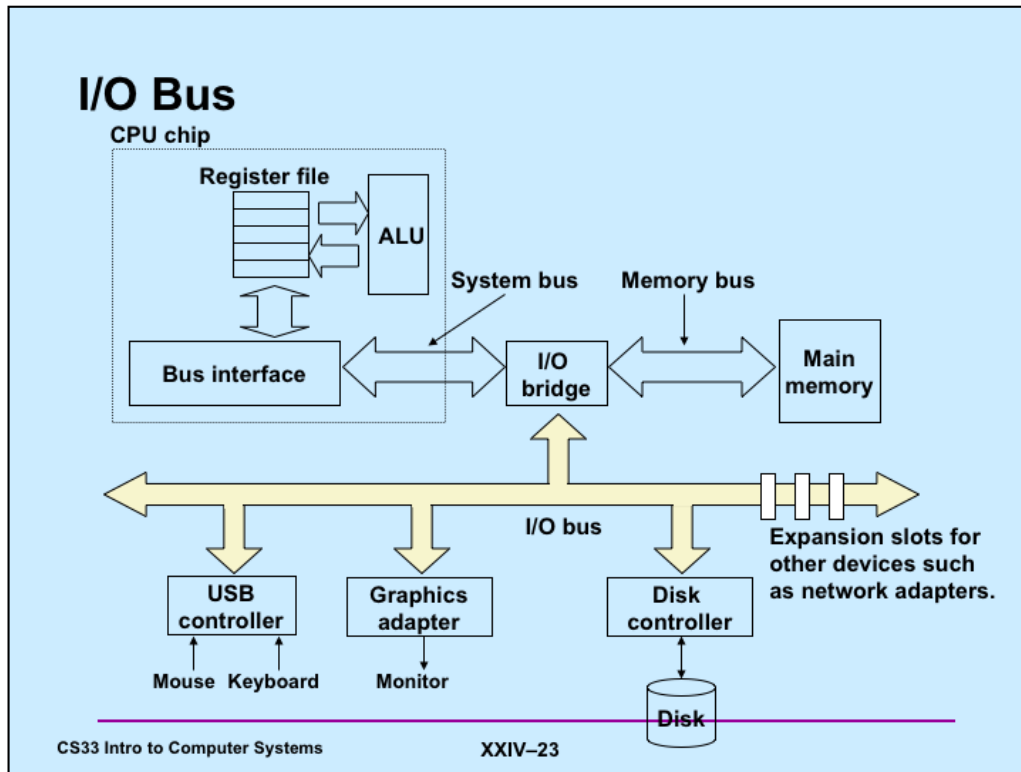
Assuming a 5-inch diameter disk spinning at 10,000 RPM, what is the approximate speed at which the outermost track is moving?

- a) faster than a speeding bullet (i.e., supersonic)
- b) roughly the speed of a pretty fast car (250 kph/155 mph)
- c) roughly the speed of a pretty slow car (50 mph)
- d) roughly the speed of a world-class marathoner (13.1 mph)

Logical Disk Blocks

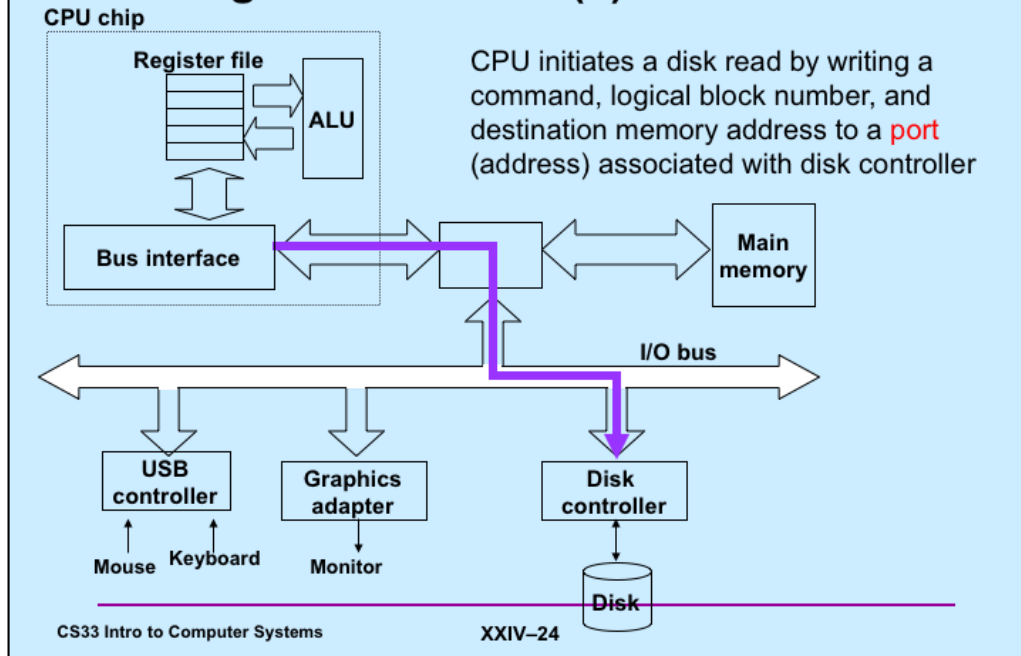
- **Modern disks present a simpler abstract view of the complex sector geometry:**
 - the set of available sectors is modeled as a sequence of b-sized **logical blocks** (0, 1, 2, ...)
- **Mapping between logical blocks and actual (physical) sectors**
 - maintained by hardware/firmware device called disk controller
 - converts requests for logical blocks into (surface, track, sector) triples
- **Allows controller to set aside spare cylinders for each zone**
 - accounts for the difference in “formatted capacity” and “maximum capacity”

Supplied by CMU.



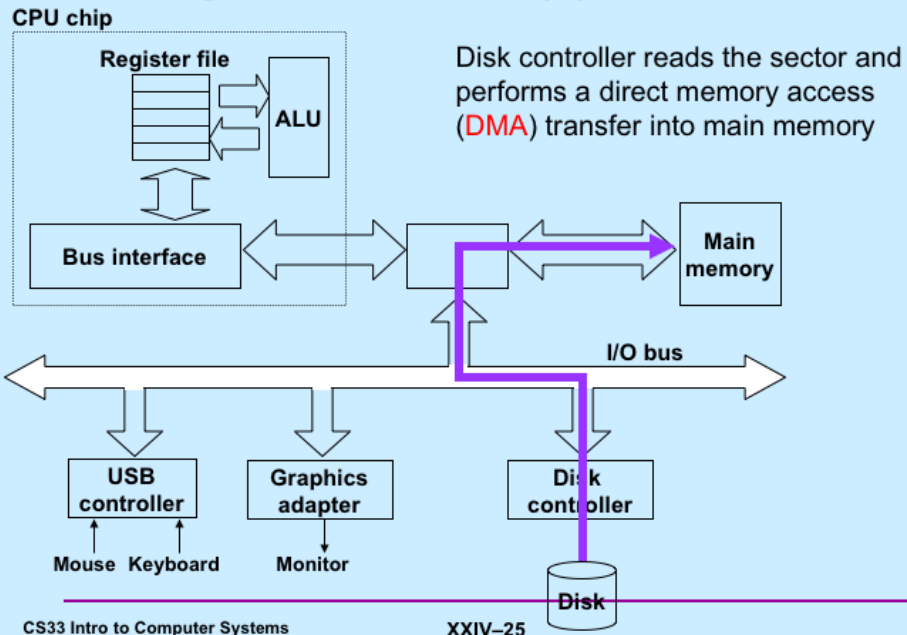
Supplied by CMU.

Reading a Disk Sector (1)



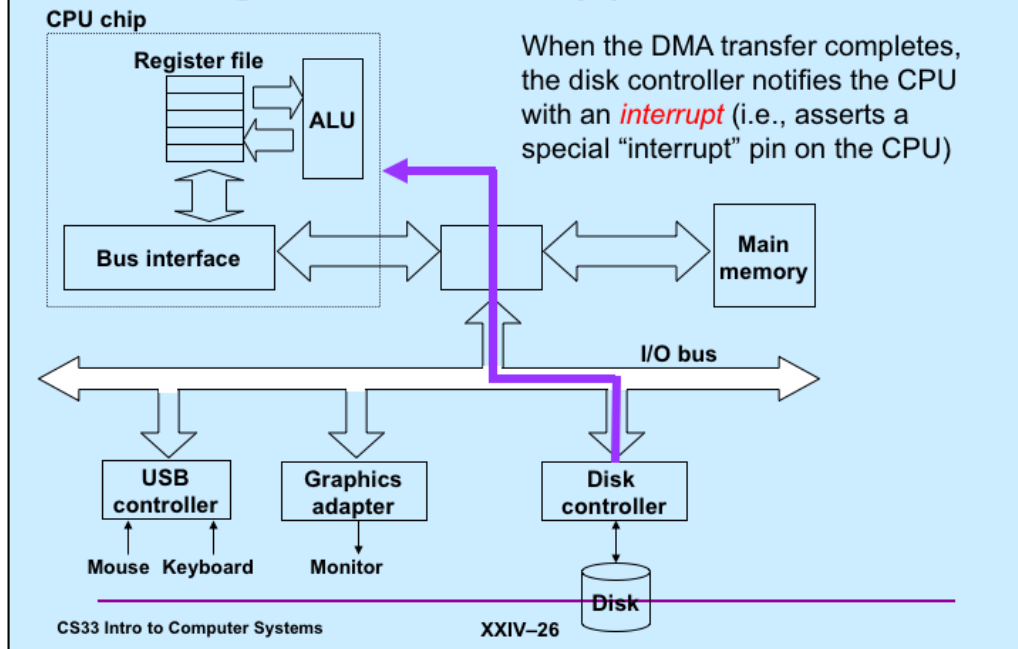
Supplied by CMU.

Reading a Disk Sector (2)



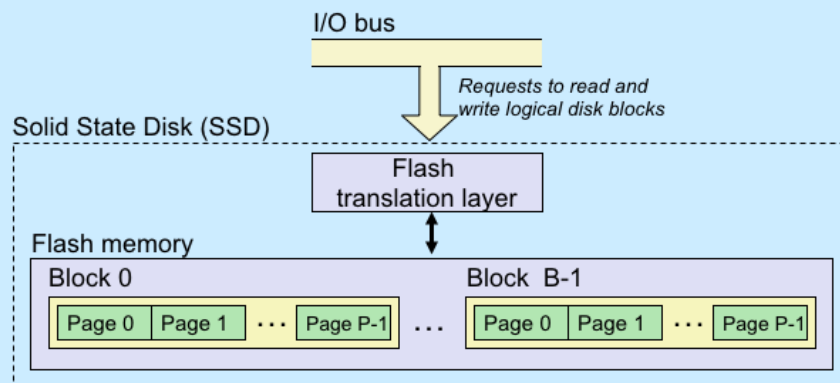
Supplied by CMU.

Reading a Disk Sector (3)



Supplied by CMU.

Solid-State Disks (SSDs)



- **Pages:** 512KB to 4KB; blocks: 32 to 128 pages
- **Data read/written in units of pages**
- **Page can be written only after its block has been erased**
- **A block wears out after 100,000 repeated writes**

Supplied by CMU.

SSD Performance Characteristics

Sequential read tput	250 MB/s	Sequential write tput	170 MB/s
Random read tput	140 MB/s	Random write tput	14 MB/s
Random read access	30 us	Random write access	300 us

- **Why are random writes so slow?**
 - erasing a block is slow (around 1 ms)
 - modifying a page triggers a copy of all useful pages in the block
 - » find a used block (new block) and erase it
 - » write the page into the new block
 - » copy other pages from old block to the new block

Supplied by CMU.


SSD Tradeoffs vs Rotating Disks

- **Advantages**
 - no moving parts → faster, less power, more rugged
- **Disadvantages**
 - have the potential to wear out
 - » mitigated by “wear-leveling logic” in flash translation layer
 - » e.g. Intel X25 guarantees 1 petabyte (10^{15} bytes) of random writes before they wear out
 - in 2010, about 100 times more expensive per byte
 - in 2017, about 6 times more expensive per byte
- **Applications**
 - smart phones, laptops
 - Apple “Fusion” drives

Supplied by CMU.

Reading a File on a Rotating Disk

- Suppose the data of a file are stored on consecutive disk sectors on one track
 - this is the best possible scenario for reading data quickly
 - » single seek required
 - » single rotational delay
 - » all sectors read in a single scan




Quiz 2

We have two files on the same (rotating) disk. The first file's data resides in consecutive sectors on one track, the second in consecutive sectors on another track. It takes a total of t seconds to read all of the first file then all of the second file.

Now suppose the files are read concurrently, perhaps a sector of the first, then a sector of the second, then the first, then the second, etc. Compared to reading them sequentially, this will take

- a) less time
- b) about the same amount of time
- c) more time



Quiz 3

You've replaced the rotating disk on your computer with a solid-state disk. The data of the two files are again each in consecutive locations. Suppose it still takes a total of t seconds to read the first file then the second. Suppose it took u seconds to read the two files concurrently on the rotating disk. It takes v seconds to read them concurrently on the SSD.

- a) $v < u$ (faster on the SSD)
- b) $v \approx u$ (about the same)
- c) $v > u$ (slower on the SSD)

Storage Trends

SRAM

Metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	19,200	2,900	320	256	100	75	60	320
access (ns)	300	150	35	15	3	2	1.5	200

DRAM

Metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	8,000	880	100	30	1	0.1	0.06	130,000
access (ns)	375	200	100	70	60	50	40	9
typical size (MB)	0.064	0.256	4	16	64	2,000	8,000	125,000

Disk

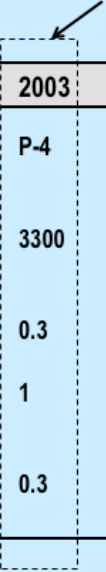
Metric	1980	1985	1990	1995	2000	2005	2010	2010:1980
\$/MB	500	100	8	0.30	0.01	0.005	0.0003	1,600,000
access (ms)	87	75	28	10	8	4	3	29
typical size (MB)	1	10	160	1,000	20,000	160,000	1,500,000	1,500,000

Supplied by CMU.

Current (late 2017) prices for SRAM vary a fair amount. As of 10/31, they range from \$11.99/MB to \$23.08/MB (the latter is via Amazon prime and includes free 2nd-day shipping).

CPU Clock Rates

Inflection point in computer history
when designers hit the “Power Wall”




	1980	1990	1995	2000	2003	2005	2010	2010:1980
CPU	8080	386	Pentium	P-III	P-4	Core 2	Core i7	---
Clock rate (MHz)	1	20	150	600	3300	2000	2500	2500
Cycle time (ns)	1000	50	6	1.6	0.3	0.50	0.4	2500
Cores	1	1	1	1	1	2	4	4
Effective cycle time (ns)	1000	50	6	1.6	0.3	0.25	0.1	10,000

Supplied by CMU.

The CPU-Memory Gap

The gap widens between DRAM, disk, and CPU speeds

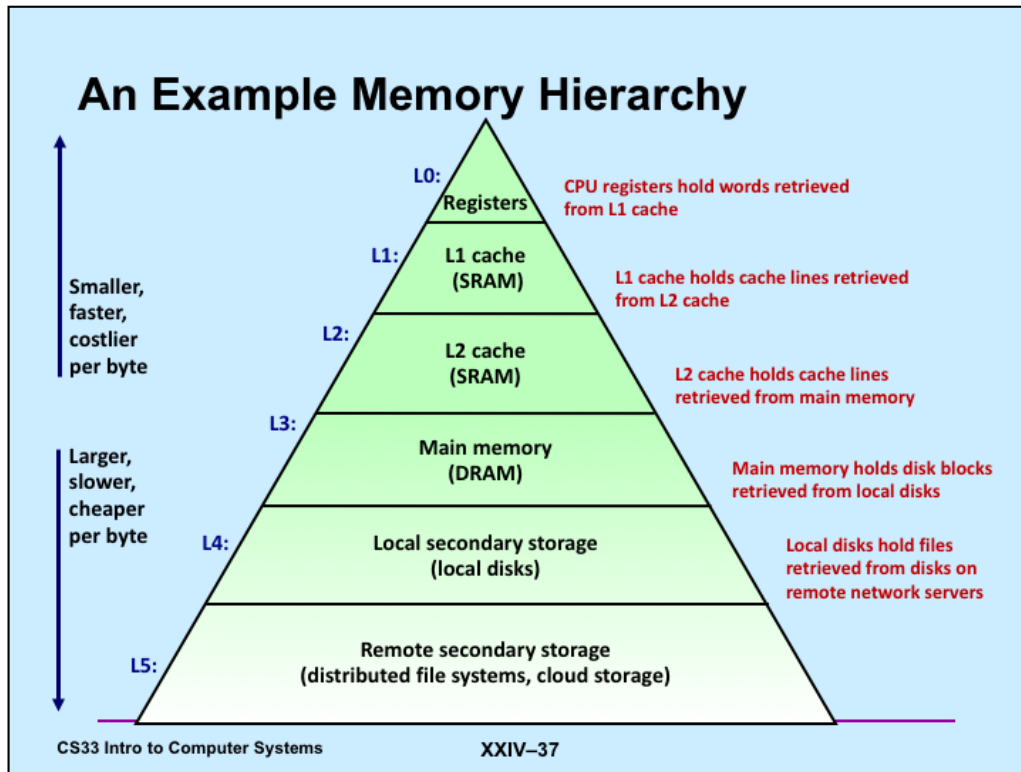


Supplied by CMU.

Memory Hierarchies

- **Some fundamental and enduring properties of hardware and software:**
 - fast storage technologies cost more per byte, have less capacity, and require more power (heat!)
 - the gap between CPU and main memory speed is widening
 - well written programs tend to exhibit good locality
- **These fundamental properties complement each other beautifully**
- **They suggest an approach for organizing memory and storage systems known as a **memory hierarchy****

Supplied by CMU.



Supplied by CMU.

Putting Things Into Perspective ...

- **Reading from:**
 - ... the L1 cache is like grabbing a piece of paper from your desk (3 seconds)
 - ... the L2 cache is picking up a book from a nearby shelf (14 seconds)
 - ... main system memory is taking a 4-minute walk down the hall to talk to a friend
 - ... a hard drive is like leaving the building to roam the earth for one year and three months

This analogy is from <http://duartes.org/gustavo/blog/post/what-your-computer-does-while-you-wait> (definitely worth reading!).