

# CS 33

## Signals Part 2

# Job Control

```
$ who
```

```
– foreground job
```

```
$ multiprocessProgram
```

```
– foreground job
```

```
^Z
```

```
stopped
```

```
$ bg
```

```
[1] multiprocessProgram &
```

```
– multiprocessProgram becomes background job 1
```

```
$ longRunningProgram &
```

```
[2]
```

```
$ fg %1
```

```
multiprocessProgram
```

```
– multiprocessProgram is now the foreground job
```

```
^C
```

```
$
```

# Process Groups

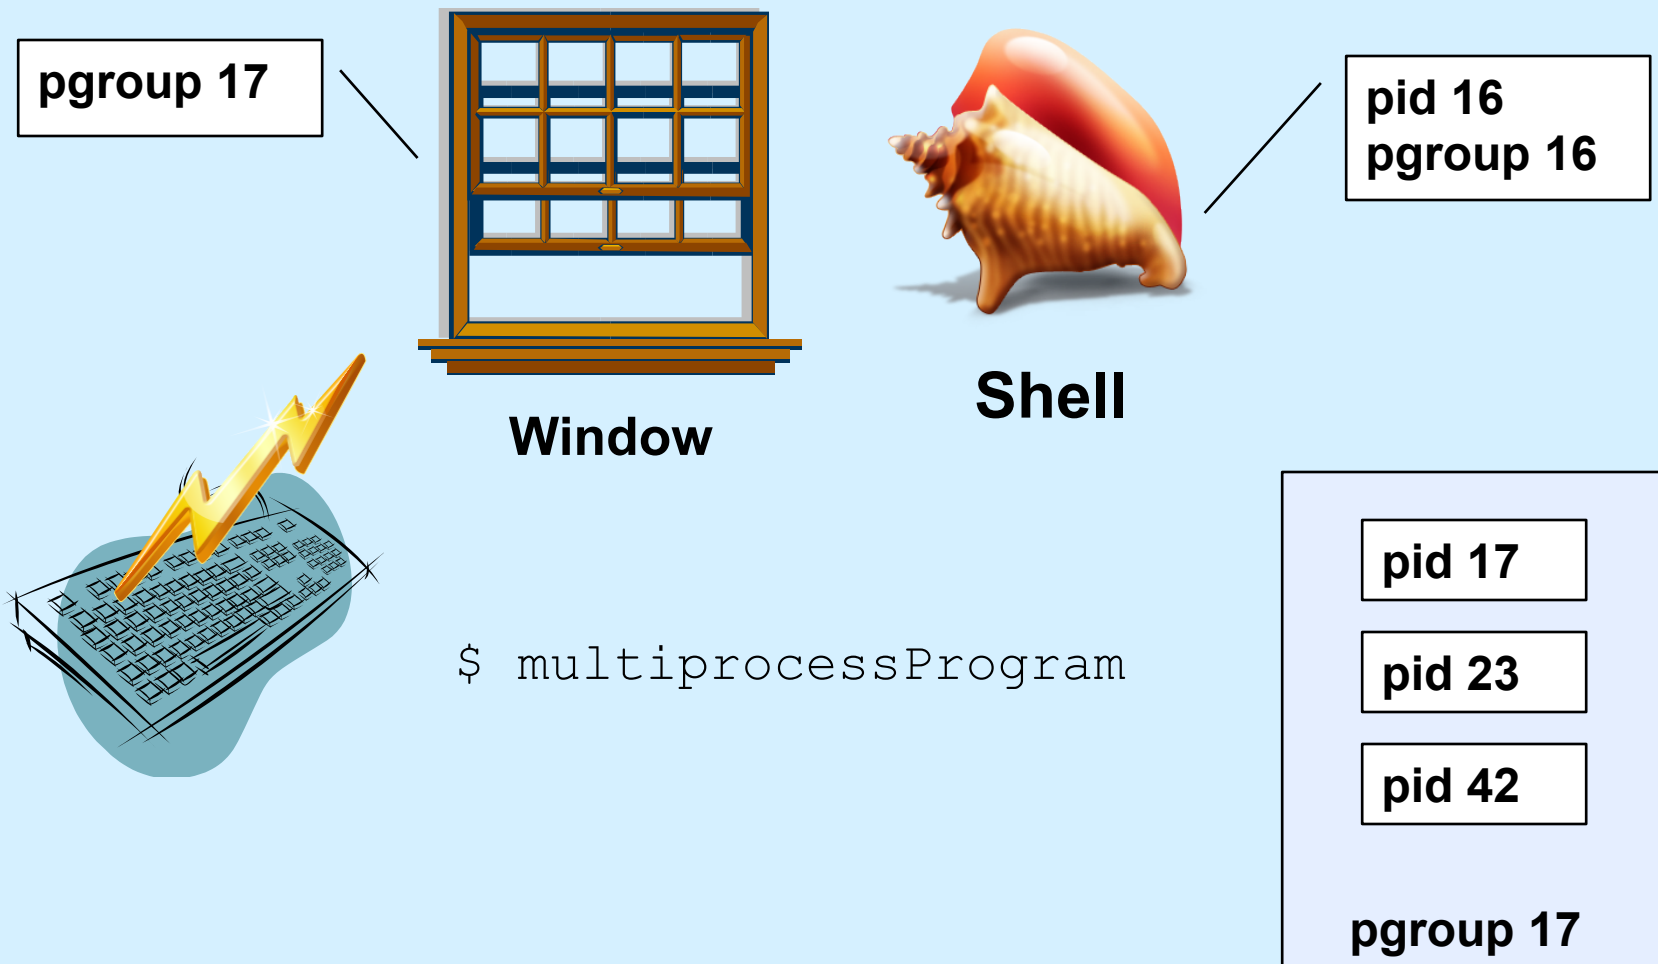
- **Set of processes sharing the window/keyboard**
  - sometimes called a *job*
- **Foreground process group/job**
  - currently associated with window/keyboard
  - receives keyboard-generated signals
- **Background process group/job**
  - not currently associated with window/keyboard
  - doesn't currently receive keyboard-generated signals

# Keyboard-Generated Signals

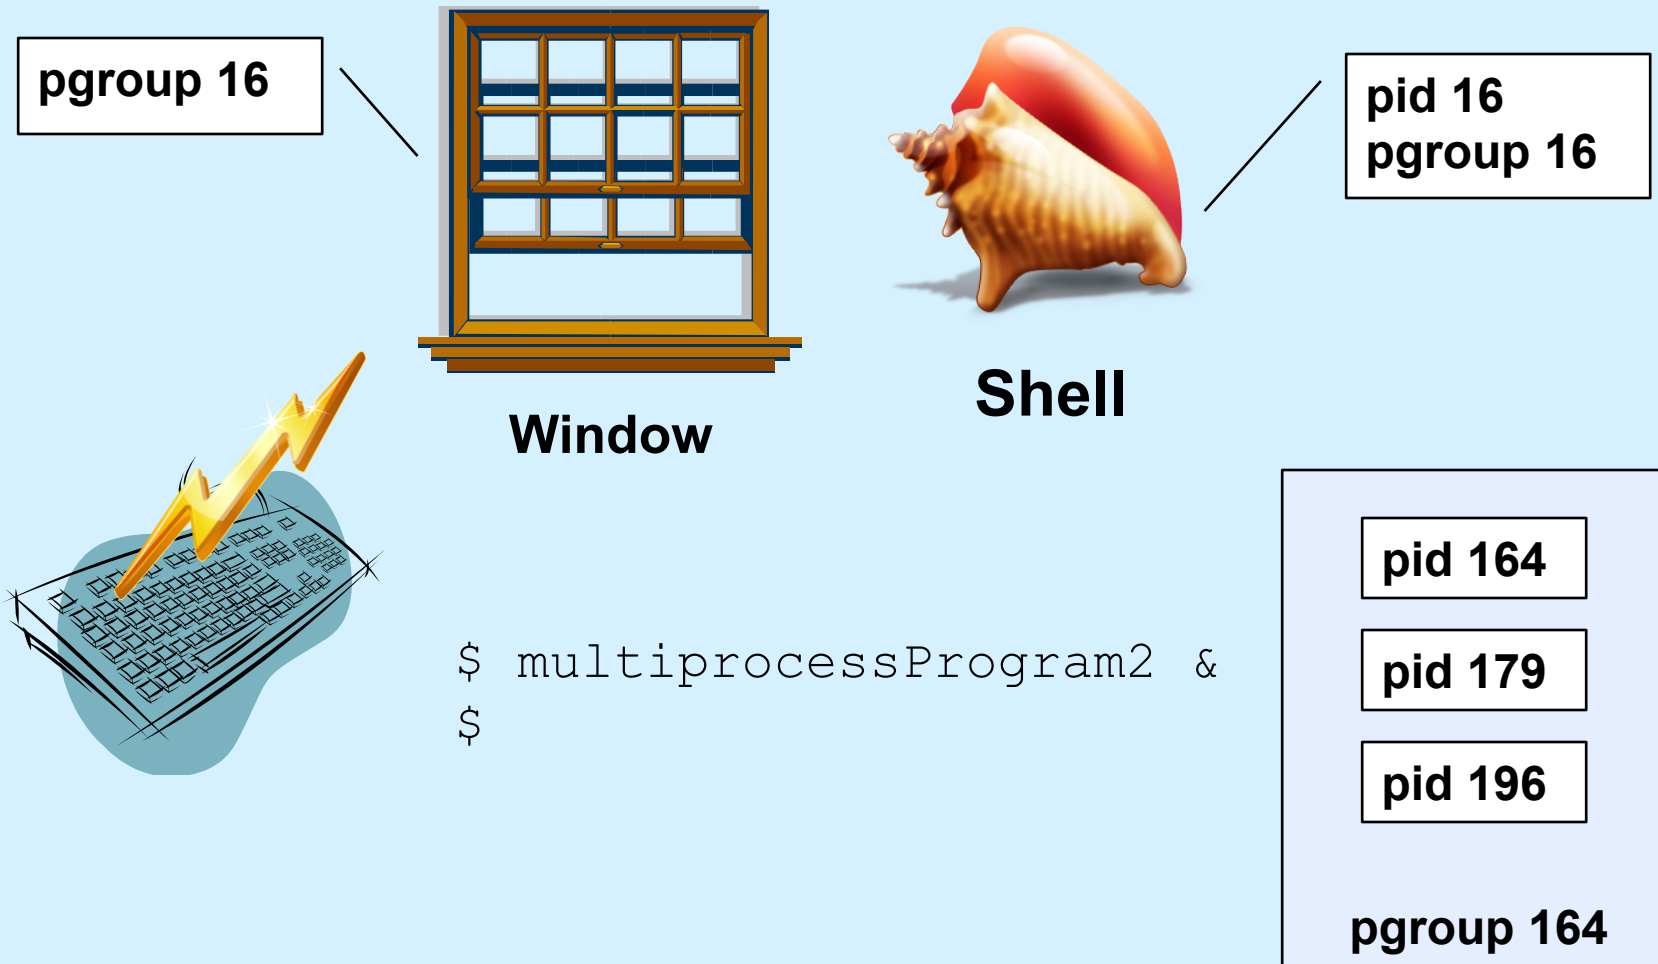
- You type ctrl-C
- How does the system know which process(es) to send the signal to?



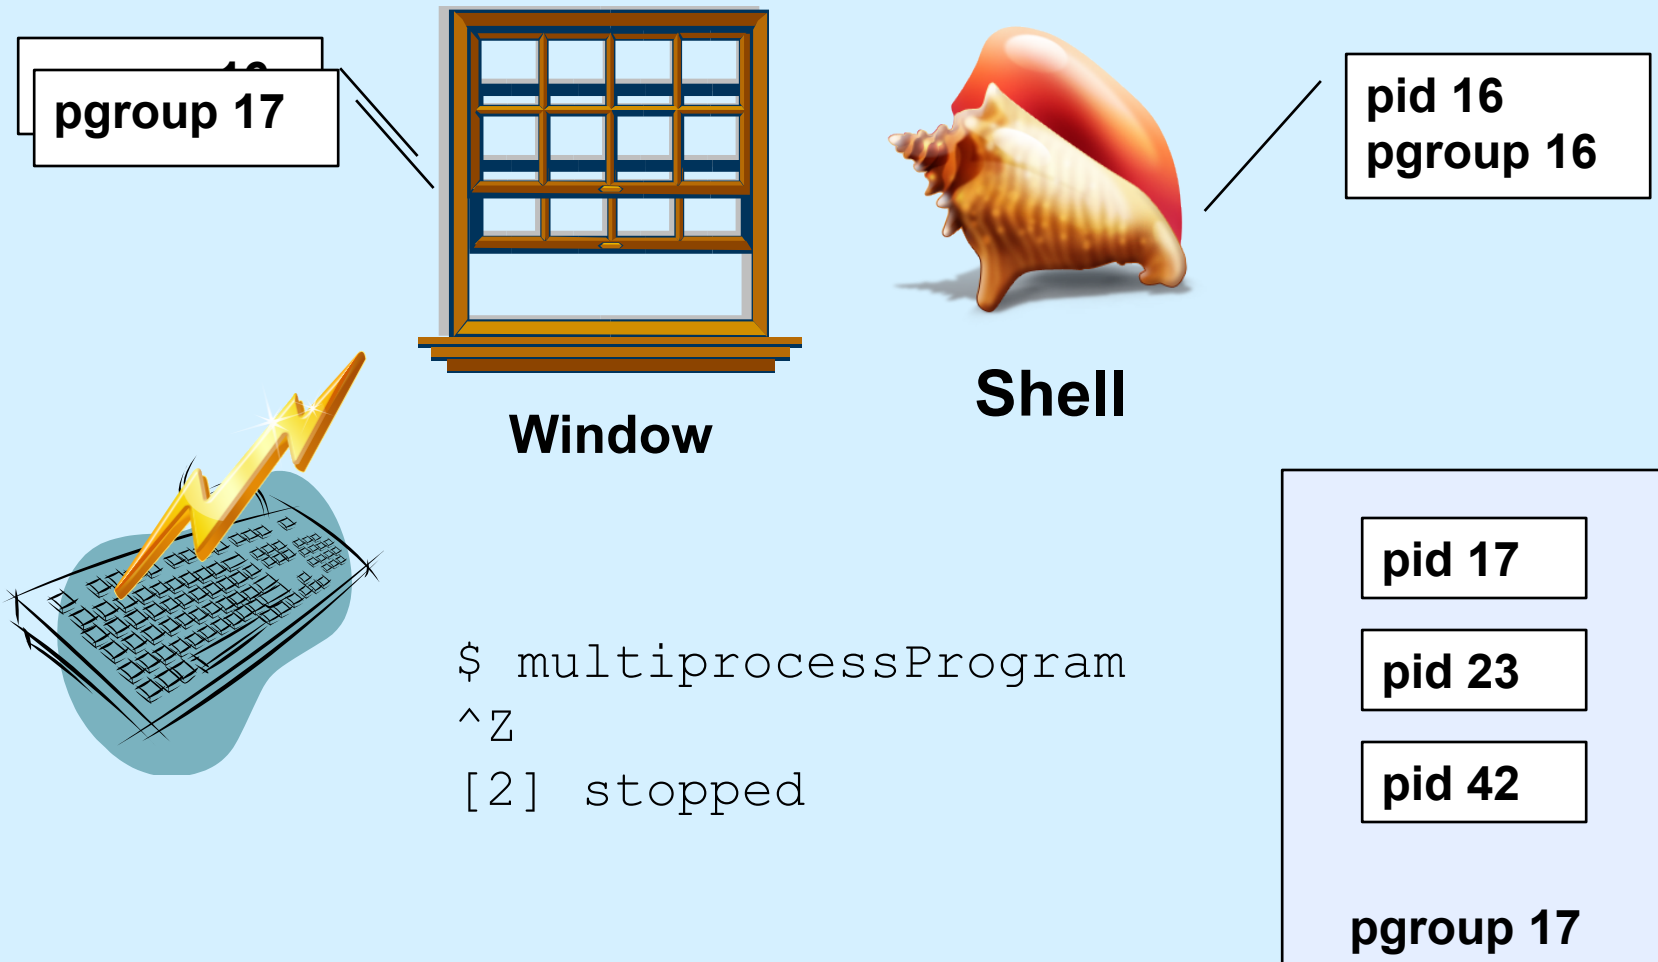
# Foreground Job



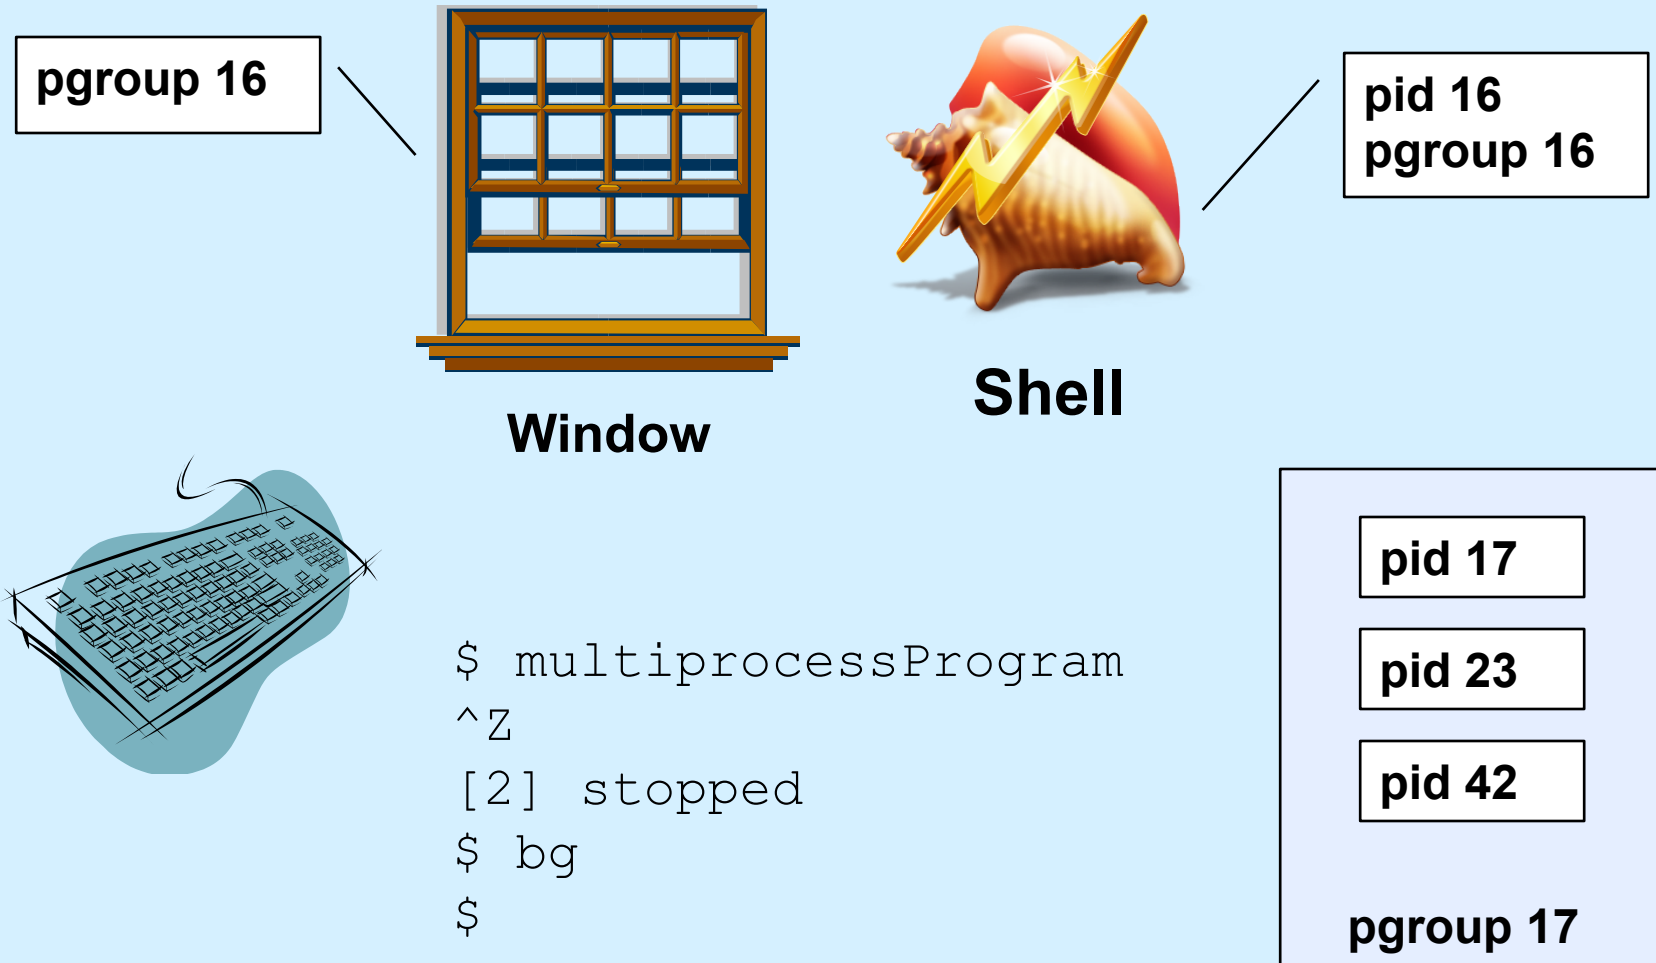
# Background Job



# Stopping a Foreground Job

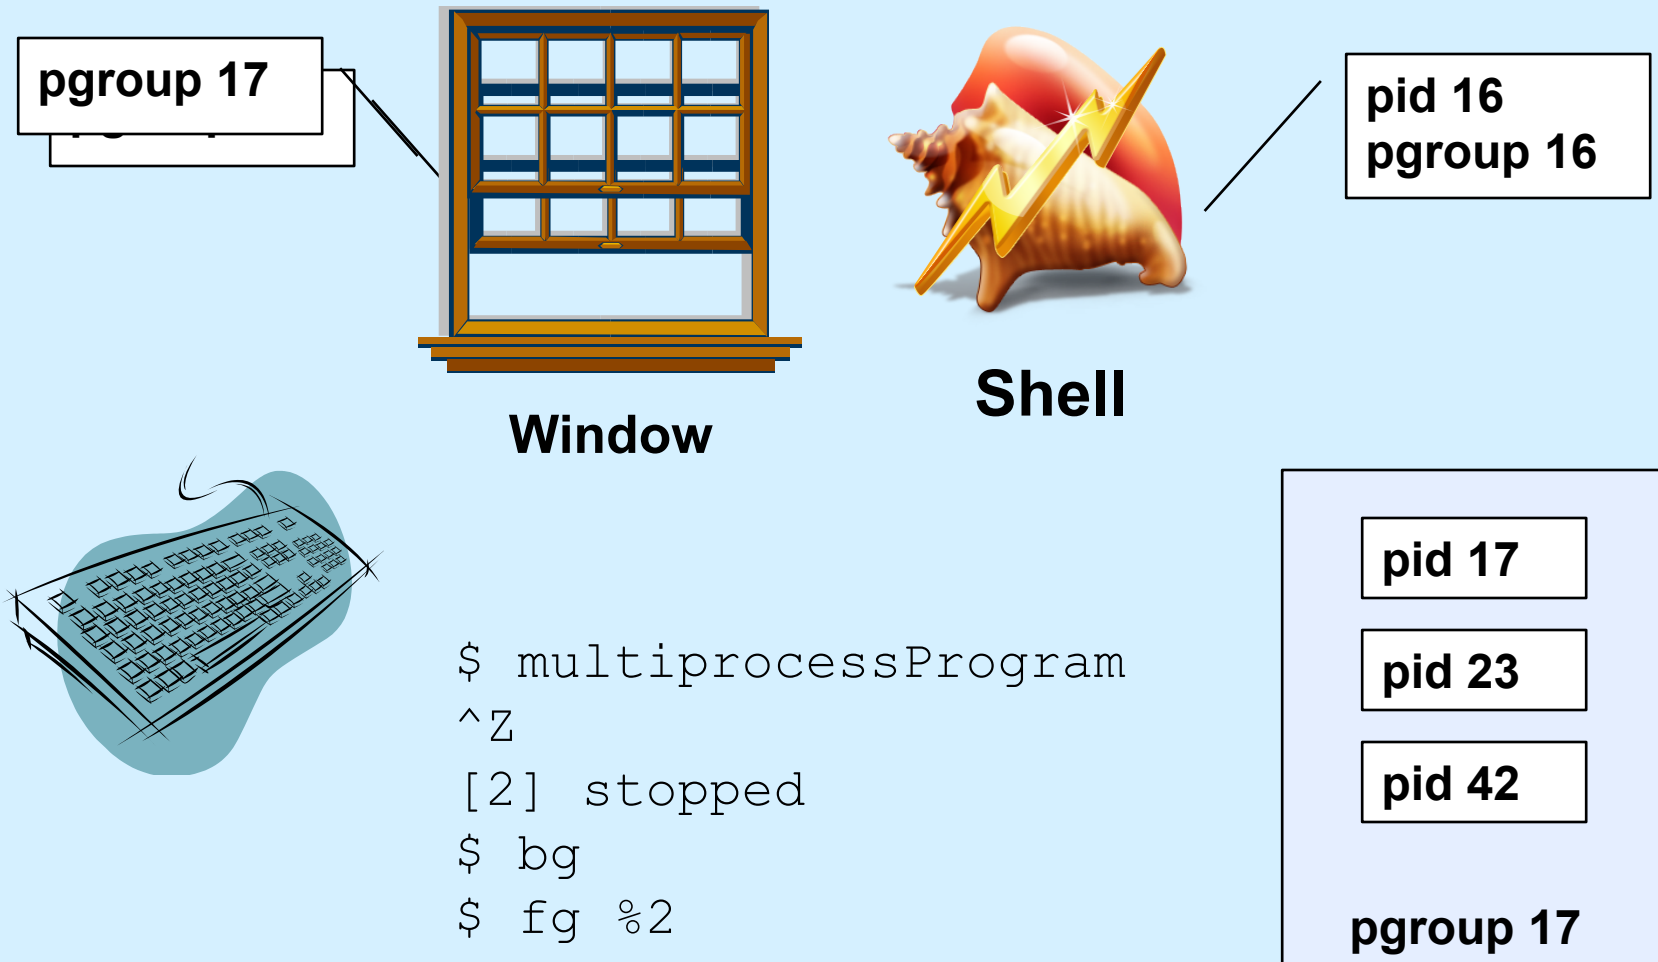


# Backgrounding a Stopped Job





# Foregrounding a Job



# Quiz 1

```
$ long_running_prog1 &  
$ long_running_prog2  
^Z  
[2] stopped  
$ ^C
```

**Which process group receives the SIGINT signal?**

- a) the one containing the shell**
- b) the one containing  
long\_running\_prog1**
- c) the one containing  
long\_running\_prog2**

# Creating a Process Group

```
if (fork() == 0) {  
    // child  
    setpgid(0, 0);  
    /* puts current process into a  
       new process group whose ID is  
       the process's pid.  
       Children of this process will be in  
       this process's process group.  
    */  
    ...  
    execv(...);  
}  
// parent
```

# Setting the Foreground Process Group

```
tcsetpgrp(fd, pgid);  
    // sets the process group of the  
    // terminal (window) referenced by  
    // file descriptor fd to be pgid
```


# Background Input and Output

- **Background process reads from keyboard**
  - the keyboard really should be reserved for foreground process
  - background process gets SIGTTIN
    - » suspends it by default
- **Background process writes to display**
  - display also used by foreground process
  - could be willing to share
  - background process gets SIGTTOU
    - » suspends it (by default)
    - » but reasonable to ignore it

# Kill: Details

- `int kill(pid_t pid, int sig)`
  - if *pid* > 0, signal *sig* sent to process *pid*
  - if *pid* == 0, signal *sig* sent to all processes in the caller's process group
  - if *pid* == -1, signal *sig* sent to all processes in the system for which sender has permission to do so
  - if *pid* < -1, signal *sig* is sent to all processes in process group *-pid*

# Process Life Cycle



# Reaping: Zombie Elimination

- **Shell must call `waitpid` on each child**
  - easy for foreground processes
  - what about background?

```
pid_t waitpid(pid_t pid, int *status, int options);
```

– ***pid* options:**

- < -1** any child process whose process group is |pid|
- 1** any child process
- 0** any child process whose process group is that of caller
- > 0** process whose ID is equal to pid

– `wait(&status)` **is equivalent to** `waitpid(-1, &status, 0)`



## (continued)

```
pid_t waitpid(pid_t pid, int *status, int options);
```

– *options* are some combination of the following

» **WNOHANG**

- return immediately if no child has exited (returns 0)

» **WUNTRACED**

- also return if a child has stopped (been suspended)

» **WCONTINUED**

- also return if a child has been continued (resumed)

# When to Call `waitpid`

- Shell reports status only when it is about to display its prompt
  - thus sufficient to check on background jobs just before displaying prompt

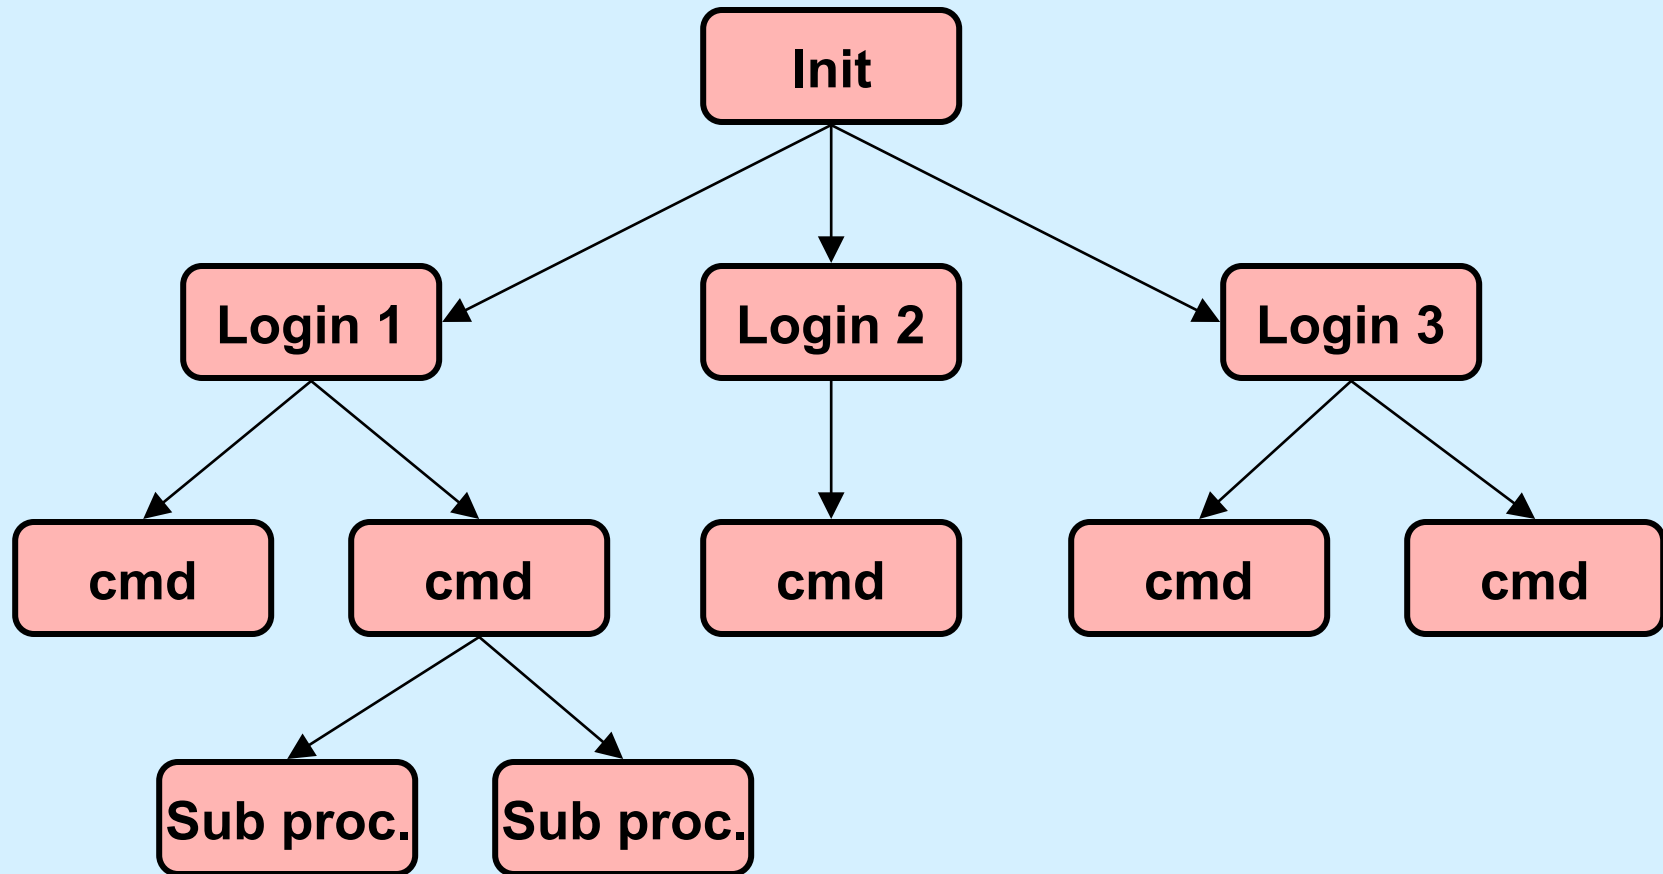
# `waitpid status`

- **`WIFEXITED(*status)`**: 1 if the process terminated normally and 0 otherwise
- **`WEXITSTATUS(*status)`**: argument to `exit`
- **`WIFSIGNALED(*status)`**: 1 if the process was terminated by a signal and 0 otherwise
- **`WTERMSIG(*status)`**: the signal which terminated the process if it terminated by a signal
- **`WIFSTOPPED(*status)`**: 1 if the process was stopped by a signal
- **`WSTOPSIG(*status)`**: the signal which stopped the process if it was stopped by a signal
- **`WIFCONTINUED(*status)`**: 1 if the process was resumed by `SIGCONT` and 0 otherwise

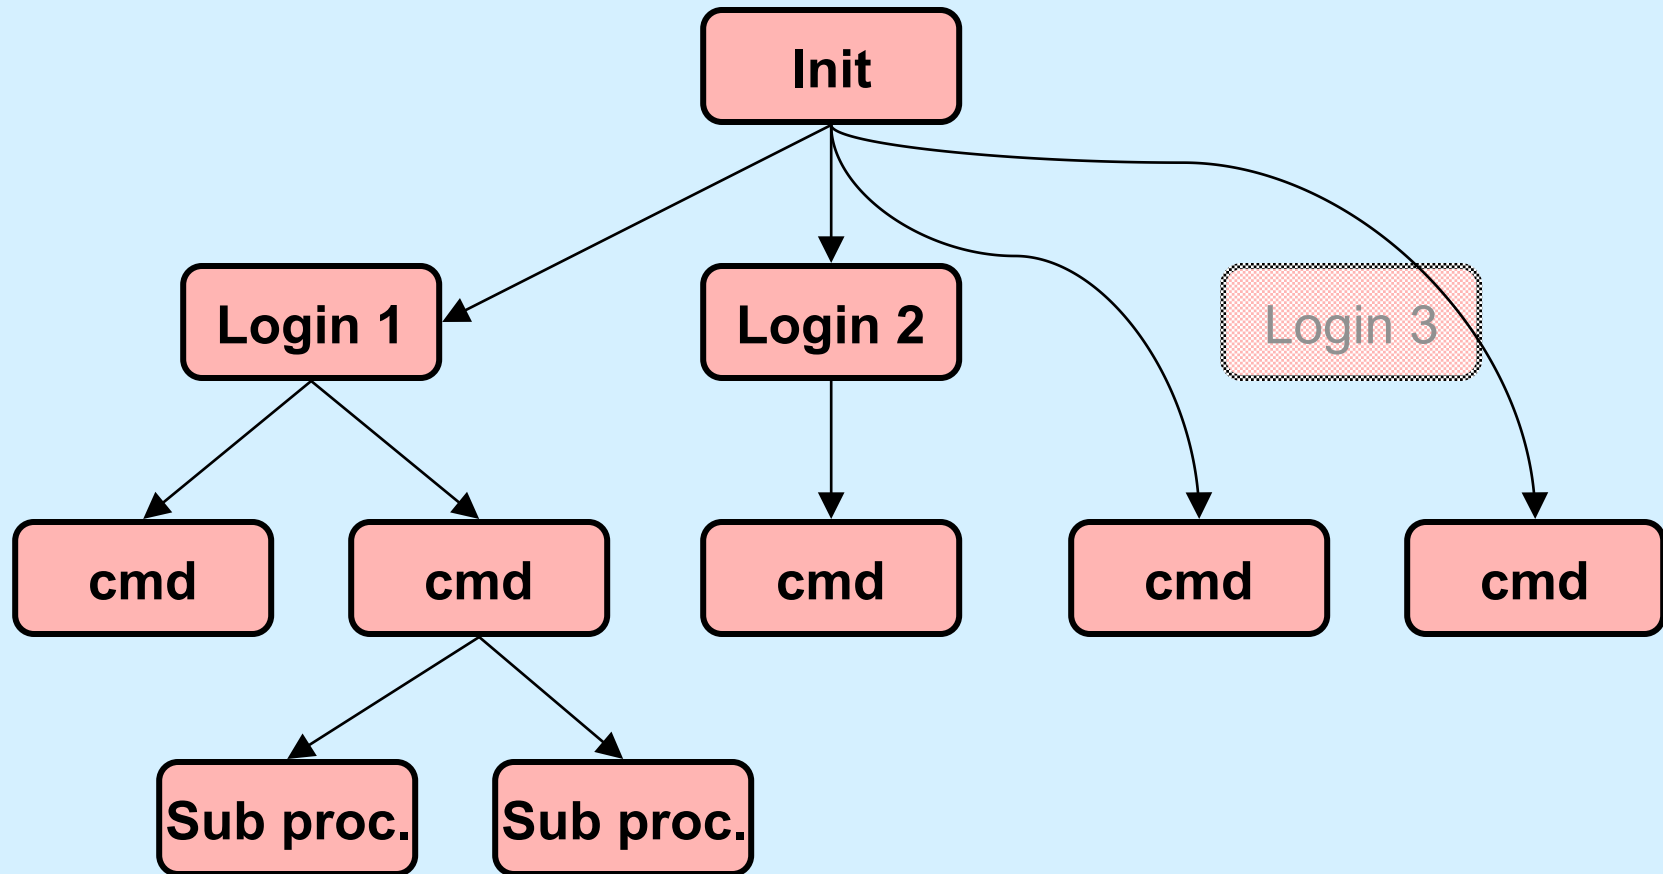
# Example (in Shell)

```
int wret, status;
while ((wret = waitpid(-1, &wstatus, WNOHANG|WUNTRACED)) > 0) {
    // examine all children who've terminated or stopped
    if (WIFEXITED(wstatus)) {
        // terminated normally
        ...
    }
    if (WIFSIGNALED(wstatus)) {
        // terminated by a signal
        ...
    }
    if (WIFSTOPPED(wstatus)) {
        // stopped
        ...
    }
}
```


# Process Relationships (1)



# Process Relationships (2)



# Process Relationships (3)



# Signals, Fork, and Exec

```
// set up signal handlers ...
if (fork() == 0) {
    // what happens if child gets signal?
    ...
    signal(SIGINT, SIG_IGN);
    signal(SIGFPE, handler);
    signal(SIGQUIT, SIG_DFL);
    execv("new prog", argv, NULL);
    // what happens if SIGINT, SIGFPE,
    // or SIGQUIT occur?
}
```

---



# Dealing with Failure

- *fork*, *execv*, *wait*, *kill* directly invoke the operating system
- Sometimes the OS says no
  - usually because you did something wrong
  - sometimes because the system has run out of resources
  - system calls return **-1** to indicate a problem

# Reporting Failure

- Integer error code placed in global variable *errno*

```
int errno;
```

- “man 3 errno” lists all possible error codes and meanings
- to print out meaning of most recent error

```
perror("message");
```

# Fork

```
int main( ) {  
    pid_t pid;  
    while(1) {  
        if ((pid = fork()) == -1) {  
            perror("fork");  
            exit(1);  
        }  
        ...  
    }  
}
```

# Exec

```
int main( ) {  
    if (fork() == 0) {  
        char *argv[] = {"garbage", 0};  
        execv("/garbage", argv);  
        /* if we get here, there was an  
           error! */  
        perror("execv: garbage");  
        exit(1);  
    }  
}
```

# Signals and Blocking System Calls

- **What if a signal is generated while a process is blocked in a system call?**
  - 1) deal with it when the system call completes
  - 2) interrupt the system call, deal with signal, resume system callor
  - 3) interrupt system call, deal with signal, return from system call with indication that something happened

# Interrupted System Calls

```
while (read(fd, buffer, buf_size) == -1) {  
    if (errno == EINTR) {  
        /* interrupted system call - try again */  
        continue;  
    }  
    /* the error is more serious */  
    perror("big trouble");  
    exit(1);  
}
```

# Timed Out, Revisited

```
void timeout(int sig) {}

int main() {
    struct sigaction act;
    sigemptyset(&act.sa_mask);
    act.sa_flags = 0;
    act.sa_handler = timeout;
    sigaction(SIGALRM, &act,
        NULL);

    alarm(10);
    char password[128];
```

```
    if (read(2, password,
        128)) == -1) {
        if (errno == EINTR) {
            fprintf(stderr,
                "Timed out\n");
            return 1;
        }
        perror("read");
        exit(1);
    }
    alarm(0);
    UsePassword(password);

    return 0;
}
```

# Quiz 2

```
int ret;  
char buf[128] = fillbuf();  
  
ret = write(1, buf, 128);
```

- **The value of ret is:**
  - a) either -1 or 128
  - b) either -1, 0, or 128
  - c) any integer in the range [-1, 128]



# Interrupted While Underway

```
remaining = total_count;
bptr = buf;
for ( ; ; ) {
    num_xfrd = write(fd, bptr,
                     remaining) ;
    if (num_xfrd == -1) {
        if (errno == EINTR) {
            /* interrupted early */
            continue;
        }
        perror("big trouble");
        exit(1);
    }

    if (num_xfrd < remaining) {
        /* interrupted after the
           first step */
        remaining -= num_xfrd;
        bptr += num_xfrd;
        continue;
    }
    /* success! */
    break;
}
```

# Asynchronous Signals (1)

```
main( ) {  
    void handler(int) ;  
    signal(SIGINT, handler) ;  
  
    ...    /* long-running buggy code */  
  
}  
  
void handler(int sig) {  
    ...    /* die gracefully */  
    exit(1) ;  
  
}
```

# Asynchronous Signals (2)

```
computation_state_t  state;    long_running_procedure( ) {  
                                while (a_long_time) {  
main( ) {                      update_state(&state);  
    void handler(int);         compute_more( );  
                                }  
    signal(SIGINT, handler);    }  
  
    long_running_procedure( );  void handler(int sig) {  
}                               display(&state);  
                                }  
                                }
```

# Asynchronous Signals (3)

```
main( ) {  
    void handler(int);  
  
    signal(SIGINT, handler);  
  
    ... /* complicated program */  
  
    myput("important message\n");  
  
    ... /* more program */  
  
}  
  
void handler(int sig) {  
  
    ... /* deal with signal */  
  
    myput("equally important "  
          "message\n");  
}
```

# Asynchronous Signals (4)

```
char buf[BSIZE];
int pos;
void myput(char *str) {
    int i;
    int len = strlen(str);
    for (i=0; i<len; i++, pos++) {
        buf[pos] = str[i];
        if ((buf[pos] == '\n') || (pos == BSIZE-1)) {
            write(1, buf, pos+1);
            pos = -1;
        }
    }
}
```

# Async-Signal Safety

- Which library routines are safe to use within signal handlers?

- abort	- dup2	- getppid	- readlink	- sigemptyset	- tcgetpgrp
- accept	- execl	- getsockname	- recv	- sigfillset	- tcseendbreak
- access	- execve	- getsockopt	- recvfrom	- sigismember	- tcsetattr
- aio_error	- _exit	- getuid	- recvmsg	- signal	- tcsetpgrp
- aio_return	- fchmod	- kill	- rename	- sigpause	- time
- aio_suspend	- fchown	- link	- rmdir	- sigpending	- timer_getoverrun
- alarm	- fcntl	- listen	- select	- sigprocmask	- timer_gettime
- bind	- fdatsync	- lseek	- sem_post	- sigqueue	- timer_settime
- cfgetispeed	- fork	- lstat	- send	- sigsuspend	- times
- cfgetospeed	- fpathconf	- mkdir	- sendmsg	- sleep	- umask
- cfsetispeed	- fstat	- mkfifo	- sendto	- sockatmark	- uname
- cfsetospeed	- fsync	- open	- setgid	- socket	- unlink
- chdir	- ftruncate	- pathconf	- setpgid	- socketpair	- utime
- chmod	- getegid	- pause	- setsid	- stat	- wait
- chown	- geteuid	- pipe	- setsockopt	- symlink	- waitpid
- clock_gettime	- getgid	- poll	- setuid	- sysconf	- write
- close	- getgroups	- posix_trace_event	- shutdown	- tcdrain	
- connect	- getpeername	- pselect	- sigaction	- tcflow	
- creat	- getpgrp	- raise	- sigaddset	- tcflush	
- dup	- getpid	- read	- sigdelset	- tcgetattr	

# Quiz 3

**Printf is not required to be async-signal safe.  
Can it be implemented so that it is?**

- a) no, it's inherently not async-signal safe**
- b) yes, but it would be so complicated, it's not done**
- c) yes, it can be easily made async-signal safe**