

# CS 33

**More Libraries**

**Pipes**

**Locking Files**

# A Problem

- **You've put together a library of useful functions**
  - **libgoodstuff.so**
- **Lots of people are using it**
- **It occurs to you that you can make it even better by adding an extra argument to a few of the functions**
  - **doing so will break all programs that currently use these functions**
- **You need a means so that old code will continue to use the old version, but new code will use the new version**


# A Solution

- **The two versions of your program coexist**
  - `libgoodstuff.so.1`
  - `libgoodstuff.so.2`
- **You arrange so that old code uses the old version, new code uses the new**
- **Most users of your code don't really want to have to care about version numbers**
  - they want always to link with `libgoodstuff.so`
  - and get the version that was current when they wrote their programs

# Versioning

```
$ gcc -fPIC -c goodstuff.c
$ ld -shared -soname libgoodstuff.so.1 \
-o libgoodstuff.so.1 goodstuff.o
$ ln -s libgoodstuff.so.1 libgoodstuff.so
$ gcc -o prog1 prog1.c -L. -lgoodstuff \
-Wl,-rpath .
$ vi goodstuff.c
$ gcc -fPIC -c goodstuff.c
$ ld -shared -soname libgoodstuff.so.2 \
-o libgoodstuff.so.2 goodstuff.o
$ rm -f libgoodstuff.so
$ ln -s libgoodstuff.so.2 libgoodstuff.so
$ gcc -o prog2 prog2.c -L. -lgoodstuff \
-Wl,-rpath .
```

# Interpositioning



# How To ...

```
int __wrap_puts(const char *s) {  
    int __real_puts(const char *);  
  
    write(2, "calling myputs: ", 16);  
    return __real_puts(s);  
}
```

# Compiling/Linking It

```
$ cat tputs.c
int main() {
    puts("This is a boring message.");
    return 0;
}
$ gcc -o tputs -Wl,--wrap=puts tputs.c myputs.c
$ ./tputs
calling myputs: This is a boring message.
$
```

# How To (Alternative Approach) ...

```
#include <dlfcn.h>

int puts(const char *s) {
    int (*pptr)(const char *);

    pptr = (int(*)())dlsym(RTLD_NEXT, "puts");

    write(2, "calling myputs: ", 16);
    return (*pptr)(s);
}
```



# What's Going On ...

- **gcc/ld**
  - compiles code
  - does static linking
    - » searches list of libraries
    - » adds references to shared objects
- **runtime**
  - program invokes *ld-linux.so* to finish linking
    - » maps in shared objects
    - » does relocation and procedure linking as required
  - *dlsym* invokes *ld-linux.so* to do more linking
    - » RTLD\_NEXT says to use the next (second) occurrence of the symbol

# Delayed Wrapping

- **LD\_PRELOAD**
  - environment variable checked by *ld-linux.so*
  - specifies additional shared objects to search (first) when program is started


# Environment Variables

- **Another form of exec**
  - `int execve(const char *filename,  
char *const argv[],  
char *const envp[]);`
- **envp is an array of strings, of the form**
  - `key=value`
- **programs can search for values, given a key**
- **example**
  - `PATH=~/bin:/bin:/usr/bin:/course/cs0330/bin`

# Example

```
$ gcc -o tputs tputs.c
$ ./tputs
This is a boring message.
$ LD_PRELOAD=./libmyputs.so.1; export LD_PRELOAD
$ ./tputs
calling myputs: This is a boring message.
$
```


# Mmapping Libraries



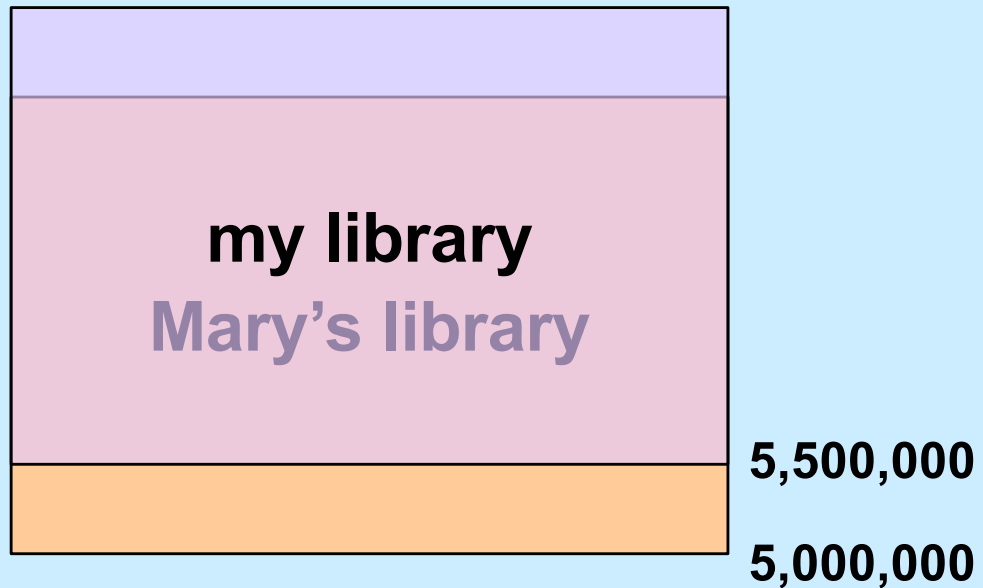
# Problem

- **How is relocation handled?**

# Pre-Relocation

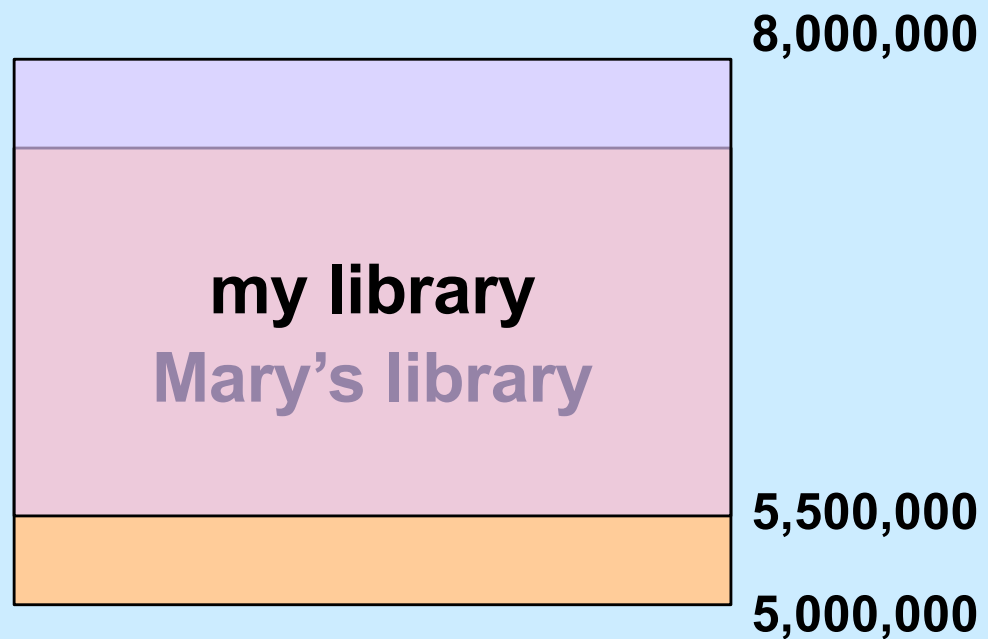


# But ...





# But ...



# Quiz 1


**We need to relocate all references to Mary's library in *my library*. What option should we give to *mmap* when we map *my library* into our address space?**

- a) the MAP\_SHARED option**
- b) the MAP\_PRIVATE option**
- c) mmap can't be used in this situation**

# Relocation Revisited

- **Modify shared code to effect relocation**
  - result is no longer shared!
- **Separate shared code from (unshared) addresses**
  - position-independent code (PIC)
  - code can be placed anywhere
  - addresses in separate private section
    - » pointed to by a register


# Mapping Shared Objects



# Mapping printf into the Address Space

- **Printf's text**
  - read-only
  - can it be shared?
    - » yes: use MAP\_SHARED
- **Printf's data**
  - read-write
  - not shared with other processes
  - initial values come from file
  - can mmap be used?
    - » MAP\_SHARED wouldn't work
      - changes made to data by one process would be seen by others
    - » MAP\_PRIVATE does work!
      - mapped region is initialized from file
      - changes are private

# Mapping printf



# Position-Independent Code

- **Produced by gcc when given the `-fPIC` flag**
- **Processor-dependent; x86-64:**
  - **each dynamic executable and shared object has:**
    - » **procedure-linkage table**
      - **shared, read-only executable code**
      - **essentially stubs for calling functions**
    - » **global-offset table**
      - **private, read-write data**
      - **relocated dynamically for each process**
    - » **relocation table**
      - **shared, read-only data**
      - **contains relocation info and symbol table**

# Global-Offset Table: Data References

Global Offset Table →

errno →

**errno address**

myglob →

**myglob address**



# Functions in Shared Objects

- Lots of them
- Many are never used
- Fix up linkages on demand

# An Example

```
int main( ) {  
    puts("Hello world\n");  
    ...  
    return 0;  
}
```

00000000000000006b0 <main>:

6b0:	55	push	%rbp
6b1:	48 89 e5	mov	%rsp,%rbp
6b4:	48 8d 3d 99 00 00 00	lea	0x99(%rip),%rdi
6bb:	e8 a0 fe ff ff	<b>callq</b>	<b>560 &lt;puts@plt&gt;</b>

...

# Before Calling puts

```
.PLT0:
    pushq GOT+8(%rip)
    jmp    *GOT+16(%rip)
    nop; nop
    nop; nop
.puts:
    jmp    *puts@GOT(%rip)
.putsnext
    pushq $putsRelOffset
    jmp    .PLT0
.PLT2:
    jmp    *name2@GOT(%rip)
.PLT2next
    pushq $name2RelOffset
    jmp    .PLT0
```

**Procedure-Linkage Table**

```
GOT:
    .quad _DYNAMIC
    .quad identification
    .quad ld-linux.so

puts:
    .quad .putsnext
name2:
    .quad .PLT2next
```

**Relocation info:**

GOT\_offset(puts) , symx(puts)

GOT\_offset(name2) , symx(name2)

**Relocation Table**

# After Calling puts

```
.PLT0:
    pushq GOT+8(%rip)
    jmp    *GOT+16(%rip)
    nop; nop
    nop; nop
.puts:
    jmp    *puts@GOT(%rip)
.putsnext
    pushq $putsRelOffset
    jmp    .PLT0
.PLT2:
    jmp    *name2@GOT(%rip)
.PLT2next
    pushq $name2RelOffset
    jmp    .PLT0
```

Procedure-Linkage Table

```
GOT:
    .quad _DYNAMIC
    .quad identification
    .quad ld-linux.so

puts:
    .quad puts
name2:
    .quad .PLT2next
```


Relocation info:

GOT\_offset(puts) , symx(puts)


GOT\_offset(name2) , symx(name2)

Relocation Table


# Interprocess Communication (IPC): Pipes




# Interprocess Communication: Same Machine I



# Interprocess Communication: Same Machine II




# Interprocess Communication: Different Machines





# Intramachine IPC

```
$cs1ab2e who | wc -l
```



# Intramachine IPC

```
$cs1ab2e who | wc -l
```

```
int fd[2];
pipe(fd);
if (fork() == 0) {
    close(fd[0]);
    close(1);
    dup(fd[1]); close(fd[1]);
    execl("/usr/bin/who", "who", 0); // who sends output to pipe
}
if (fork() == 0) {
    close(fd[1]);
    close(0);
    dup(fd[0]); close(fd[0]);
    execl("/usr/bin/wc", "wc", "-l", 0); // wc's input is from pipe
}
close(fd[1]); close(fd[0]);
// ...
```



# Intermachine Communication

- **Can pipes be made to work across multiple machines?**

- **covered soon ...**

- » **what happens when you type**

- `who | ssh cs1ab3a wc -l`

- ?**


# Sharing Files

- **You're doing a project with a partner**
- **You code it as one 15,000-line file**
  - the first 7,500 lines are yours
  - the second 7,500 lines are your partner's
- **You edit the file, changing 6,000 lines**
  - it's now 5am
- **Your partner completes her changes at 5:01am**
- **At 5:02am you look at the file**
  - your partner's changes are there
  - yours are not

# Lessons

- **Never work with a partner**
- **Use more than one file**
- **Read up on git**
- **Use an editor and file system that support file locking**


# What We Want ...



# Types of Locks

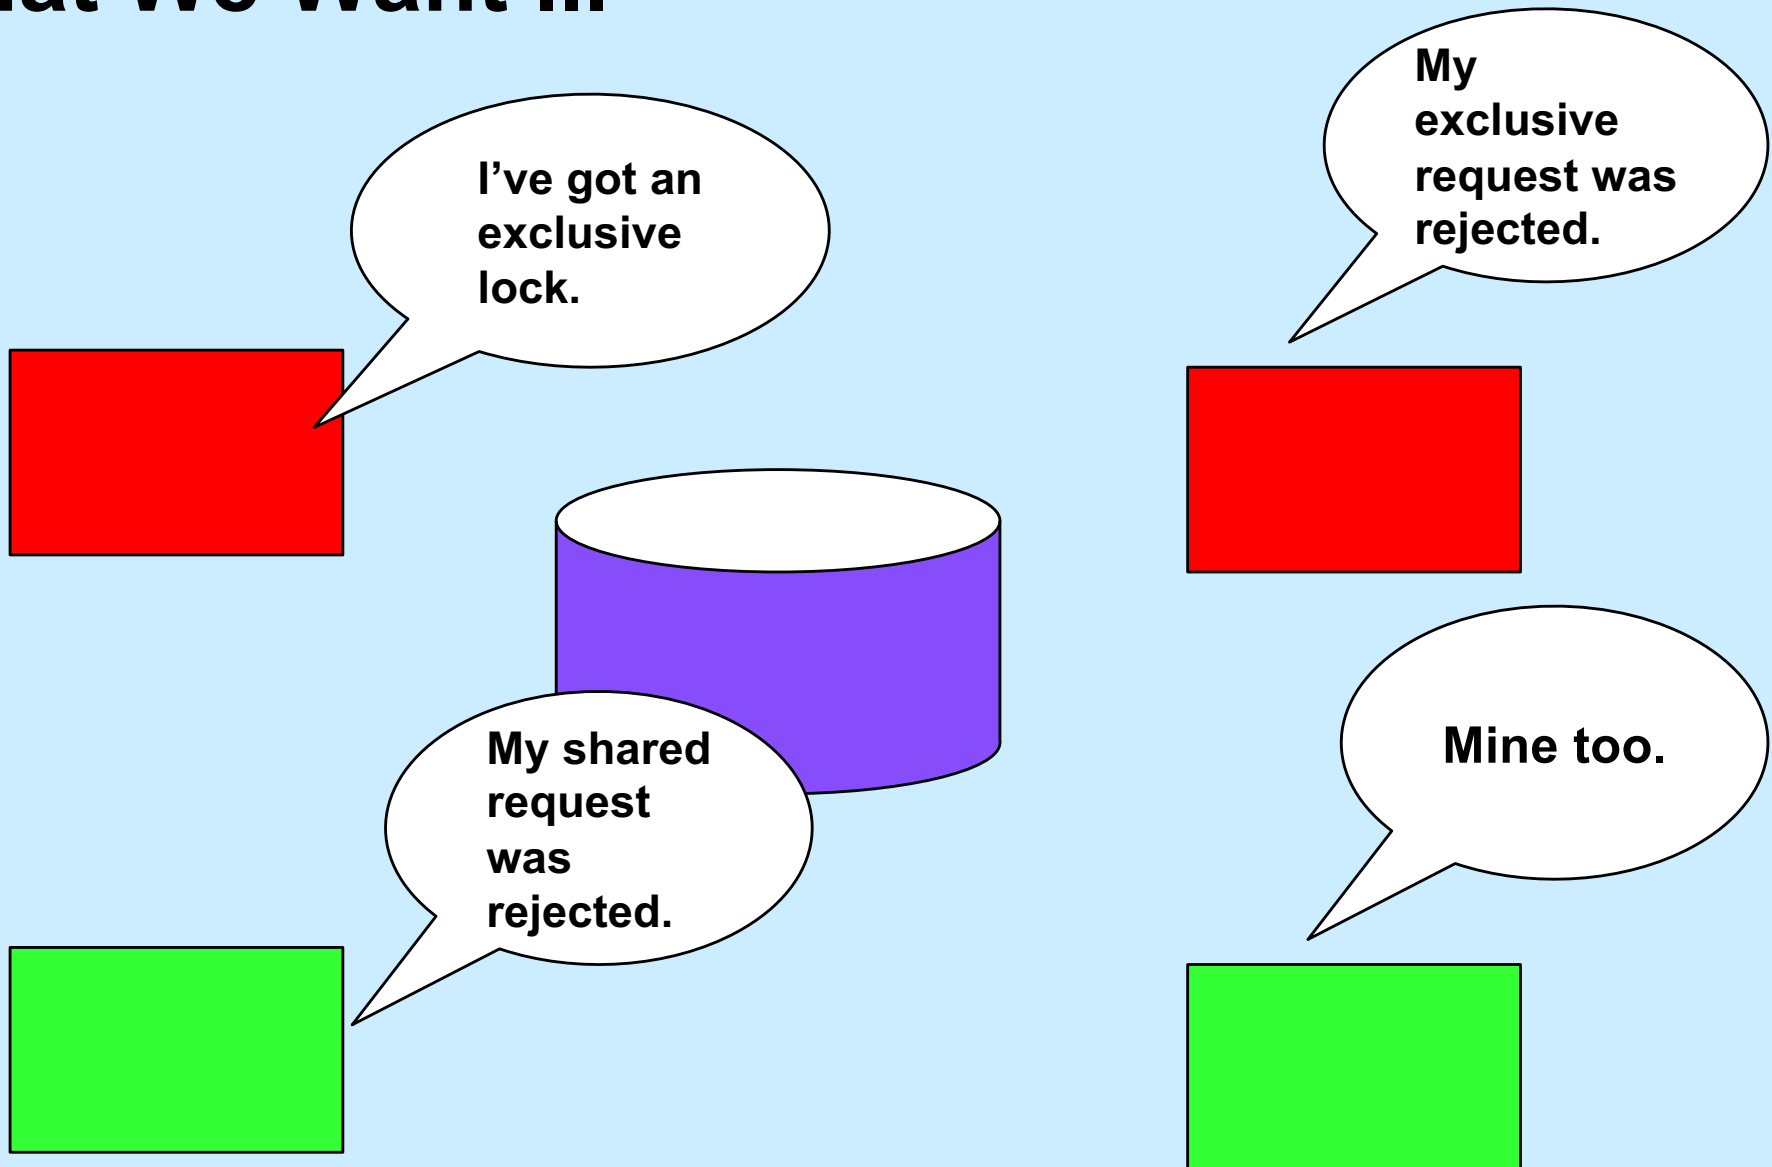
- **Shared (readers) locks**
  - any number may have them at same time
  - may not be held when an exclusive lock is held
- **Exclusive (writers) locks**
  - only one at a time
  - may not be held when a shared lock is held

# What We Want ...





# What We Want ...



# Locking Files

- Early Unix didn't support file locking
- How did people survive?

- `open("file.lck", O_RDWR|O_CREAT|O_EXCL, 0666);`
    - » operation fails if *file.lck* exists, succeeds (and creates *file.lck*) otherwise
    - » requires cooperative programs

# Locking Files (continued)

- How it's done in “modern” Unix
  - “advisory locks” may be placed on files
    - » may request shared (readers) or exclusive (writers) lock
      - *fcntl* system call
    - » either succeeds or fails
    - » *open*, *read*, *write* always work, regardless of locks
    - » a lock applies to a specified range of bytes, not necessarily the whole file
    - » requires cooperative programs

# Locking Files (still continued)

- **How to:**

```
struct flock fl;
fl.l_type = F_RDLCK;           // read lock
// fl.l_type = F_WRLCK;       // write lock
// fl.l_type = F_UNLCK;       // unlock
fl.l_whence = SEEK_SET;        // starting where
fl.l_start = 0;                // offset
fl.l_len = 0;                  // how much? (0 = whole file)
fd = open("file", O_RDWR);
if (fcntl(fd, F_SETLK, &fl) == -1)
    if ((errno == EACCES) || (errno == EAGAIN))
        // didn't get lock
    else
        // something else is wrong
else
    // got the lock!
```

---

# Locking Files (yet still continued)

- **Making locks mandatory:**
  - if the file's permissions have group execute permission off and set-group-ID on, then locking is enforced
    - » *read, write* fail if file is locked by someone other than the caller
  - however ...
    - » difficult to implement on distributed file systems (such as used at Brown CS)

# Quiz 2

- Your program currently has a shared lock on a portion of a file. It would like to “upgrade” the lock to be an exclusive lock. Would there be any problems with adding an option to *fcntl* that would allow the holder of a shared lock to wait until it’s possible to upgrade to an exclusive lock, then do the upgrade?
  - a) at least one major problem
  - b) either no problems whatsoever or some easy-to-deal-with problems