



# CS 33

## Network Programming

# Communicating Over the Internet



# The Internet




# Names and Addresses

- **cslab1c.cs.brown.edu**
  - the name of a computer on the internet
  - mapped to an internet address
- **www.nyt.com**
  - the name of a service on the internet
  - mapped to a number of internet addresses
- **How are names mapped to addresses?**
  - domain name service (DNS): a distributed database
- **How are the machines corresponding to internet addresses found?**
  - with the aid of various routing protocols

# Internet Addresses

- **IP (internet protocol) address**
  - one per network interface
  - **32 bits (IPv4)**
    - » 5527 per acre of RI
    - » 25 per acre of Texas
  - **128 bits (IPv6)**
    - » 1.6 billion per cubic mile of a sphere whose radius is the mean distance from the Sun to the (former) planet Pluto
- **Port number**
  - one per application instance per machine
  - **16 bits**
    - » port numbers less than 1024 are reserved for privileged applications



# Notation

- **Addresses (assume IPv4: 32-bit addresses)**
  - written using dot notation
    - » 128.48.37.1
      - dots separate bytes
  - address plus port (1426):
    - » 128.48.37.1:1426

# Reliability

- **Two possibilities**
  - **don't worry about it**
    - » **just send it**
      - **if it arrives at its destination, that's good!**
        - **no verification**
  - **worry about it**
    - » **keep track of what's been successfully communicated**
      - **receiver "acks"**
    - » **retransmit until**
      - **data is received**
    - or**
    - **it appears that "the network is down"**

# Reliability vs. Unreliability

- **Reliable communication**
    - good for
      - » email
      - » texting
      - » distributed file systems
      - » web pages
    - bad for
      - » streaming audio
      - » streaming video
- } a little noise is better than a long pause



# The Data Abstraction

- **Byte stream**
  - sequence of bytes
    - » as in pipes
  - any notion of a larger data aggregate is the responsibility of the programmer
- **Discrete records**
  - sequence of variable-size “records”
  - boundaries between records maintained
  - receiver receives discrete records, as sent by sender

# What's Supported

- **Stream**
  - byte-stream data abstraction
  - reliable transmission
- **Datagram**
  - discrete-record data abstraction
  - unreliable transmission

# Quiz 1

The following code is used to transmit data over a reliable byte-stream communication channel. Assume `sizeof(data)` is large.


```
// sender
record_t data=getData();
write(fd, &data,
    sizeof(data));
```

```
// receiver
read(fd, &data,
    sizeof(data));
useData(data);
```

Does it work?

- a) always
- b) always, assuming no network problems
- c) sometimes
- d) never

# Sockets



- You tell the system what you want by setting up the socket
- The system deals with all the other details

# Socket Parameters

- **Styles of communication:**
  - stream: reliable, two-way byte streams
  - datagram: unreliable, two-way record-oriented
  - and others
- **Communication domains**
  - **UNIX**
    - » endpoints (sockets) named with file-system pathnames
    - » supports stream and datagram
    - » trivial protocols: strictly for intra-machine use
  - **Internet**
    - » endpoints named with IP addresses
    - » supports stream and datagram
  - others
- **Protocols**
  - the means for communicating data
  - e.g., TCP/IP, UDP/IP

# Setting Things Up

- **Socket (communication endpoint) is given a name**
    - *bind* system call
  - **Datagram communication**
    - use *sendto* system call to send data to named recipient
    - use *recvfrom* system call to receive data and name of sender
  - **Stream communication**
    - client connects to server
      - » server uses *listen* and *accept* system calls to receive connections
      - » client uses *connect* system call to make connections
    - data transmitted using *send* or *write* system calls
    - data received using *recv* or *read* system calls
-

# Datagrams in the Internet Domain (1)

- **Steps**

- 1) **create socket**

```
int socket(int domain, int type,  
           int protocol);
```

```
fd = socket(AF_INET, SOCK_DGRAM, 0);
```

# Datagrams in the Internet Domain (2)

```
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;   /* internet address */
};

struct in_addr {
    uint32_t        s_addr;    /* address in network byte order */
};

struct sock_addr_in my_addr;

my_addr.sin_family = AF_INET;
inet_pton(AF_INET, "10.116.72.109", &my_addr.sin_addr.s_addr);
my_addr.sin_port = htons(3333);
```



# Datagrams in the Internet Domain (3)

## 2) bind name to socket

```
if (bind(fd, (struct sockaddr *)&my_addr,  
        sizeof(my_addr)) < 0) {  
    perror("bind");  
    exit(1);  
}
```

# Datagrams in the Internet Domain (4)

## 3) receive data

```
ssize_t recvfrom(int fd, void *buf,  
                ssize_t len,  
                int flags, struct sockaddr *from,  
                socklen_t *from_len);
```

```
struct sockaddr_in from_addr;  
int from_len = sizeof(from_addr);
```

```
recvfrom(fd, buf, sizeof(buf), 0,  
        (struct sockaddr *)&from_addr,  
        &from_len);
```

# Datagrams in the Internet Domain (5)

## 4) send data

```
ssize_t sendto(int fd, const void *buf,  
               ssize_t len, int flags,  
               const struct sockaddr *to,  
               socklen_t to_len);
```

```
sendto(fd, buf, sizeof(buf), 0,  
       (struct sockaddr *)&from_addr,  
       from_len);
```

# Quiz 2

Suppose a process on one machine sends a datagram to a process on another machine. The sender uses *sendto* and the receiver uses *recvfrom*. There's a momentary problem with the network and the datagram doesn't make it to the receiving process. Its call to *recvfrom*

- a) doesn't return
- b) returns -1 (indicating an error)
- c) returns 0
- d) returns some other value

# Using DNS

- Translate names to addresses using *getaddrinfo*
  - looks up name in DNS, gets list of possible addresses

# getaddrinfo()

- **int** getaddrinfo(  
    **const char** \*node,  
    **const char** \*service,  
    **const struct addrinfo** \*hints,  
    **struct addrinfo** \*\*res);

```
struct addrinfo {  
    int                ai_flags;  
    int                ai_family;  
    int                ai_socktype;  
    int                ai_protocol;  
    socklen_t          *ai_addrlen;  
    struct sockaddr    *ai_addr;  
    char               *ai_canonname;  
    struct addrinfo    *ai_next;  
};
```

# Using *getaddrinfo* (1)

```
struct addrinfo hints, **res, *rp;  
// zero out hints  
memset(&hints, 0, sizeof(hints));  
hints.ai_family = AF_INET;  
    // want IPv4  
hints.ai_socktype = SOCK_DGRAM;  
    // want datagram communication  
  
getaddrinfo("cslab1a.cs.brown.edu", "3333",  
            &hints, &res);
```

## Using *getaddrinfo* (2)

```
for (rp = res; rp != NULL; rp = rp->ai_next) {  
    // try each interface till we find one that works  
    if ((sock = socket(rp->ai_family, rp->ai_socktype,  
        rp->ai_protocol)) < 0) {  
        continue;  
    }  
    if (communicate(sock, ...)) // try using the socket  
        break; // it worked!  
    close(sock); // didn't work  
}  
if (rp == NULL) {  
    fprintf(stderr, "Could not contact %s\n", argv[1]);  
    exit(1);  
}  
  
freeaddrinfo(res); // free up storage allocated for list
```



# Client-Server Interaction

- **Client sends requests to server**
- **Server responds**
- **Server may deal with multiple clients at once**
- **Client may contact multiple servers**


# Reliable Communication

- **The promise ...**
  - what is sent is received
  - order is preserved
- **Set-up is required**
  - two parties agree to communicate
  - within the implementation of the protocol:
    - » each side keeps track of what is sent, what is received
    - » received data is acknowledged
    - » unack'd data is re-sent
- **The standard scenario**
  - server receives connection requests
  - client makes connection requests

# Streams in the Inet Domain (1)

- **Server steps**
  - 1) **create socket**


```
sfd = socket(AF_INET, SOCK_STREAM, 0);
```



# Streams in the Inet Domain (2)

- **Server steps**
  - 2) **bind name to socket**

```
bind(sfd,  
    (struct sockaddr *) &my_addr, sizeof(my_addr));
```



# Some Details ...

- **Server may have multiple interfaces; we want to be able to receive on all of them**

```
struct sockaddr_in {  
    sa_family_t sin_family;  
    in_port_t sin_port;  
    struct in_addr sin_addr;  
} my_addr;
```

```
my_addr.sin_family = AF_INET;  
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
my_addr.sin_port = htons(port);
```




**“Wildcard”  
address**

# Streams in the Inet Domain (3)

- **Server steps**

- 3) put socket in “listening mode”


```
int listen(int sfd, int MaxQueueLength);
```



# Streams in the Inet Domain (4)

- **Cient steps**
  - 1) **create socket**

```
cfd = socket(AF_INET, SOCK_STREAM, 0);
```




# Streams in the Inet Domain (5)

- **Client steps**
  - 2) bind name to socket

```
bind(cfd,  
    (struct sockaddr *) &my_addr, sizeof(my_addr));
```

**cfd**




**128.137.23.6:43**



# Streams in the Inet Domain (6)

- Client steps
  - 3) connect to server


```
connect(cfd, (struct sockaddr *)&server_addr,  
        sizeof(server_addr));
```



# Streams in the Inet Domain (7)

- **Server steps**
  - 4) **accept connection**

```
fd = accept((int)sfd, (struct sockaddr *)addr,  
            (int *)&addrlen);
```



# Inet Stream Example (1)

- **Server side**

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

int main(int argc, char *argv[ ]) {
    struct sockaddr_in my_addr;
    int lsock;
    void serve(int);
    if (argc != 2) {
        fprintf(stderr, "Usage: tcpServer port\n");
        exit(1);
    }
}
```

# Inet Stream Example (2)

```
// Step 1: establish a socket for TCP
if ((lsock = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket");
    exit(1);
}
```

# Inet Stream Example (3)

```
/* Step 2: set up our address */
memset(&my_addr, 0, sizeof(my_addr));
my_addr.sin_family = AF_INET;
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
my_addr.sin_port = htons(atoi(argv[1]));

/* Step 3: bind the address to our socket */
if (bind(lsock, (struct sockaddr *)&my_addr,
        sizeof(my_addr)) < 0) {
    perror("bind");
    exit(1);
}
```

# Inet Stream Example (4)

```
/* Step 4: put socket into "listening mode" */
if (listen(lsock, 100) < 0) {
    perror("listen");
    exit(1);
}
while (1) {
    int csock;
    struct sockaddr_in client_addr;
    int client_len = sizeof(client_addr);

    /* Step 5: receive a connection */
    csock = accept(lsock,
        (struct sockaddr *)&client_addr, &client_len);
    printf("Received connection from %s#%hu\n",
        inet_ntoa(client_addr.sin_addr), client_addr.sin_port);
```

# Inet Stream Example (5)

```
switch (fork( )) {  
  case -1:  
    perror("fork");  
    exit(1);  
  case 0:  
    // Step 6: create a new process to handle connection  
    serve(csock);  
    exit(0);  
  default:  
    close(csock);  
    break;  
}  
}  
}
```

# Inet Stream Example (6)

```
void serve(int fd) {  
    char buf[1024];  
    int count;  
  
    // Step 7: read incoming data from connection  
    while ((count = read(fd, buf, 1024)) > 0) {  
        write(1, buf, count);  
    }  
    if (count == -1) {  
        perror("read");  
        exit(1);  
    }  
    printf("connection terminated\n");  
}
```



# Inet Stream Example (7)

- **Client side**

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#include <string.h>
// + more includes ...

int main(int argc, char *argv[]) {
    int s, sock;
    struct addrinfo hints, *result, *rp;

    char buf[1024];
    if (argc != 3) {
        fprintf(stderr, "Usage: tcpClient host port\n");
        exit(1);
    }
```

# Inet Stream Example (8)

```
// Step 1: find the internet address of the server
memset(&hints, 0, sizeof(hints));
hints.ai_family = AF_UNSPEC;
hints.ai_socktype = SOCK_STREAM;

if ((s=getaddrinfo(argv[1], argv[2], &hints, &result)) != 0) {
    fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(s));
    exit(1);
}
```

# Inet Stream Example (9)

```
// Step 2: set up socket for TCP and connect to server
for (rp = result; rp != NULL; rp = rp->ai_next) {
    // try each interface till we find one that works
    if ((sock = socket(rp->ai_family, rp->ai_socktype,
        rp->ai_protocol)) < 0) {
        continue;
    }
    if (connect(sock, rp->ai_addr, rp->ai_addrlen) >= 0) {
        break;
    }
    close(sock);
}
if (rp == NULL) {
    fprintf(stderr, "Could not connect to %s\n", argv[1]);
    exit(1);
}
freeaddrinfo(result);
```

# Inet Stream Example (10)

```
// Step 3: send data to the server
while(fgets(buf, 1024, stdin) != 0) {
    if (write(sock, buf, strlen(buf)) < 0) {
        perror("write");
        exit(1);
    }
}
return 0;
}
```

# Quiz 3

**The previous slide contains**

```
write(sock, buf, strlen(buf))
```

**If data is lost and must be retransmitted**

- a) write returns an error so the caller can retransmit the data.**
- b) nothing happens as far as the application code is concerned, the data is retransmitted automatically.**

# Quiz 4

**A previous slide contains**

```
write(sock, buf, strlen(buf))
```

**We lose the connection to the other party (perhaps a network cable is cut).**

- a) write returns an error so the caller can reconnect, if desired.**
- b) nothing happens as far as the application code is concerned, the connection is reestablished automatically.**