


# CS 33

## Event-Based Programming Multithreaded Programming I

# Stream Relay



# Select System Call

```
int select(  
    int nfd,           // size of fd_sets  
    fd_set *readfds,   // descriptors of interest  
                        // for reading  
    fd_set *writefds,  // descriptors of interest  
                        // for writing  
    fd_set *excpfds,   // descriptors of interest  
                        // for exceptional events  
    struct timeval *timeout  
                        // max time to wait  
);
```

# Relay Sketch

```
void relay(int left, int right) {
    fd_set rd, wr;
    int maxFD = max(left, right) + 1;
    FD_ZERO(&rd); FD_SET(left, &rd); FD_SET(right, &rd);
    FD_ZERO(&wr); FD_SET(left, &wr); FD_SET(right, &wr);
    while (1) {
        select(maxFD, &rd, &wr, 0, 0);
        if (FD_ISSET(left, &rd))
            read(left, bufLR, BSIZE);
        if (FD_ISSET(right, &rd))
            read(right, bufRL, BSIZE);
        if (FD_ISSET(right, &wr))
            write(right, bufLR, BSIZE);
        if (FD_ISSET(left, &wr))
            write(left, bufRL, BSIZE);
    }
}
```

# Relay (1)

```
void relay(int left, int right) {  
    fd_set rd, wr;  
    int left_read = 1, right_write = 0;  
    int right_read = 1, left_write = 0;  
    int sizeLR, sizeRL, wret;  
    char bufLR[BSIZE], bufRL[BSIZE];  
    char *bufpR, *bufpL;  
    int maxFD = max(left, right) + 1;
```

# Relay (2)

```
while (1) {  
    FD_ZERO(&rd);  
    FD_ZERO(&wr);  
    if (left_read)  
        FD_SET(left, &rd);  
    if (right_read)  
        FD_SET(right, &rd);  
    if (left_write)  
        FD_SET(left, &wr);  
    if (right_write)  
        FD_SET(right, &wr);  
  
    select(maxFD, &rd, &wr, 0, 0);
```

# Relay (3)

```
if (FD_ISSET(left, &rd)) {  
    sizeLR = read(left, bufLR, BSIZE);  
    left_read = 0;  
    right_write = 1;  
    bufpR = bufLR;  
}  
if (FD_ISSET(right, &rd)) {  
    sizeRL = read(right, bufRL, BSIZE);  
    right_read = 0;  
    left_write = 1;  
    bufpL = bufRL;  
}
```

# Relay (4)

```
    if (FD_ISSET(right, &wr)) {
        if ((wret = write(right, bufpR, sizeLR)) == sizeLR) {
            left_read = 1; right_write = 0;
        } else {
            sizeLR -= wret; bufpR += wret;
        }
    }
    if (FD_ISSET(left, &wr)) {
        if ((wret = write(left, bufpL, sizeRL)) == sizeRL) {
            right_read = 1; left_write = 0;
        } else {
            sizeRL -= wret; bufpL += wret;
        }
    }
}
return 0;
}
```



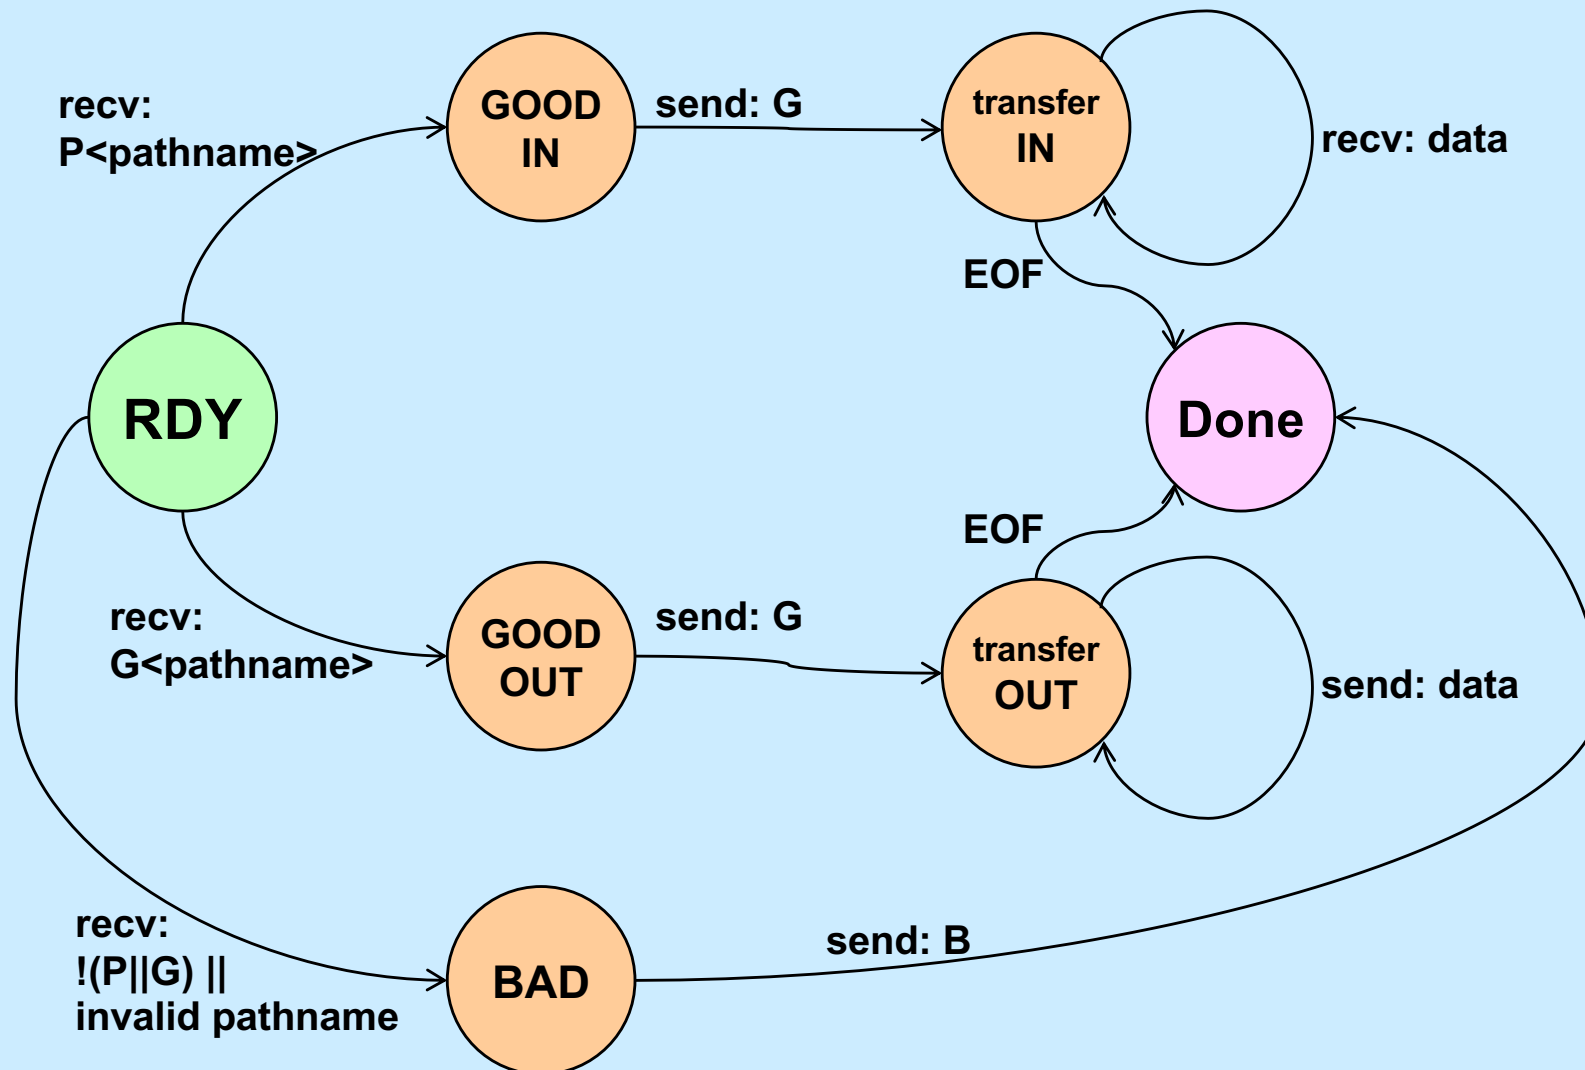
# A Really Simple Protocol

- **Transfer a file**
  - layered on top of TCP
    - » reliable
    - » indicates if connection is closed
- **To send a file**

P<null-terminated pathname><contents of file>
- **To retrieve a file**

G<null-terminated pathname>

# Server State Machine



# Keeping Track of State

```
typedef struct client {
    int fd;          // file descriptor of local file being transferred
    int size;        // size of out-going data in buffer
    char buf[BSIZE];
    enum state {RDY, BAD, GOOD, TRANSFER} state;
    /*
       states:
       RDY: ready to receive client's command (P or G)
       BAD: client's command was bad, sending B response + error msg
       GOOD: client's command was good, sending G response
       TRANSFER: transferring data
    */
    enum dir {IN, OUT} dir;
    /*
       IN: client has issued P command
       OUT: client has issued G command
    */
} client_t;
```

# Keeping Track of Clients

```
client_t clients[MAX_CLIENTS];  
for (i=0; i < MAX_CLIENTS; i++)  
    clients[i].fd = -1; // illegal value  
  
listen(lsock, max_queue_len);  
fd_set rd, wr;  
FD_ZERO(&rd);  
FD_SET(lsock, &rd);  
FD_ZERO(&wr);  
  
fd_set trd = rd;  
fd_set twr = wr;
```

# Main Server Loop

```
while(1) {
    select(maxfd, &trd, &twr, 0, 0);
    if (FD_ISSET(lsock, &trd)) {
        // a new connection
        new_client(lsock);
    }
    for (i=lsock+1; i<maxfd; i++) {
        if (FD_ISSET(i, &trd)) {
            // ready to read
            read_event(i);
        }
        if (FD_ISSET(i, &twr)) {
            // ready to write
            write_event(i);
        }
    }
    trd = rd; twr = wr;
}
```

---

# New Client

```
// Accept a new connection on listening socket  
// fd. Return the connected file descriptor
```

```
int new_client(int fd) {  
    int cfd = accept(fd, 0, 0);  
    clients[cfd].state = RDY;  
    FD_SET(cfd, &rd);  
    return cfd;  
}
```

# Read Event (1)

```
// File descriptor fd is ready to be read. Read it, then handle
// the input
void read_event(int fd) {
    client_t *c = &clients[fd];
    int ret = read(fd, c->buf, BSIZE);
    switch (c->state) {
    case RDY:
        if (c->buf[0] == 'G') {
            // GET request (to fetch a file)
            c->dir = OUT;
            if ((c->fd = open(&c->buf[1], O_RDONLY)) == -1) {
                // open failed; send negative response and error message
                c->state = BAD;
                c->buf[0] = 'B';
                strncpy(&c->buf[1], strerror(errno), BSIZE-2);
                c->buf[BSIZE-1] = 0;
                c->size = strlen(c->buf)+1;
            }
        }
    }
```

---

# Read Event (2)

```
else {  
    // open succeeded; send positive response  
    c->state = GOOD;  
    c->size = 1;  
    c->buf[0] = 'G';  
}  
// prepare to send response to client  
FD_SET(fd, &wr);  
FD_CLR(fd, &rd);  
break;  
}
```



# Read Event (3)

```
if (c->buf[0] == 'P') {
    // PUT request (to create a file)
    c->dir = IN;
    if ((c->fd = open(&c->buf[1],
        O_RDWR|O_CREAT|O_TRUNC, 0666)) == -1) {
        // open failed; send negative response and error message
        ...
    } else {
        // open succeeded; send positive response
        ...
    }
    // prepare to send response to client
    FD_SET(fd, &wr);
    FD_CLR(fd, &rd);
    break;
}
```

# Read Event (4)

```
case TRANSFER:
    // should be in midst of receiving file contents from client
    if (ret == 0) {
        // eof: all done
        close(c->fd);
        close(fd);
        FD_CLR(fd, &rd);
        break;
    }
    if (write(c->fd, c->buf, ret) == -1) {
        // write to file failed: terminate connection to client
        ...
        break;
    }
    // continue to read more data from client
    break;
}
```

# Write Event (1)

```
// File descriptor fd is ready to be written to. Write to it, then,  
// depending on current state, prepare for the next action.
```

```
void write_event(int fd) {  
    client_t *c = &clients[fd];  
    int ret = write(fd, c->buf, c->size);  
    if (ret == -1) {  
        // couldn't write to client; terminate connection  
        close(c->fd);  
        close(fd);  
        FD_CLR(fd, &wr);  
        c->fd = -1;  
        perror("write to client");  
        return;  
    }  
    switch (c->state) {
```

# Write Event (2)

**case** BAD:

```
// finished sending error message; now terminate client connection
close(c->fd);
close(fd);
FD_CLR(fd, &wr);
c->fd = -1;
break;
```

# Write Event (3)

**case** GOOD:

```
c->state = TRANSFER;
```

```
if (c->dir == IN) {
```

```
    // finished response to PUT request
```

```
    FD_SET(fd, &rd);
```

```
    FD_CLR(fd, &wr);
```

```
    break;
```

```
}
```

```
// otherwise finished response to GET request, so proceed
```

```
// to read file and start transfer out
```

```
// fd should remain in wr
```

# Write Event (4)

```
case TRANSFER:
    // should be in midst of transferring file contents to client
    if ((c->size = read(c->fd, c->buf, BSIZE)) == -1) {
        ...
        break;
    } else if (c->size == 0) {
        // no more file to transfer; terminate client connection
        close(c->fd);
        close(fd);
        FD_CLR(fd, &wr);
        c->fd = -1;
        break;
    }
    // continue to write more data to client
    break;
}
```

# Problems


- **Works fine as long as the protocol is followed correctly**
  - can client (malicious or incompetent) cause server to misbehave?
- **How can the server limit the number of clients?**
- **How does server limit file access?**

# Multithreaded Programming

- **A thread is a virtual processor**
  - an independent agent executing instructions
- **Multiple threads**
  - multiple independent agents executing instructions




# Why Threads?



- Many things are easier to do with threads
- Many things run faster with threads

# A Simple Example



# Life Without Threads

```
void relay(int left, int right) {
    fd_set rd, wr;
    int left_read = 1, right_write = 0;
    int right_read = 1, left_write = 0;
    int sizeLR, sizeRL, wret;
    char bufLR[BSIZE], bufRL[BSIZE];
    char *bufpR, *bufpL;
    int maxFD = max(left, right) + 1;

    fcntl(left, F_SETFL, O_NONBLOCK);
    fcntl(right, F_SETFL, O_NONBLOCK);

    while(1) {
        FD_ZERO(&rd);
        FD_ZERO(&wr);
        if (left_read)
            FD_SET(left, &rd);
        if (right_read)
            FD_SET(right, &rd);
        if (left_write)
            FD_SET(left, &wr);
        if (right_write)
            FD_SET(right, &wr);


        select(maxFD, &rd, &wr, 0, 0);
```

```
        if (FD_ISSET(left, &rd)) {
            sizeLR = read(left, bufLR, BSIZE);
            left_read = 0;
            right_write = 1;
            bufpR = bufLR;
        }
        if (FD_ISSET(right, &rd)) {
            sizeRL = read(right, bufRL, BSIZE);
            right_read = 0;
            left_write = 1;
            bufpL = bufRL;
        }
        if (FD_ISSET(right, &wr)) {
            if ((wret = write(right, bufpR, sizeLR)) == sizeLR) {
                left_read = 1; right_write = 0;
            } else {
                sizeLR -= wret; bufpR += wret;
            }
        }
        if (FD_ISSET(left, &wr)) {
            if ((wret = write(left, bufpL, sizeRL)) == sizeRL) {
                right_read = 1; left_write = 0;
            } else {
                sizeRL -= wret; bufpL += wret;
            }
        }
    }
    return 0;
}
```


# Life With Threads

```
void copy(int source, int destination) {  
    struct args *targs = args;  
    char buf[BSIZE];  
  
    while(1) {  
        int len = read(source, buf, BSIZE);  
        write(destination, buf, len);  
    }  
}
```


# Processes vs. Threads



**Process 1**




**Process 2**




**Process 3**


# Single-Threaded Database Server



# Multithreaded Database Server




# Single-Core Chips






# Dual-Core Chips



# Multi-Core Chips



# Good News/Bad News


## Good news

- multi-threaded programs can take advantage of multi-core chips (single-threaded programs cannot)

## Bad news

- it's not easy
  - » must have parallel algorithm
    - employing at least as many threads as processors
    - threads must keep processors busy
      - doing useful work

# Matrix Multiplication Revisited



# Standards

- **POSIX 1003.4a → 1003.1c → 1003.1j**
- **Microsoft**
  - **Win32/64**

# Creating Threads

```
long A[M][N], B[N][P], C[M][P];  
...  
for (i=0; i<M; i++)    // create worker threads  
    pthread_create(&thr[i], 0, matmult, i);  
  
...
```

```
void *matmult(void *arg) {  
    long i = (long)arg;  
    // compute row i of the product C of A and B  
    ...  
}
```

# When Is It Done?

```
long A[M][N], B[N][P], C[M][P];  
...  
for (i=0; i<M; i++)    // create worker threads  
    pthread_create(&thr[i], 0, matmult, i));  
  
for (i=0; i<M; i++)    // wait for termination  
    pthread_join(thr[i], 0);  
  
printResult(C); // shouldn't do this until  
                // workers have terminated
```

# Example (1)

```
#include <stdio.h>
#include <pthread.h>
#include <string.h>
```

```
#define M    3
#define N    4
#define P    5
```

```
long A[M][N];
long B[N][P];
long C[M][P];
```

```
void *matmult(void *);
```

```
main( ) {
    long i;
    pthread_t thr[M];
    int error;

    // initialize the matrices
    ...
}
```



## Example (2)

```
for (i=0; i<M; i++) { // create worker threads
    if (error = pthread_create(
        &thr[i],
        0,
        matmult,
        (void *)i)) {
        fprintf(stderr, "pthread_create: %s", strerror(error));
        exit(1);
    }
}

for (i=0; i<M; i++) // wait for workers to finish their jobs
    pthread_join(thr[i], 0)

/* print the results ... */
}
```

## Example (3)

```
void *matmult(void *arg) {  
    long row = (long) arg;  
    long col;  
    long i;  
    long t;  
  
    for (col=0; col < P; col++) {  
        t = 0;  
        for (i=0; i<N; i++)  
            t += A[row][i] * B[i][col];  
        C[row][col] = t;  
    }  
    return (0);  
}
```

# Compiling It

```
% gcc -o mat mat.c -pthread
```

# Termination

```
pthread_exit((void *) value);
```

```
return((void *) value);
```

```
pthread_join(thread, (void **) &value);
```

# Detached Threads

```
start_servers( ) {  
    pthread_t thread;  
    int i;  
  
    for (i=0; i<nr_of_server_threads; i++) {  
        pthread_create(&thread, 0, server, 0);  
        pthread_detach(thread);  
    }  
    ...  
}  
  
void *server(void * arg) {  
    ...  
}
```

# Complications

```
void relay(int left, int right) {  
    pthread_t LRthread, RLthread;  
  
    pthread_create(&LRthread,  
        0,  
        copy,  
        left, right);           // Can't do this ...  
    pthread_create(&RLthread,  
        0,  
        copy,  
        right, left);           // Can't do this ...  
}
```

# Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```

```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```

# Multiple Arguments

```
typedef struct args {  
    int src;  
    int dest;  
} args_t;
```

```
void relay(int left, int right) {  
    args_t LRargs, RLargs;  
    pthread_t LRthread, RLthread;  
    ...  
    pthread_create(&LRthread, 0, copy, &LRargs);  
    pthread_create(&RLthread, 0, copy, &RLargs);  
}
```


## Quiz 1

**Does this work?**

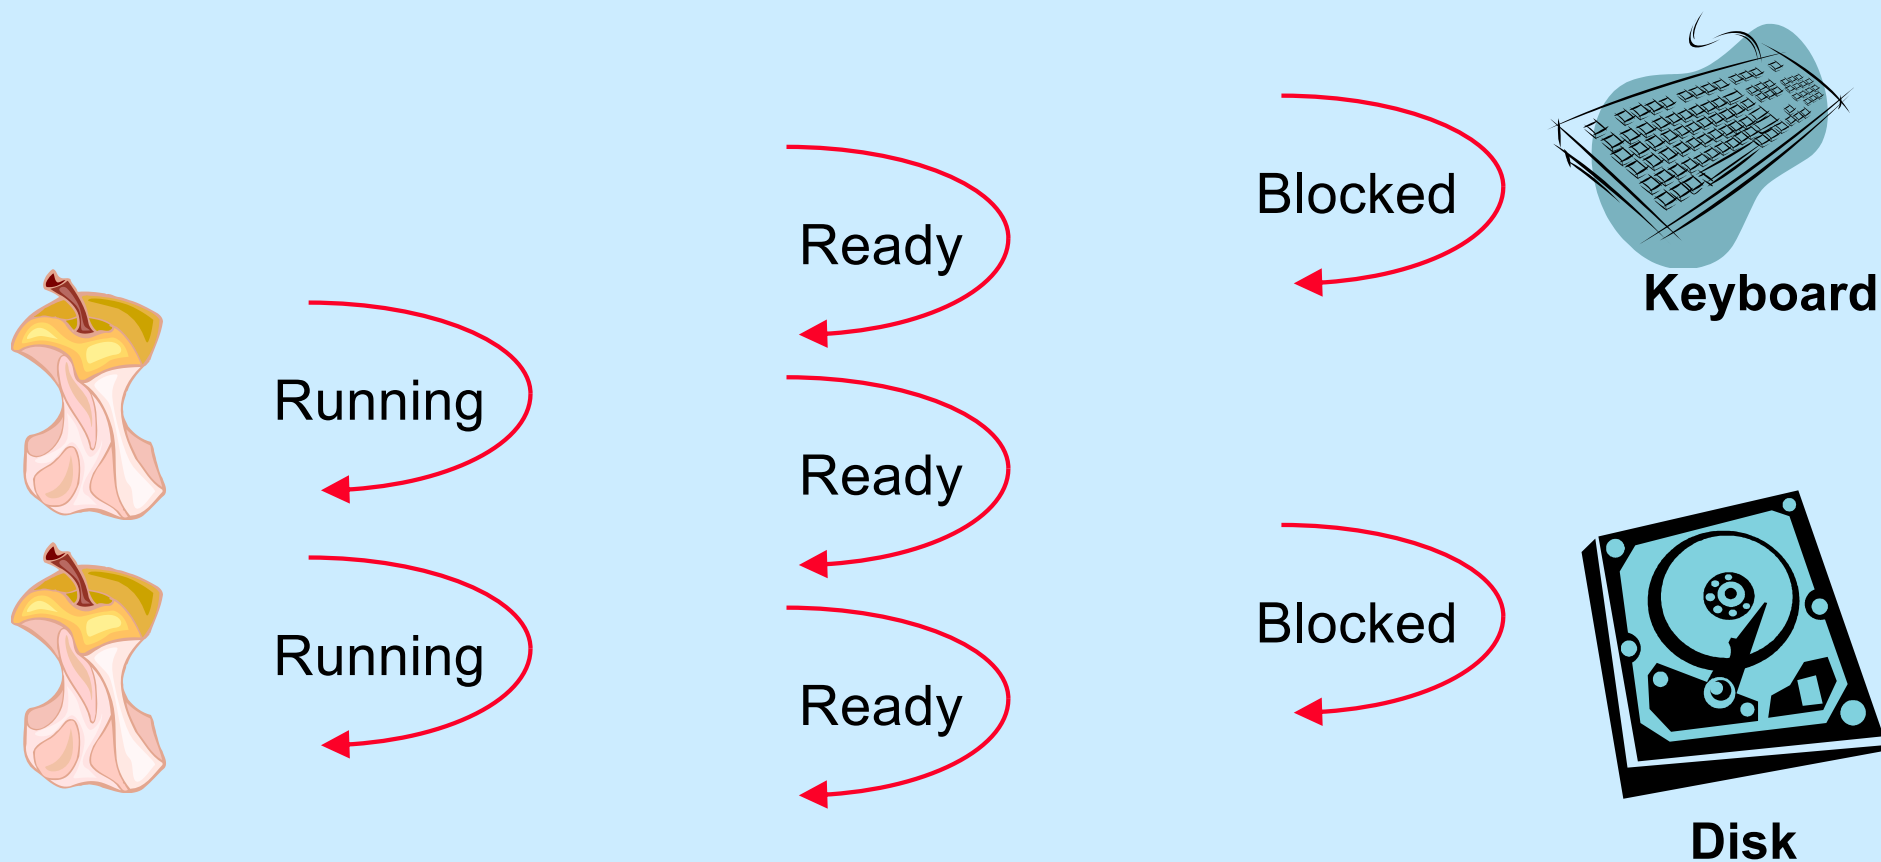
- a) yes**
- b) no**



# Execution



# Multiplexing Processors



# Quiz 2

```
pthread_create(&tid, 0, tproc, (void *)1);  
pthread_create(&tid, 0, tproc, (void *)2);
```

```
printf("T0\n");
```

```
...
```

```
void *tproc(void *arg) {  
    printf("T%d\n", (long) arg);  
    return 0;  
}
```

**In which order are things printed?**

- a) T0, T1, T2
- b) T1, T2, T0
- c) T2, T1, T0
- d) indeterminate

# Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return 0;
}
```

# Cost of Threads

```
int main(int argc, char *argv[]) {
    ...
    val = niters/nthreads;

    for (i=0; i<nthreads; i++)
        pthread_create(&thread, 0, work, (void *)val);
    pthread_exit(0);
    return 0;
}

void *work(void *arg) {
    long n = (long)arg; int i, j; volatile long x;

    for (i=0; i<n; i++) {
        x = 0;
        for (j=0; j<1000; j++)
            x = x*j;
    }
    return 0;
}
```

## Quiz 3

This code runs in time  $n$  on a 4-core processor when  $nthreads$  is 8. It runs in time  $p$  on the same processor when  $nthreads$  is 400.

- a)  $n \ll p$  (slower)
- b)  $n \approx p$  (same speed)
- c)  $n \gg p$  (faster)

# Problem

```
pthread_create(&thread, 0, start, 0);
```

```
...
```

```
void *start(void *arg) {  
    long BigArray[128*1024*1024];  
    ...  
    return 0;  
}
```

# Thread Attributes

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
  
...  
  
/* establish some attributes */  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```

# Stack Size

```
pthread_t thread;  
pthread_attr_t thr_attr;  
  
pthread_attr_init(&thr_attr);  
pthread_attr_setstacksize(&thr_attr, 130*1024*1024);  
  
...  
  
pthread_create(&thread, &thr_attr, startroutine, arg);  
  
...
```