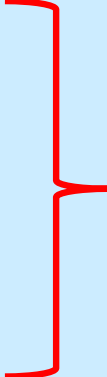


CS 33

Final Words

Malloc and Threads

- Multiple threads
 - One heap
- 
- Bottleneck?**

Solution 1

- **Divvy up the heap among the threads**
 - each thread has its own heap
 - no mutexes required
 - no bottleneck
- **How much heap does each thread get?**
- **What if one thread mallocs something and another thread frees it?**

Solution 2

- **Multiple “arenas”**
 - each with its own mutex
 - thread allocates from the first one it can find whose mutex was unlocked
 - » if none, then creates new one
 - deallocations go back to original arena

Solution 3

- **Global heap plus per-thread heaps**
 - threads pull storage from global heap
 - freed storage goes to per-thread heap
 - » unless things are imbalanced
 - then thread moves storage back to global heap
 - mutex on only the global heap
- **What if one thread allocates and another frees storage?**

Malloc/Free Implementations

- **ptmalloc**
 - based on solution 2
 - in glibc (i.e., used by default)
- **tcmalloc**
 - based on solution 3
 - from Google
- **Which is best?**

Test Program

```
const unsigned int N=64, nthreads=32, iters=100000000;
int main() {
    void *tfunc(void *);
    pthread_t thread[nthreads];
    for (int i=0; i<nthreads; i++) {
        pthread_create(&thread[i], 0, tfunc, (void *)i);
        pthread_detach(thread[i]);
    }
    pthread_exit(0);
}

void *tfunc(void *arg) {
    long i;
    for (i=0; i<iters; i++) {
        long *p = (long *)malloc(sizeof(long) * ((i%N)+1));
        free(p);
    }
    return 0;
}
```

Compiling It ...

```
% gcc -o ptalloc alloc.cc -lpthread
```

```
% gcc -o tcalloc alloc.cc -lpthread -ltcmalloc
```


Running It (2014) ...

```
$ time ./ptalloc
real    0m5.142s
user    0m20.501s
sys     0m0.024s
$ time ./tcalloc
real    0m1.889s
user    0m7.492s
sys     0m0.008s
```

What's Going On?

```
$ strace -c -f ./ptalloc
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
100.00	0.040002	13	3007	520	futex

```
...
```

```
$ strace -c -f ./tcalloc
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
0.00	0.000000	0	59	13	futex

```
...
```

Test Program 2, part 1

```
#define N 64
#define npairs 16
#define allocsPerIter 1024
const long iters = 8*1024*1024/allocsPerIter;
#define BufSize 10240
typedef struct buffer {
    int *buf[BufSize];
    unsigned int nextin;
    unsigned int nextout;
    sem_t empty;
    sem_t occupied;
    pthread_t pthread;
    pthread_t cthread;
} buffer_t;
```

Test Program 2, part 2

```
int main() {
    long i;
    buffer_t b[npairs];
    for (i=0; i<npairs; i++) {
        b[i].nextin = 0;
        b[i].nextout = 0;
        sem_init(&b[i].empty, 0, BufSize/allocsPerIter);
        sem_init(&b[i].occupied, 0, 0);
        pthread_create(&b[i].pthread, 0, prod, &b[i]);
        pthread_create(&b[i].cthread, 0, cons, &b[i]);
    }
    for (i=0; i<npairs; i++) {
        pthread_join(b[i].pthread, 0);
        pthread_join(b[i].cthread, 0);
    }
    return 0;
}
```

Test Program 2, part 3

```
void *prod(void *arg) {
    long i, j;
    buffer_t *b = (buffer_t *)arg;
    for (i = 0; i<iters; i++) {
        sem_wait(&b->empty);
        for (j = 0; j<allocsPerIter; j++) {
            b->buf[b->nextin] = malloc(sizeof(int) * ((j%N)+1));
            if (++b->nextin >= BufSize)
                b->nextin = 0;
        }
        sem_post(&b->occupied);
    }
    return 0;
}
```

Test Program 2, part 4

```
void *cons(void *arg) {
    long i, j;
    buffer_t *b = (buffer_t *)arg;
    for (i = 0; i<iters; i++) {
        sem_wait(&b->occupied);
        for (j = 0; j<allocsPerIter; j++) {
            free(b->buf[b->nextout]);
            if (++b->nextout >= BufSize)
                b->nextout = 0;
        }
        sem_post(&b->empty);
    }
    return 0;
}
```

Running It (2014) ...

```
$ time ./ptalloc2
real    0m1.087s
user    0m3.744s
sys     0m0.204s
$ time ./tcalloc2
real    0m3.535s
user    0m11.361s
sys     0m2.112s
```

What's Going On?

```
$ strace -c -f ./ptalloc2
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

-----	-----	-----	-----	-----	-----
-------	-------	-------	-------	-------	-------

94.96	2.347314	44	53653	14030	futex
-------	----------	----	-------	-------	-------

```
...
```

```
$ strace -c -f ./tcalloc2
```

```
...
```

% time	seconds	usecs/call	calls	errors	syscall
--------	---------	------------	-------	--------	---------

-----	-----	-----	-----	-----	-----
-------	-------	-------	-------	-------	-------

93.86	6.604632	36	185731	45222	futex
-------	----------	----	--------	-------	-------

```
...
```


Running it (2015) ...

```
sphere $ time ./ptalloc
```

```
real      0m2.373s
```

```
user      0m9.152s
```

```
sys       0m0.008s
```

```
sphere $ time ./tcalloc
```

```
real      0m4.868s
```

```
user      0m19.444s
```

```
sys       0m0.020s
```

Running it (2015) ...

```
kui $ time ./ptalloc
```

```
real      0m2.787s
```

```
user      0m11.045s
```

```
sys       0m0.004s
```

```
kui $ time ./tcalloc
```

```
real      0m1.701s
```

```
user      0m6.584s
```

```
sys       0m0.004s
```

Running it (2015) ...

```
cslab0a $ time ./ptalloc
```

```
real      0m2.234s
```

```
user      0m8.468s
```

```
sys       0m0.000s
```

```
cslab0a $ time ./tcalloc
```

```
real      0m4.938s
```

```
user      0m19.584s
```

```
sys       0m0.000s
```

What's Going On?

- On kui:
 - `libtcmalloc.so` -> `libtcmalloc.so.4.1.0`
- On other machines:
 - `libtcmalloc.so` -> `libtcmalloc.so.4.2.2`

However (2015) ...

```
cslab0a $ time ./ptalloc2
```

```
real      0m0.466s
```

```
user      0m1.504s
```

```
sys       0m0.212s
```

```
cslab0a $ time ./tcalloc2
```

```
real      0m1.516s
```

```
user      0m5.212s
```

```
sys       0m0.328s
```

It's 2019

- **tcmalloc no longer exists**
 - no explanation from Google, it's simply gone
- **ptmalloc continues to improve**

Thread Scheduling

- **The OS multiplexes threads on the available processors/cores**
 - **share the processors equally**
 - » **time slicing: each thread gets a fixed amount of time before it's forced to yield the processor to another thread (if there is one)**
 - **some threads are more important than others**
 - » **priorities: higher-priority threads get the processor in preference to lower-priority threads**


A Scheduling Issue

- You and four friends each contribute \$1000 towards a server
 - you, rightfully, feel you own 20% of it
- Your friends are into threads, you're not
 - they run 5-threaded programs
 - you run a 1-threaded program
- The scheduler treats all threads equally
- Their programs each get $5/21$ of the processor
- Your programs get $1/21$ of the processor
 - (you should have paid more attention to the fractal threads lab)

Lottery Scheduling

- **25 lottery tickets are distributed equally to you and your four friends**
 - you give 5 tickets to your one thread
 - they give one ticket each to their threads
- **A lottery is held for every scheduling decision**
 - your thread is 5 times more likely to win than the others

Metered Processors




Algorithm

- **Each thread has a meter, which runs only when the thread is running on the processor**
- **At every clock tick**
 - **give processor to thread that's had the least processor time as shown on its meter**
 - **in case of tie, thread with lowest ID wins**

Issue

- **Some threads may be more important than others**

Metered Processors (RI Variation)



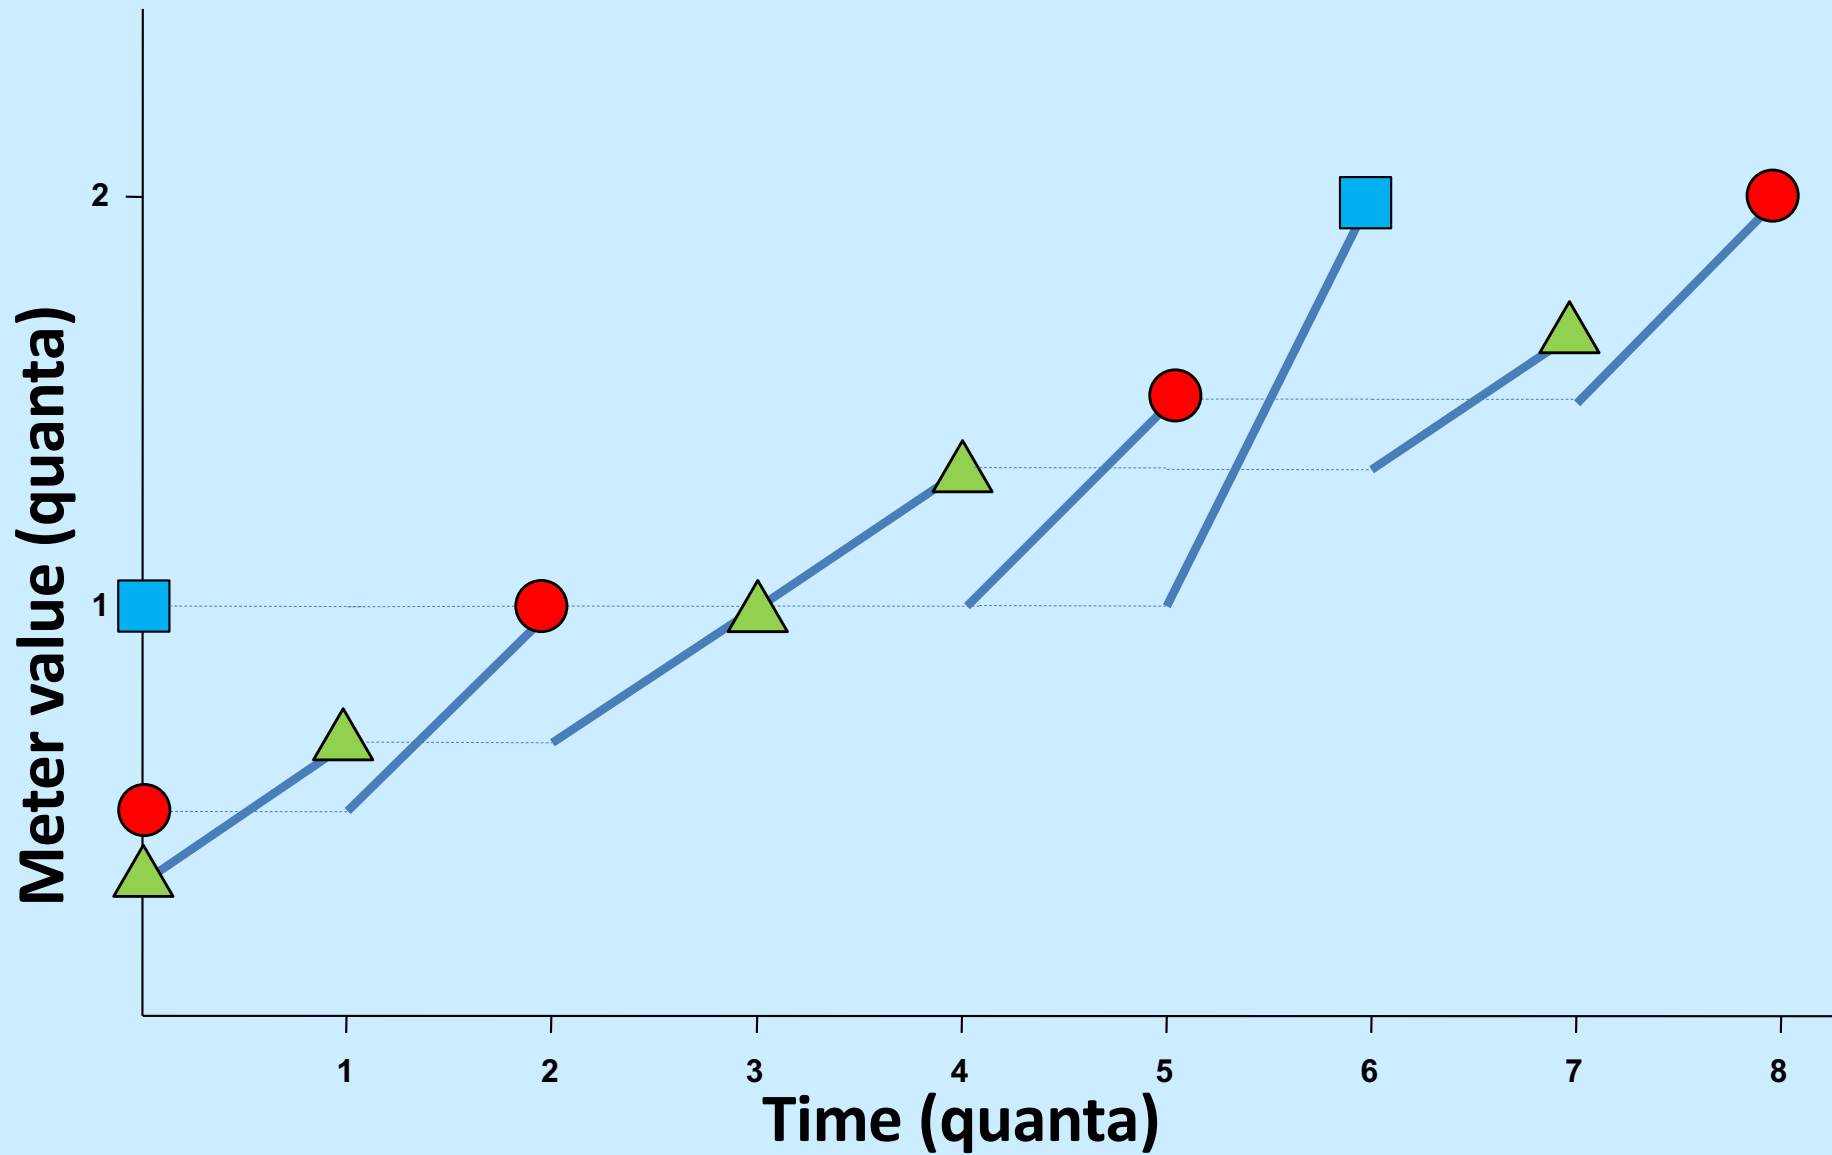
Details ...

- **Each thread pays a bribe**
 - the greater the bribe, the slower the meter runs
 - to simplify bribing, you buy “tickets”
 - » one ticket is required to get a fair meter
 - » two tickets get a meter running at half speed
 - » three tickets get a meter running at 1/3 speed
 - » etc.

New Algorithm

- Each thread has a (*possibly crooked*) meter, which runs only when the thread is running on the processor
- At every clock tick
 - give processor to thread that's had the least processor time as shown on its meter
 - in case of tie, thread with lowest ID wins

Example



You'll Soon Finish CS 33 ...

- You might
 - celebrate



- take another systems course

- » 32
 - » 131
 - » 138
 - » 166
 - » 167



- become a 33 TA



Systems Courses Next Semester

- **CS 32 (Intro to Software Engineering)**
 - you've mastered low-level systems programming
 - now do things at a higher level
 - learn software-engineering techniques using Java, XML, etc.
- **CS 131 (Fundamentals of Computer Systems)**
 - an overview of how computer systems work
- **CS 138 (Distributed Systems)**
 - you now know how things work on one computer
 - what if you've got lots of computers?
 - some may have crashed, others may have been taken over by your worst (and smartest) enemy
- **CS 166 (Computer Systems Security)**
 - liked buffer?
 - you'll really like 166
- **CS 167/169 (Operating Systems)**
 - still mystified about what the OS does?
 - write your own!

Critical Review

- **Do it online**

- https://brown.co1.qualtrics.com/jfe/form/SV_cOBf7p7go_cbkCI5
- password: KLFDS

The End

Well, not quite ...

Database is due on 12/13.

The TAs and I will hold hours all this week.

I'll hold hours 3-4 today, 2-4 Wednesday, 2-5 Friday

Happy coding and happy holidays!